fit@hcmus

University of Science – VNU-HCM
Faculty of Information Technology
CSC15105 – Text Mining and Application

Slot 01 - Introduction to Text Mining

Presenter:

Dr. LE Thanh Tung

Content

1. Text Mining
2. Machine Learning
3. Evaluation Metrics
4. Programming Language

Introduction

- Text mining is the process of transforming unstructured **text** into a structured format to identify meaningful patterns and new insights

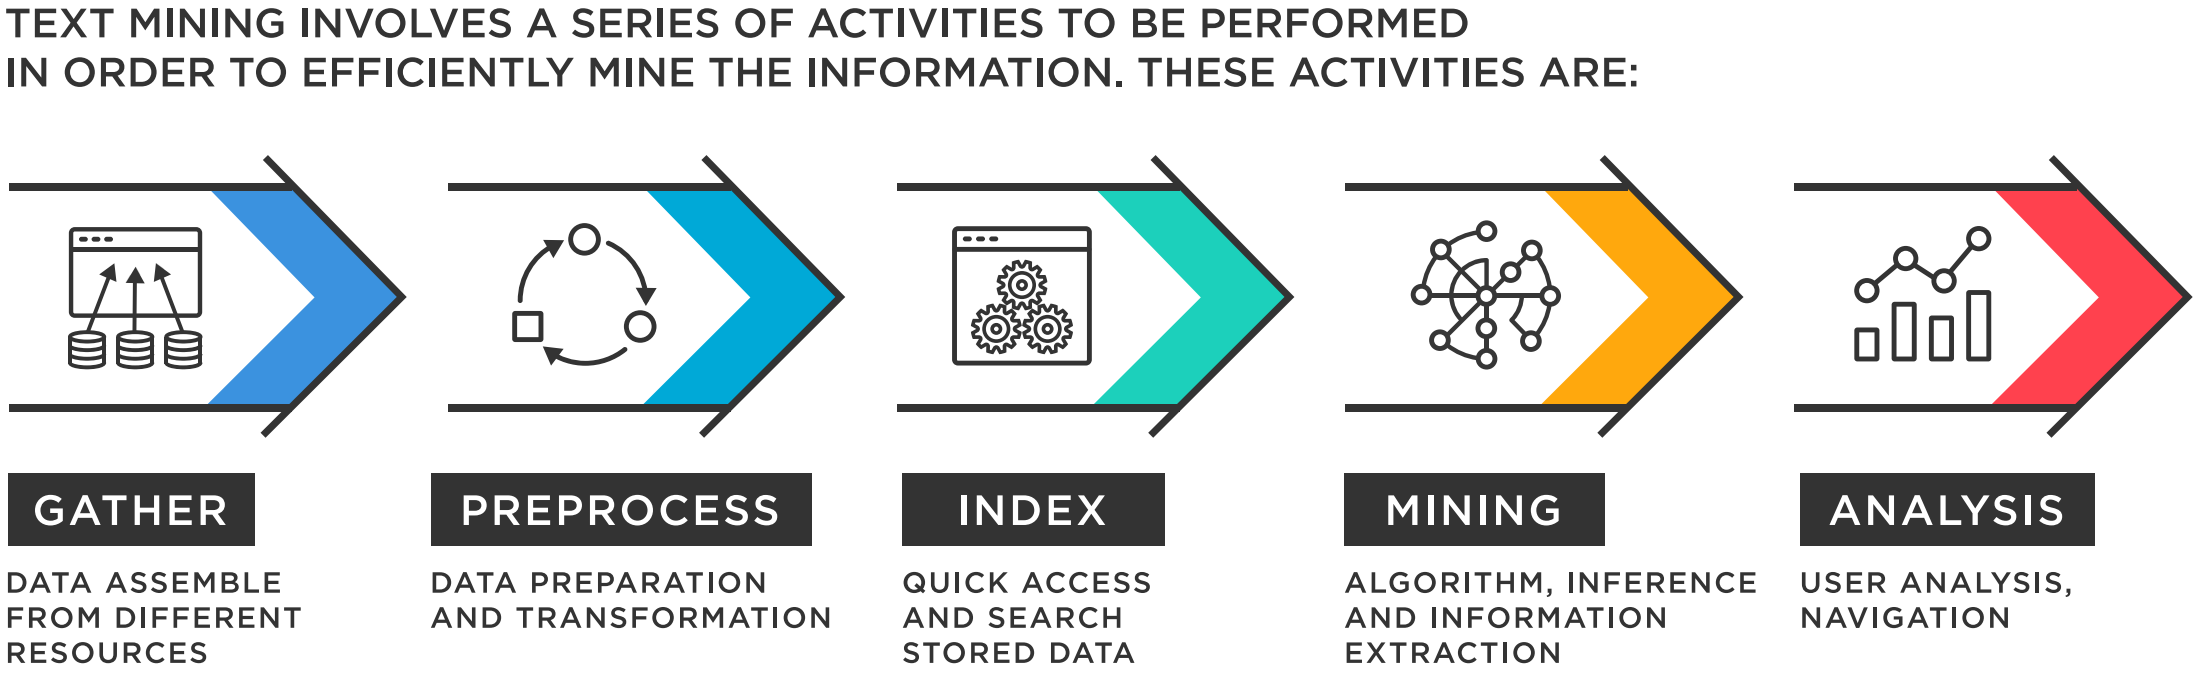

Introduction

- Gupta & Lehal (2009) have regarded text mining as new interdisciplinary area which is an amalgamation of data mining, information retrieval, machine learning, computer linguistic and statistics

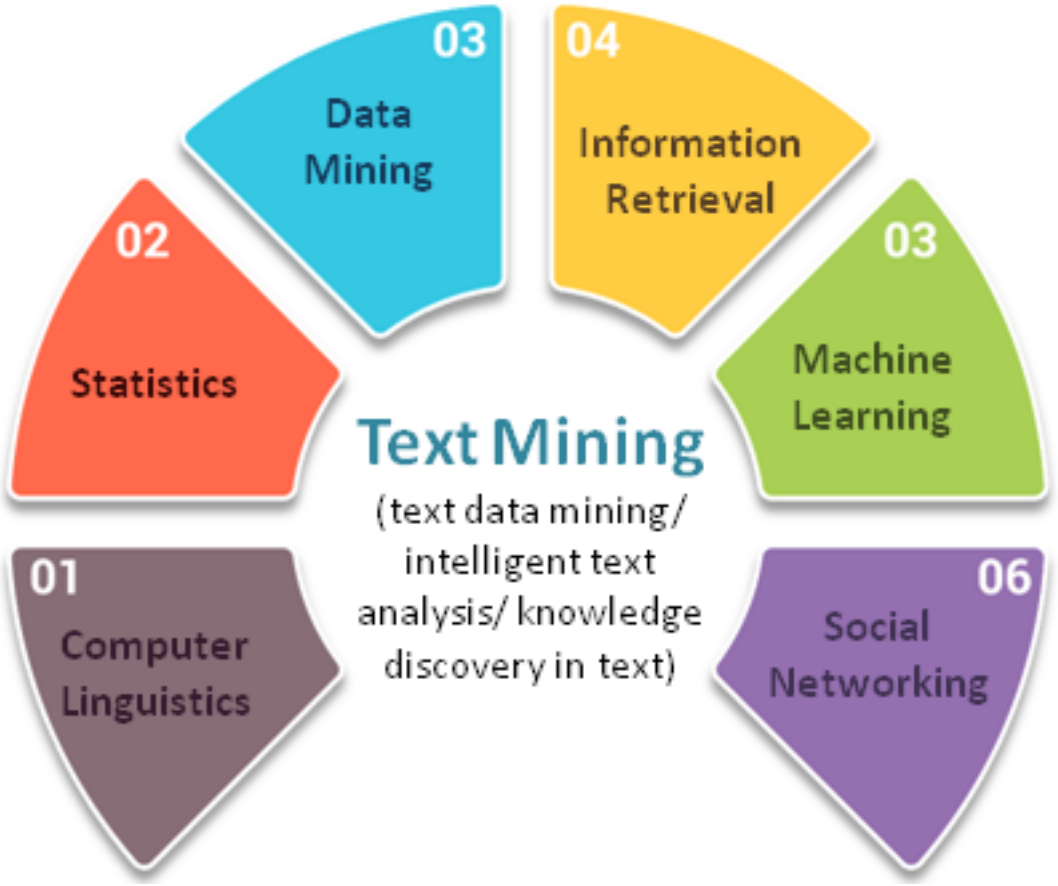

Introduction

- Machine Learning is a discipline of artificial intelligence (AI) that provides machines with the ability to automatically **learn from data and past experiences** while identifying patterns to make predictions with minimal human intervention

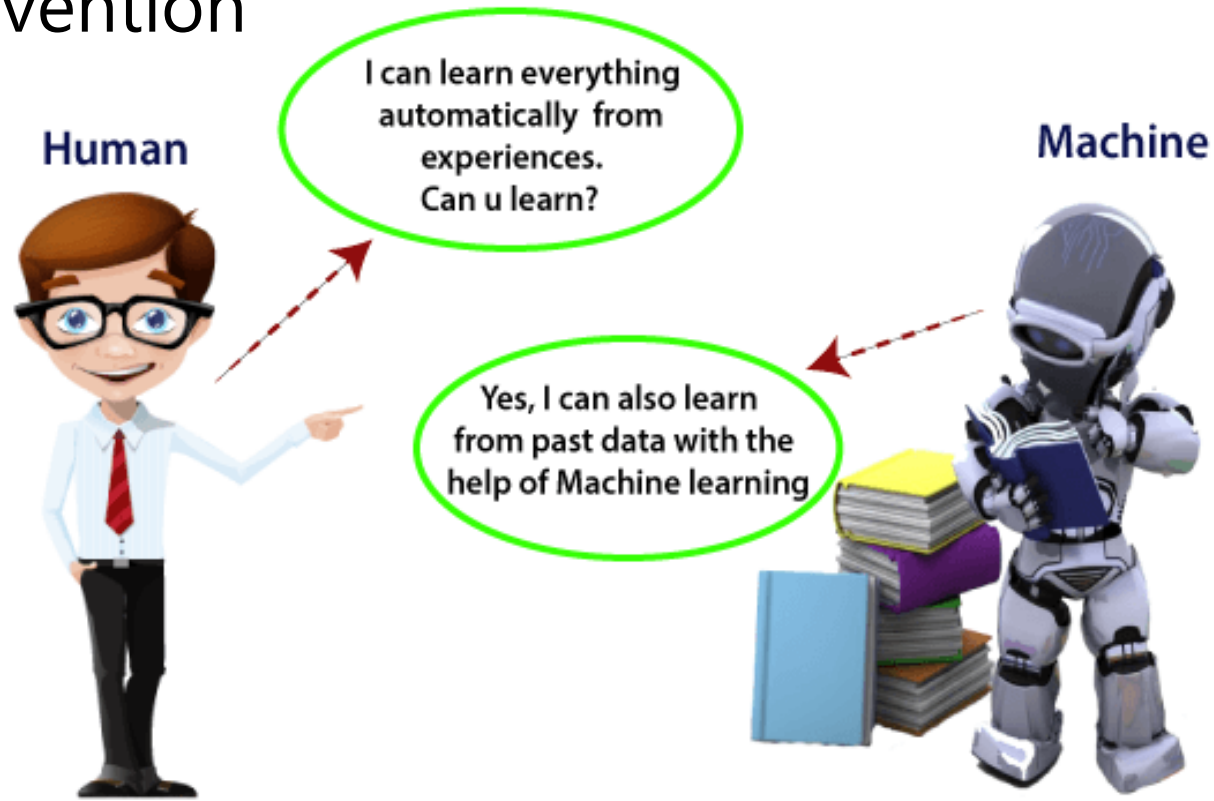

Introduction

- Data mining is the process of analyzing a large batch of information to **discern trends and patterns**

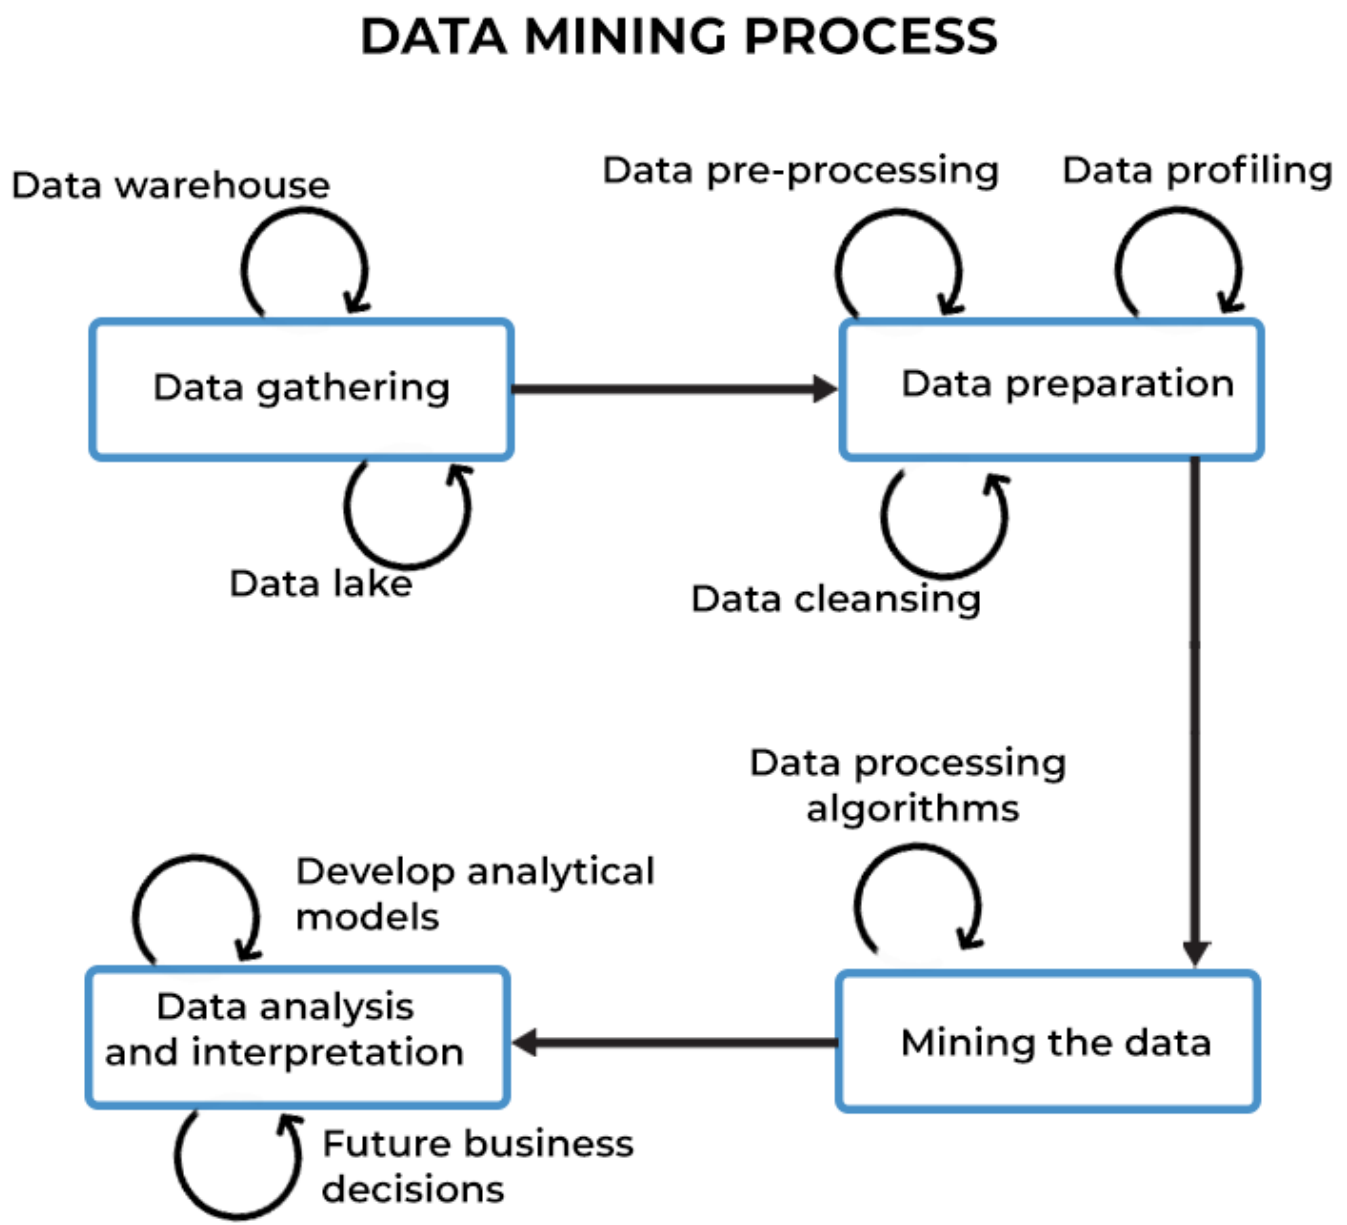

Introduction

- Information Retrieval is finding material (usually documents) of an unstructured nature (usually text) that satisfies an information need from within large collections (usually stored on computers)

- Two main approaches are matching words in the query against the database index (keyword searching) and traversing the database using hypertext or hypermedia links

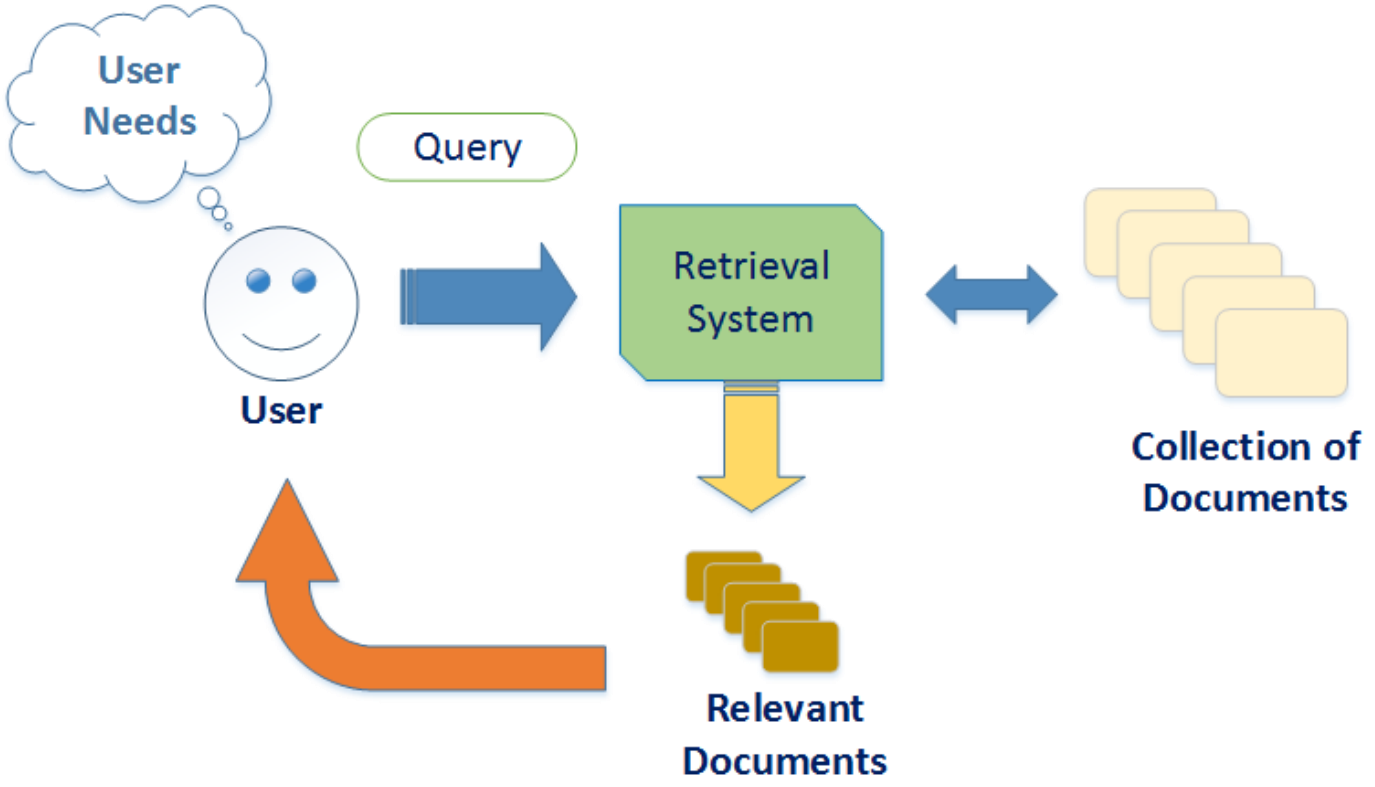

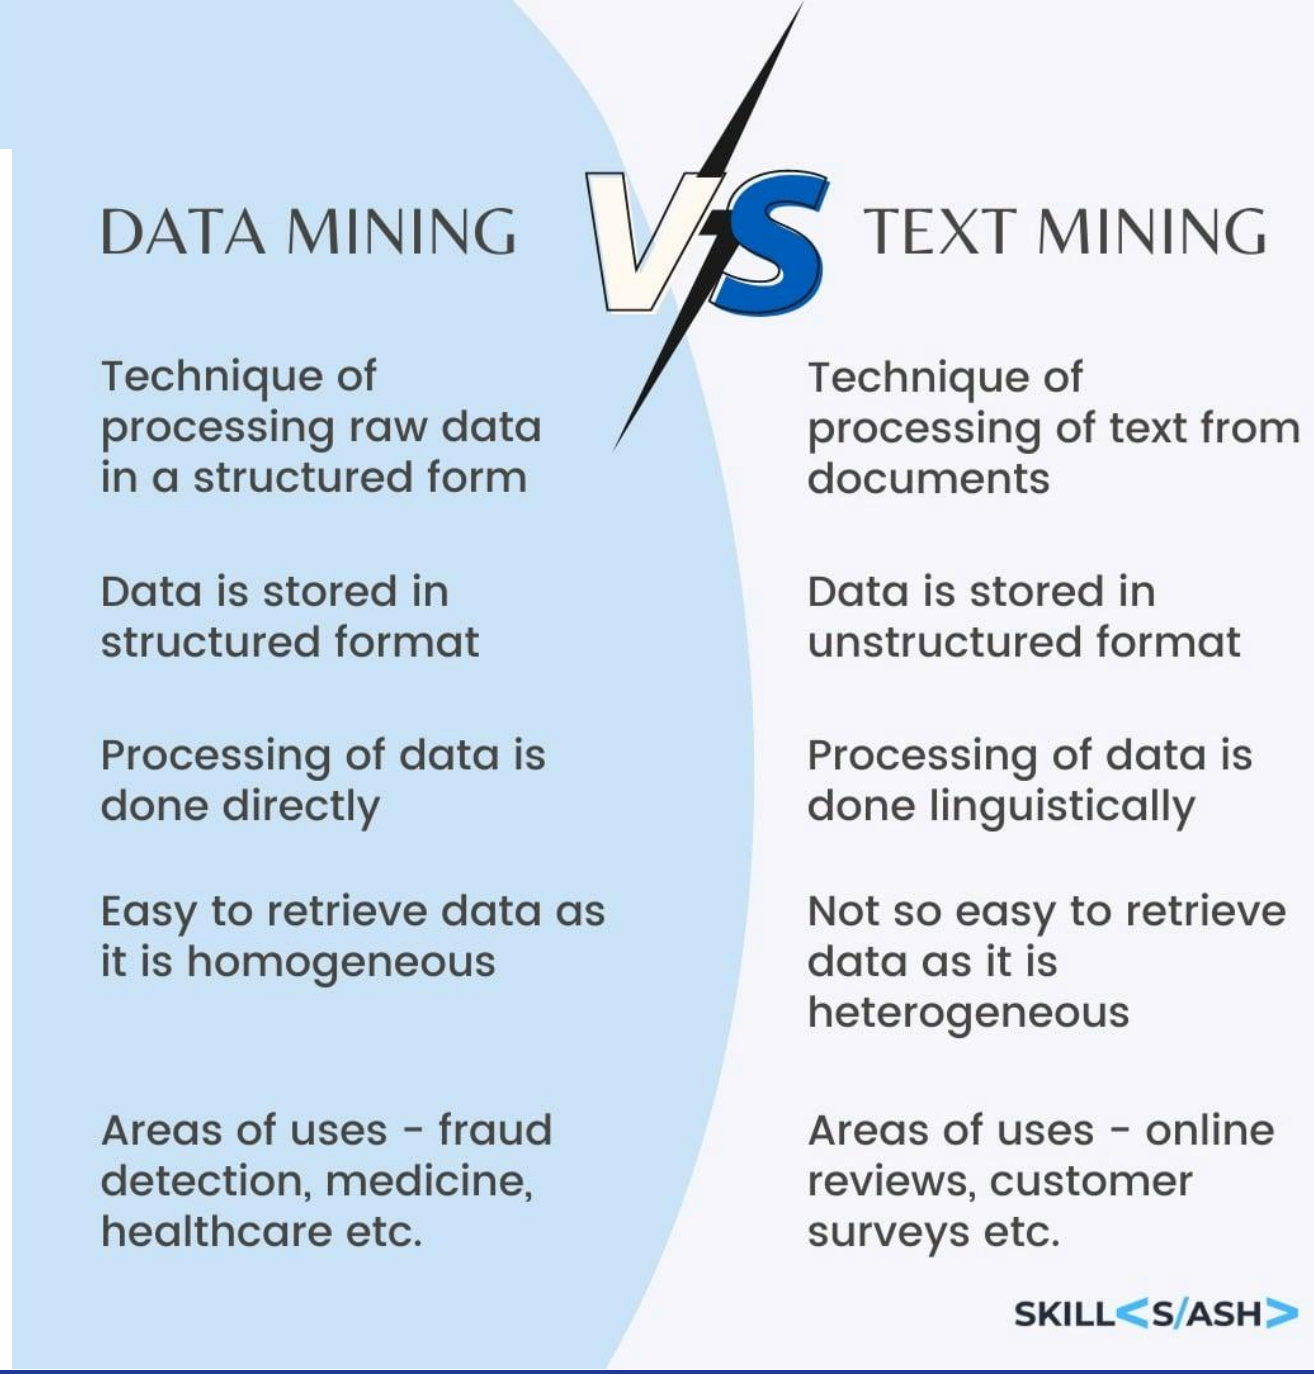
DATA MINING
VS
TEXT MINING
Technique of processing raw data in a structured form
Technique of processing of text from documents
Data is stored in structured format
Data is stored in unstructured format
Processing of data is done directly
Processing of data is done linguistically
Easy to retrieve data as it is homogeneous
Not so easy to retrieve data as it is heterogeneous
Areas of uses - fraud detection, medicine, healthcare etc.
Areas of uses - online reviews, customer surveys etc.
SKILL<S/ASH>

Text Analytics Application

General Architecture

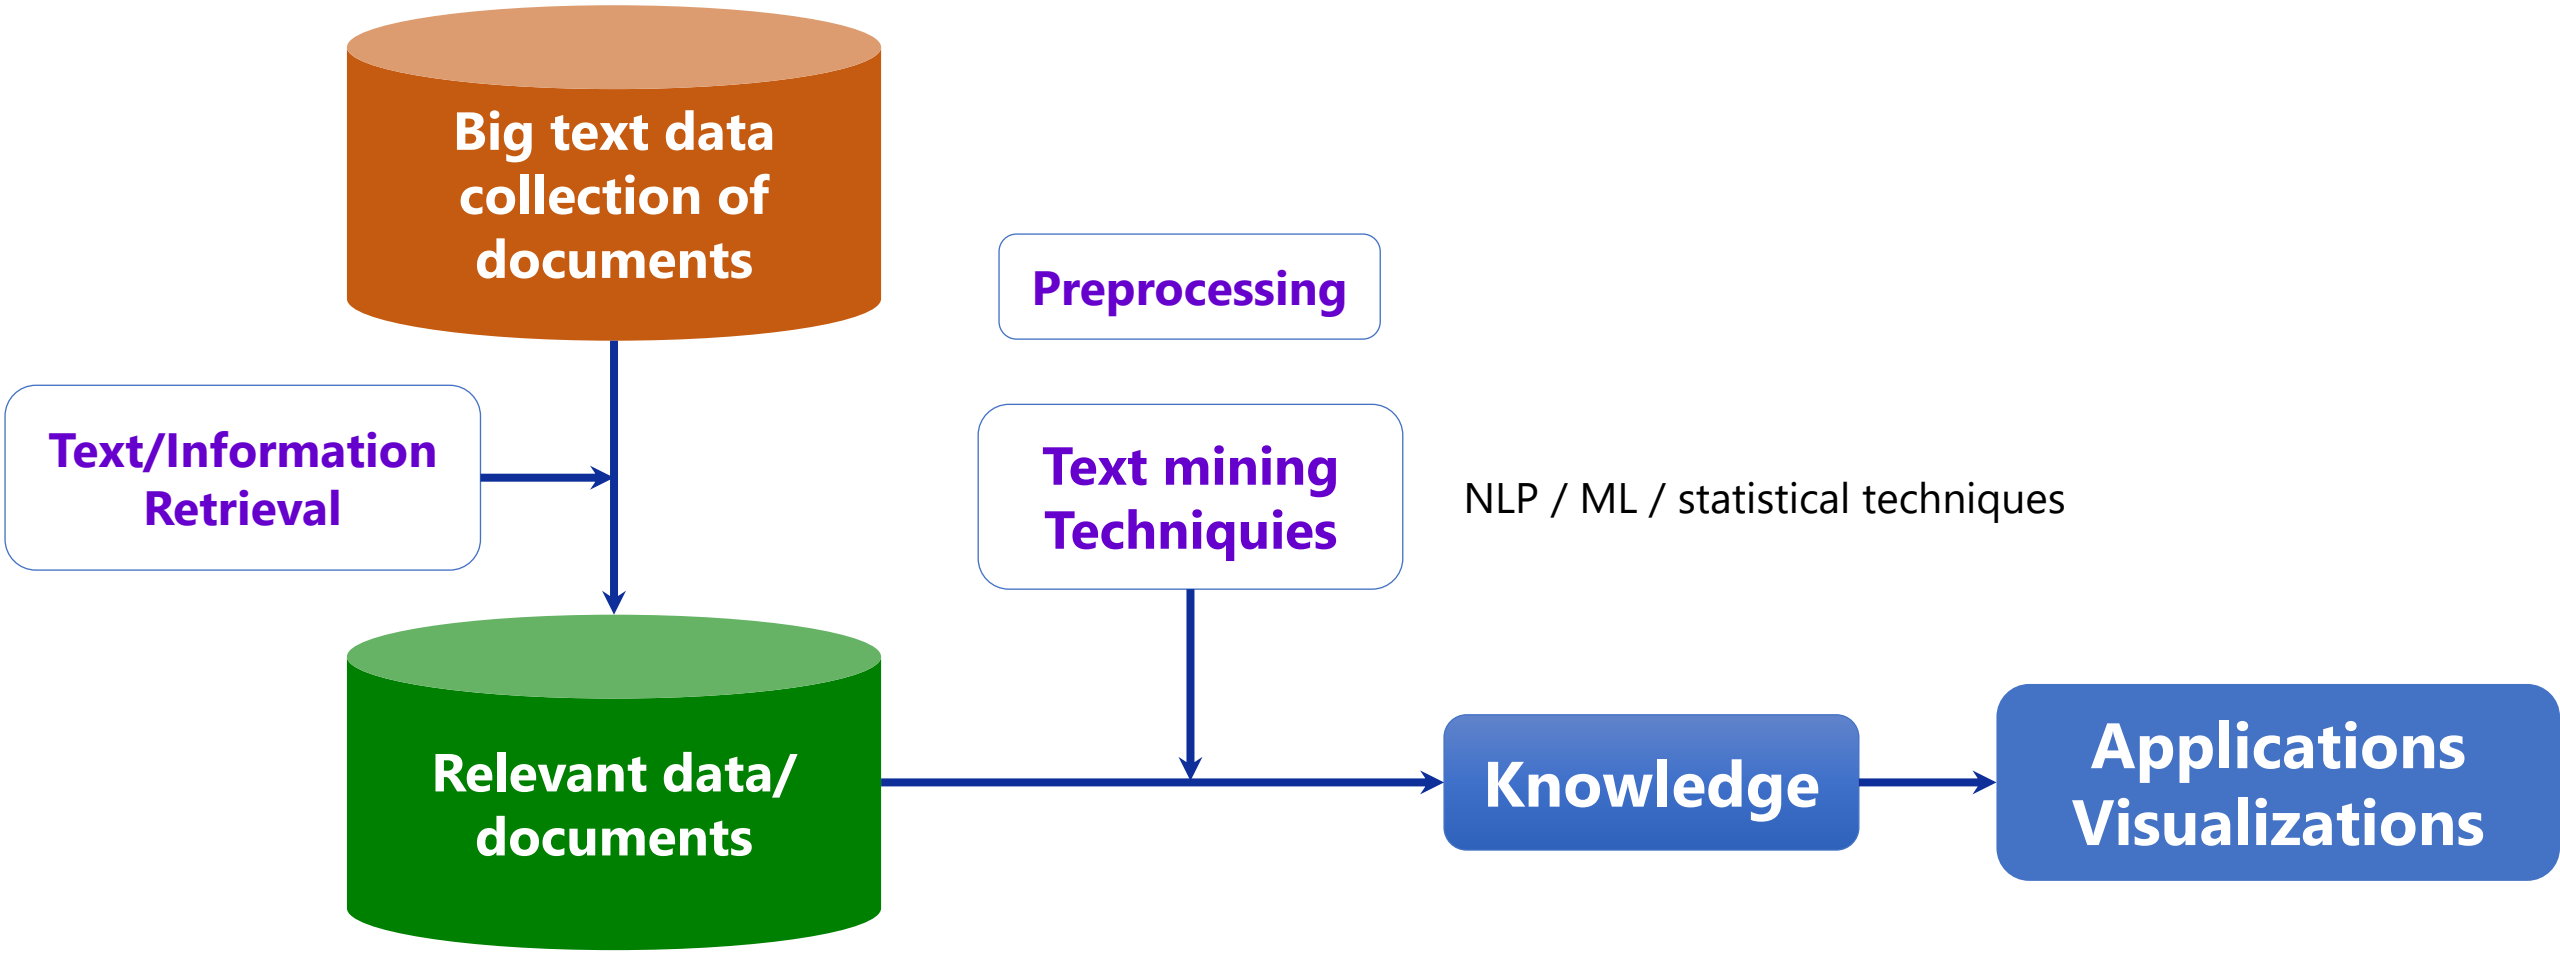

General Architecture

MENDELEY

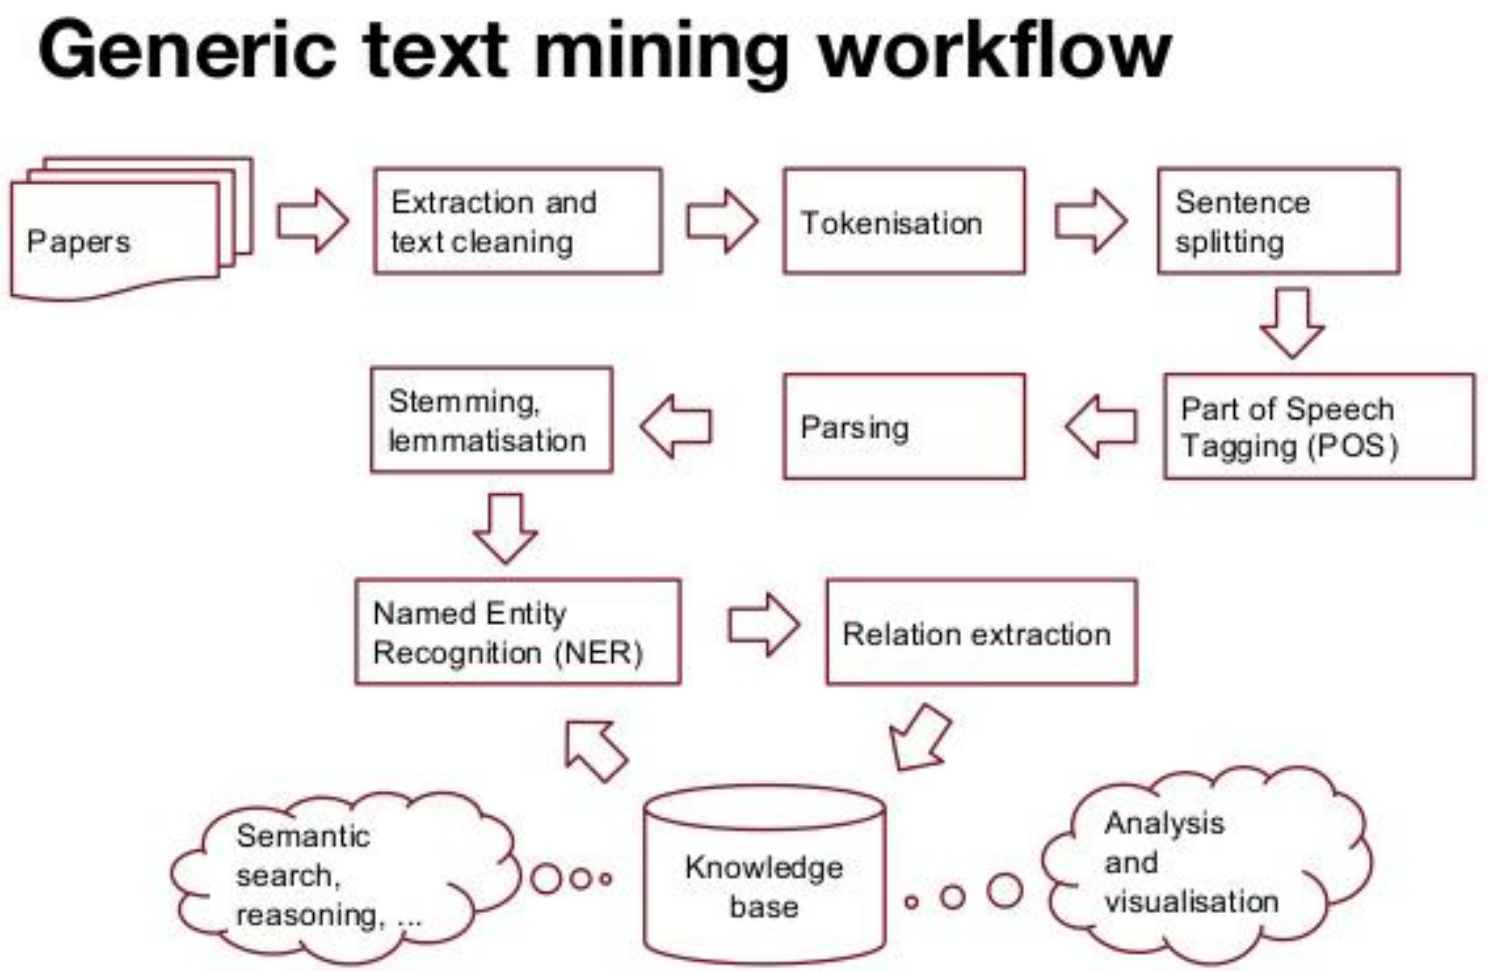

Introduction

- Machine learning tasks
  - Build a **model** from some **data**
    - choose how to map raw data to feature vectors
    - choose a model form
    - choose parameter values in the model
  - **Test** or **validate** the model
    - Evaluate the model on unseen data to assess its performance

Learning Algorithms

- The data is formed by a specific distribution.
- The goal of learning is to find the data distribution which is generalized in this kind of data

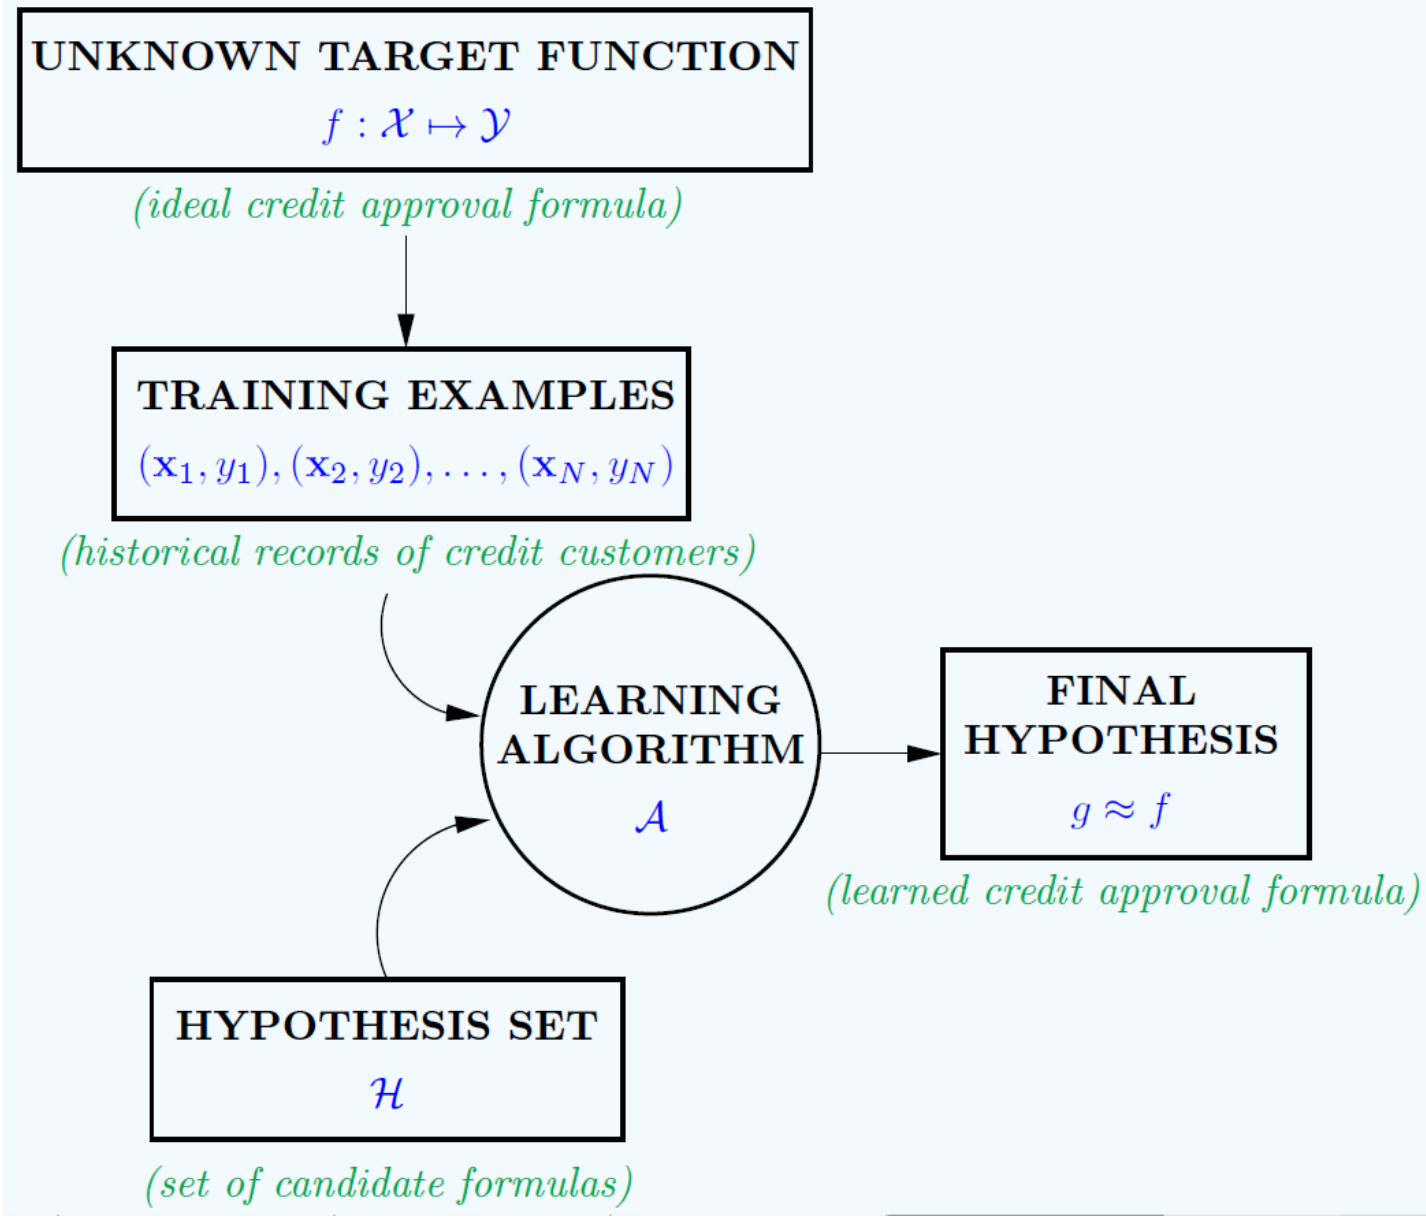

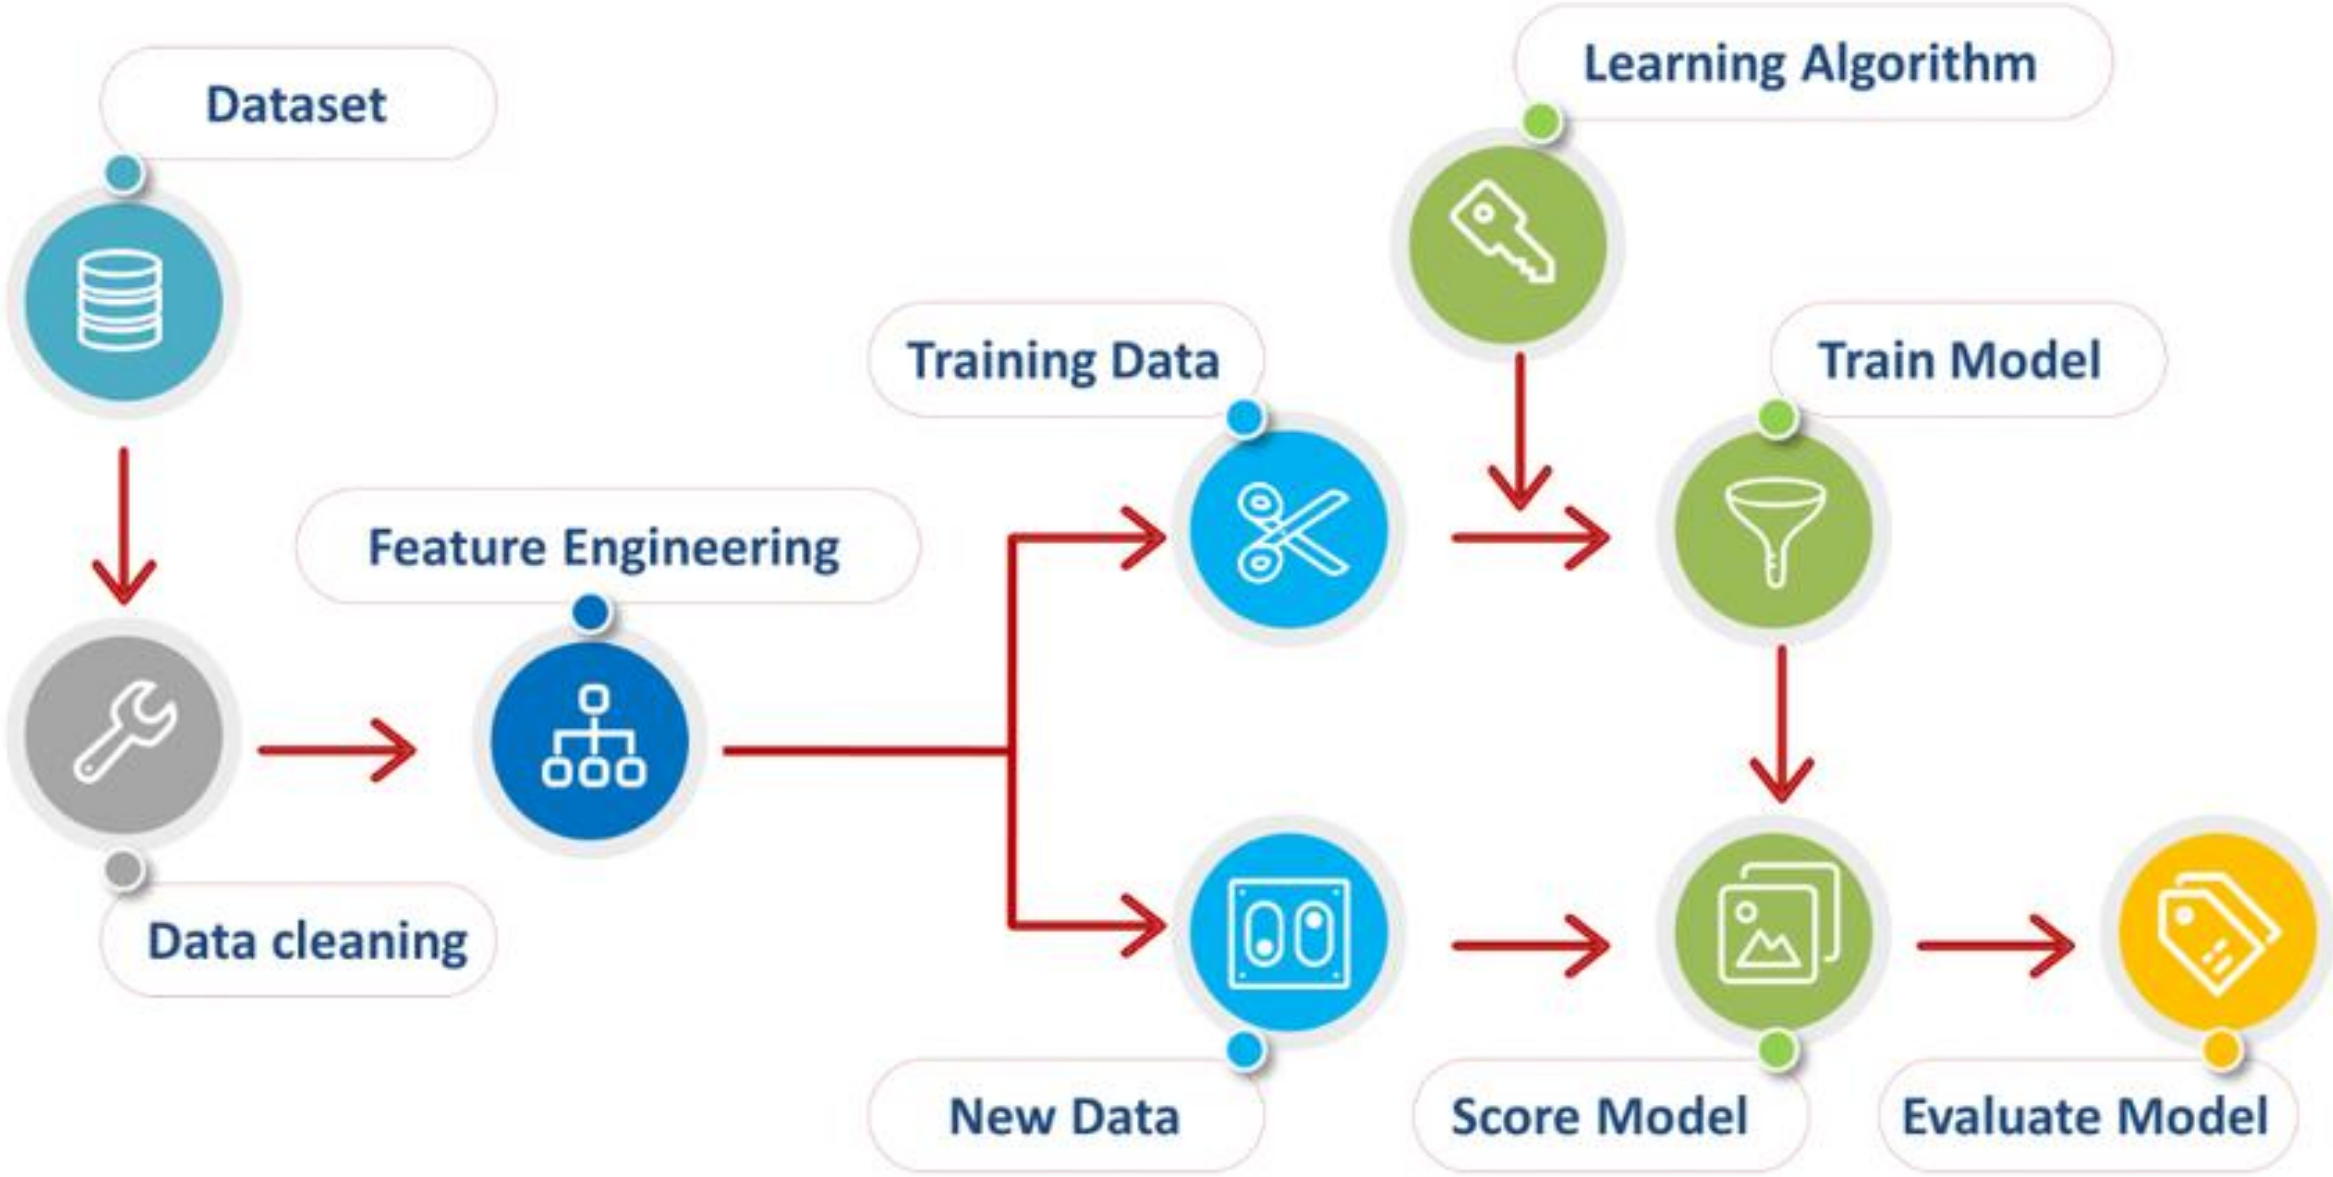
Dataset
Learning Algorithm
Training Data
Train Model
Feature Engineering
Data cleaning
New Data
Score Model
Evaluate Model

- Feature
  - A measurable property or characteristic of phenomenon being observed.

Human Learning:

We learn through

Examples

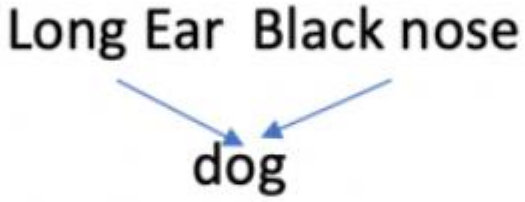

Diagrams

Comparisons

- Feature
  - A measurable property or characteristic of phenomenon being observed.

Machine Learning:

Sample

Label

dog

cat

horse

- Feature
  - A measurable property or characteristic of phenomenon being observed.

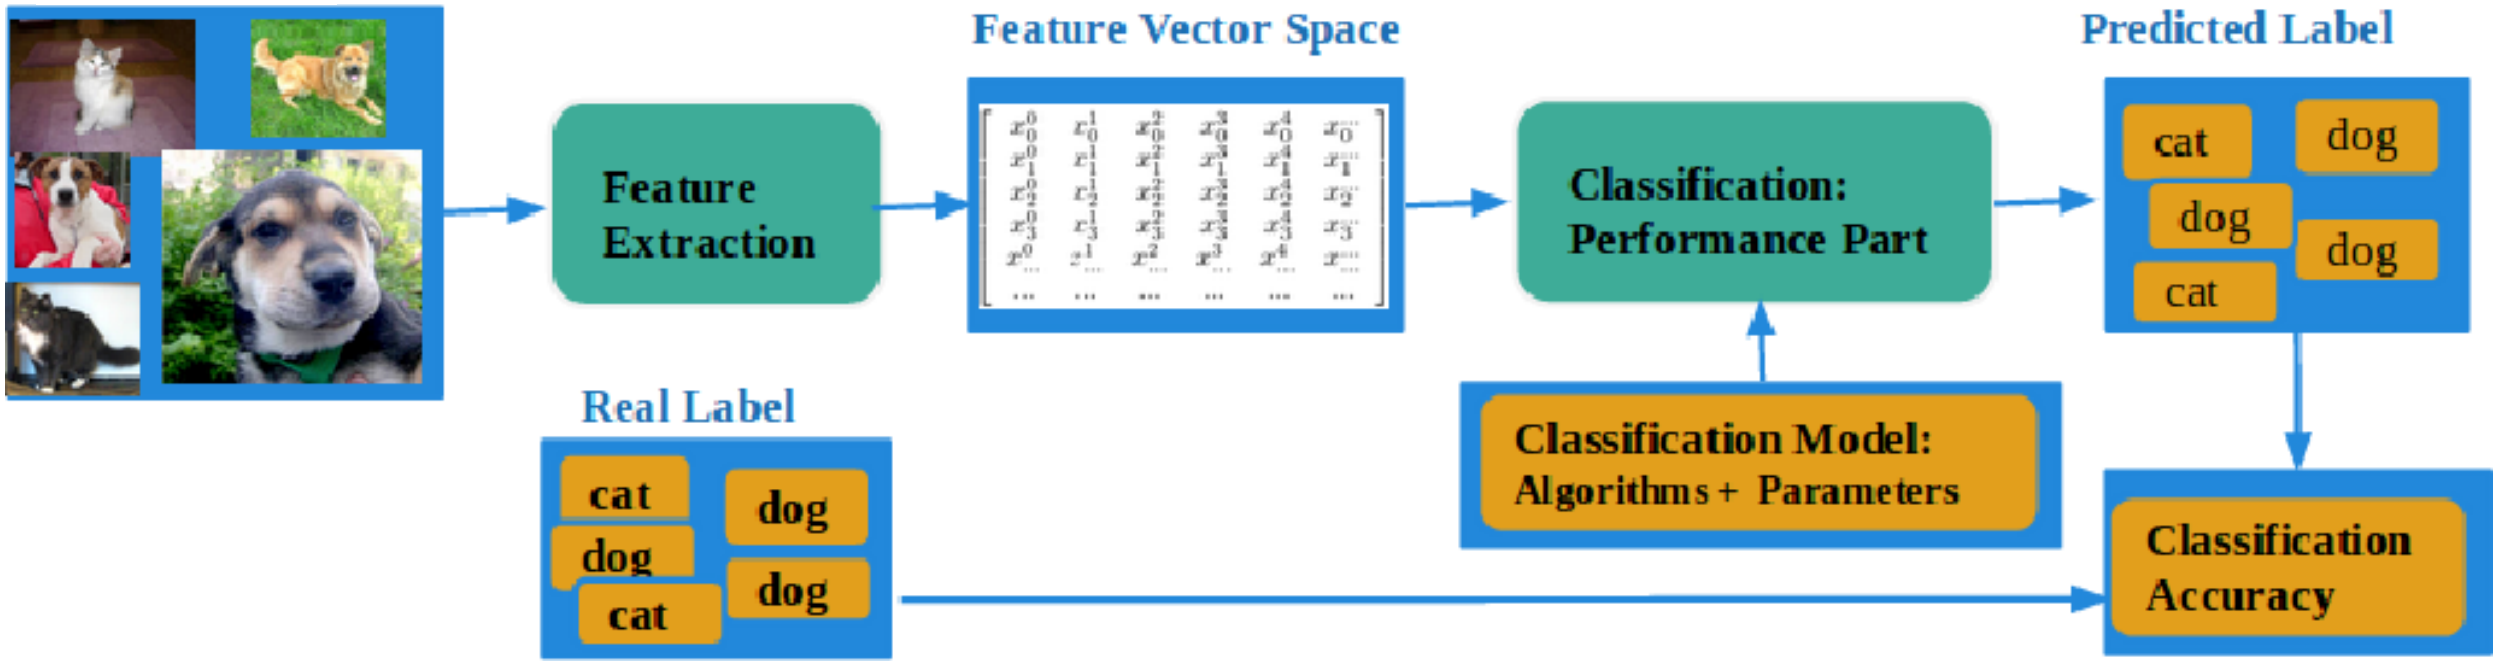

- Feature
  - A measurable property or characteristic of phenomenon being observed.

	Color	Weight	my Rank	...
	Red	200g	1st	
	Yellow	300g	3rd	

Predictive Tasks

- Build a Machine Learning model:
  - Training
  - Evaluating
  - Testing

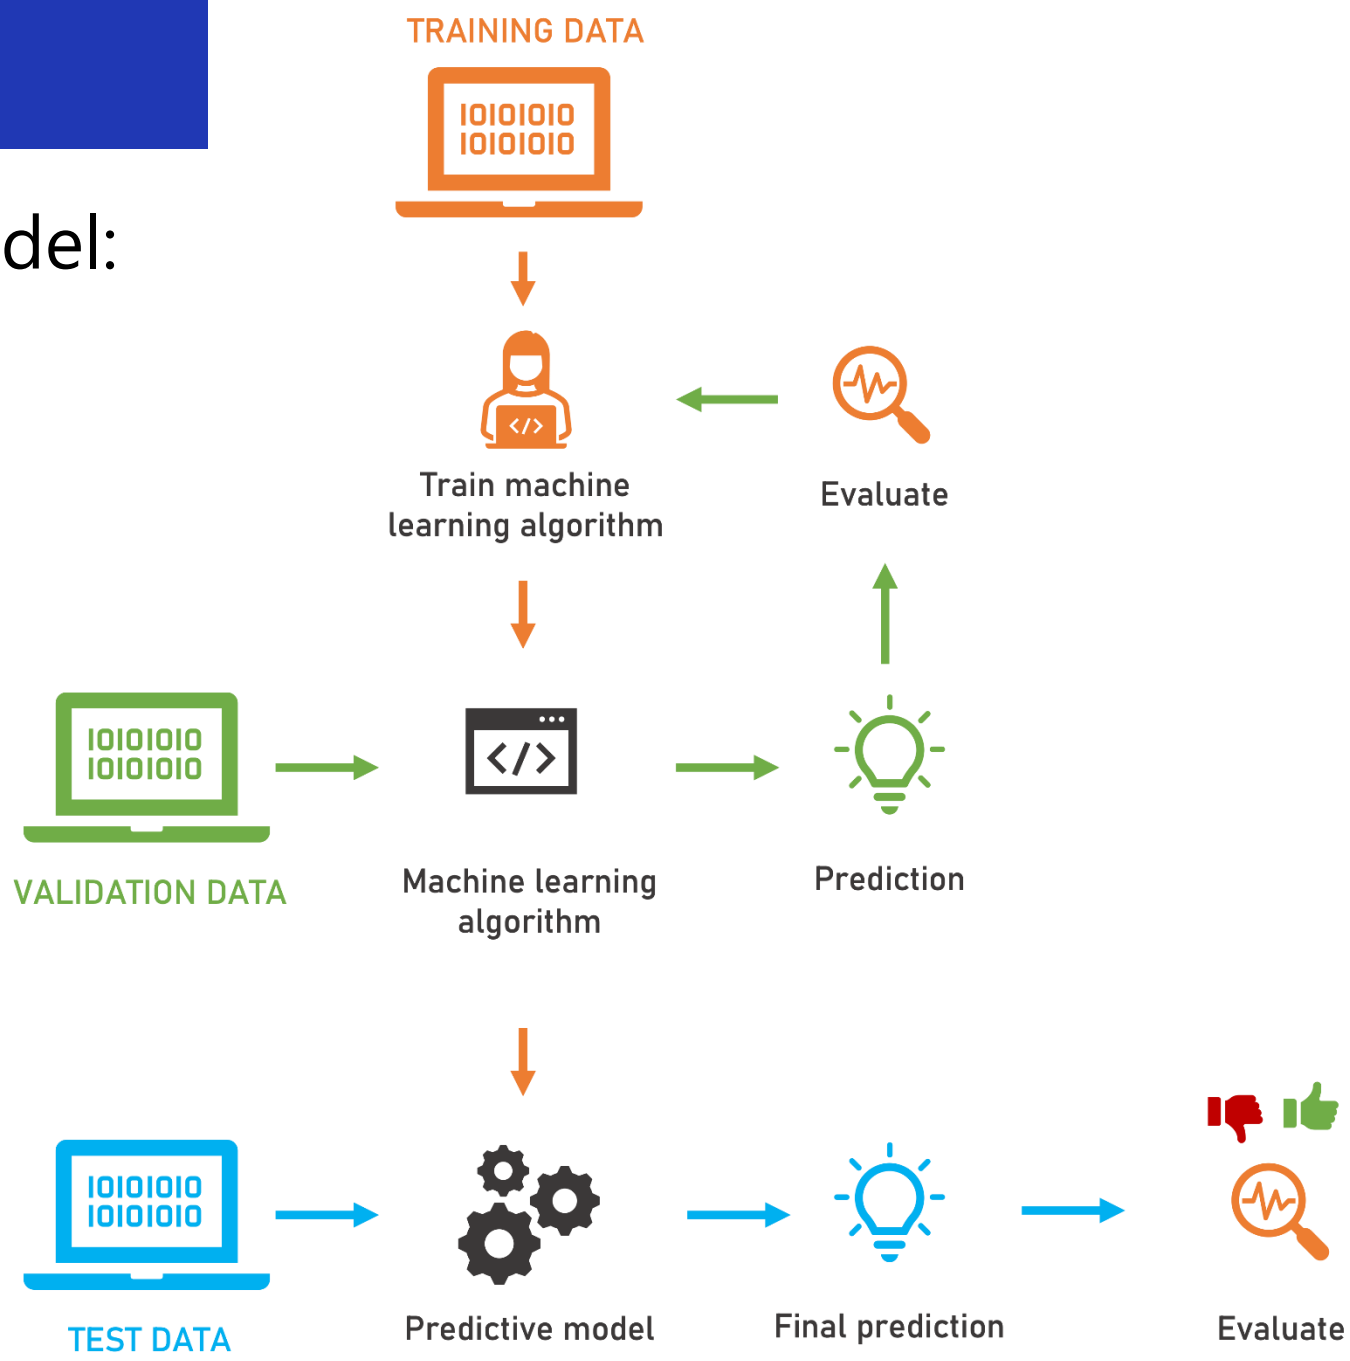

Learning Algorithms

- The goal of learning is to find the hidden pattern/distribution from historic data

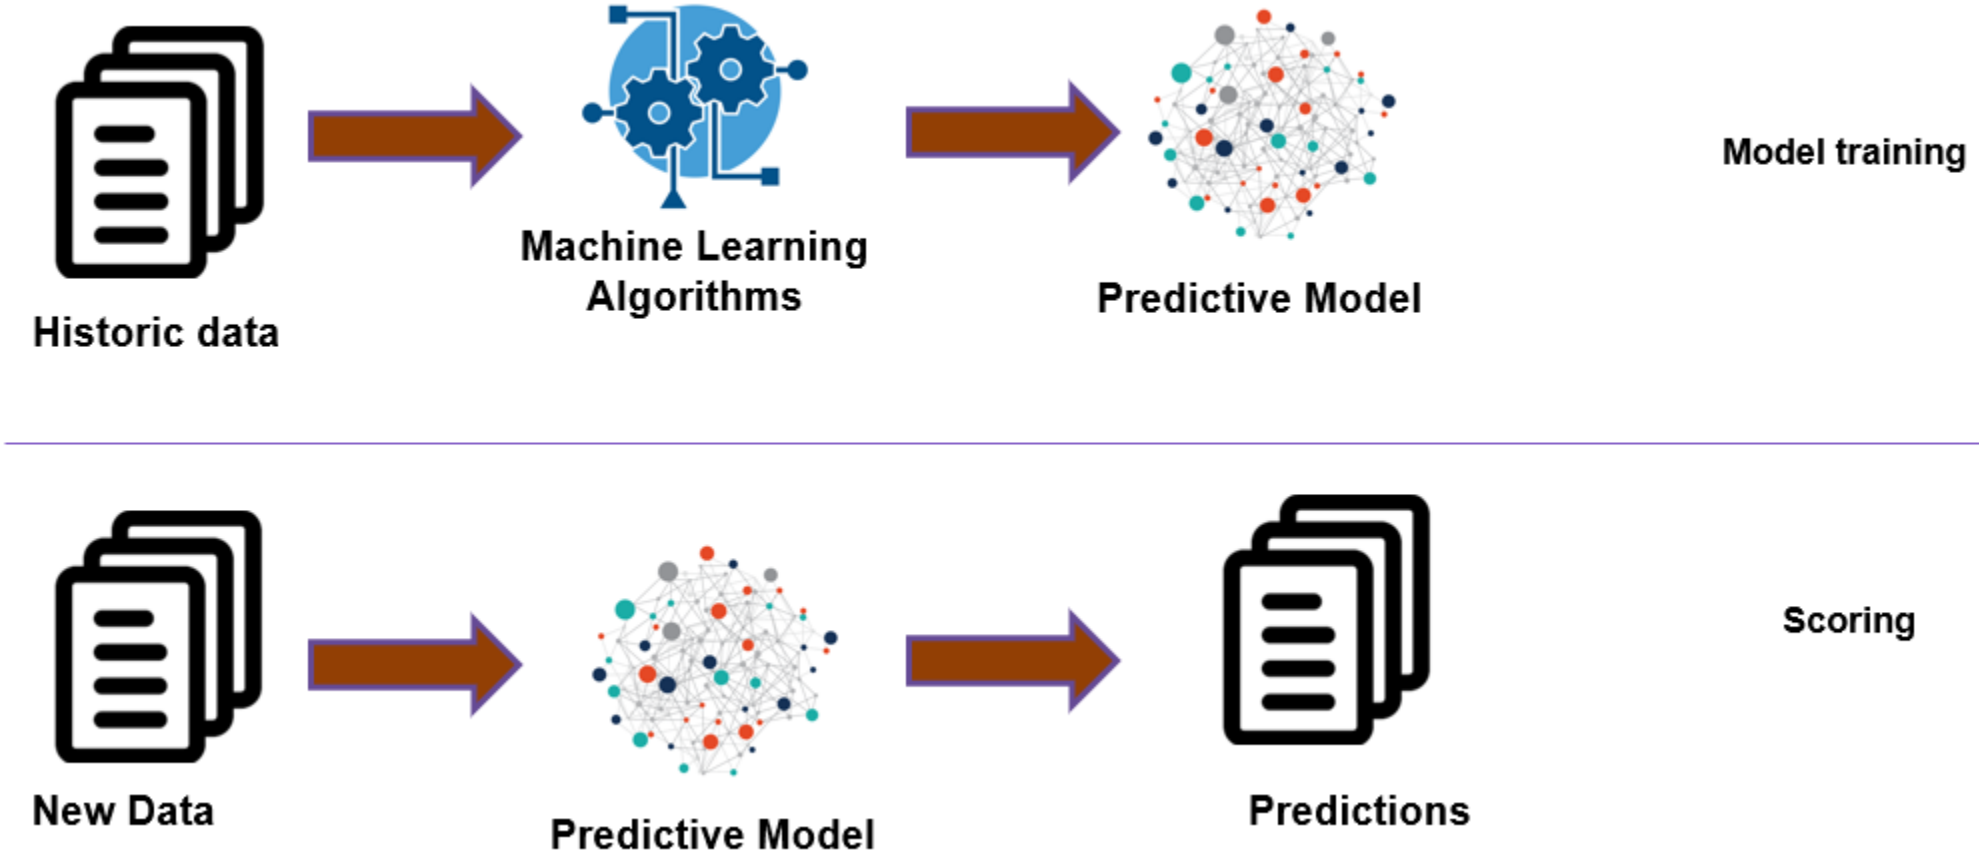

Type of Machine Learning

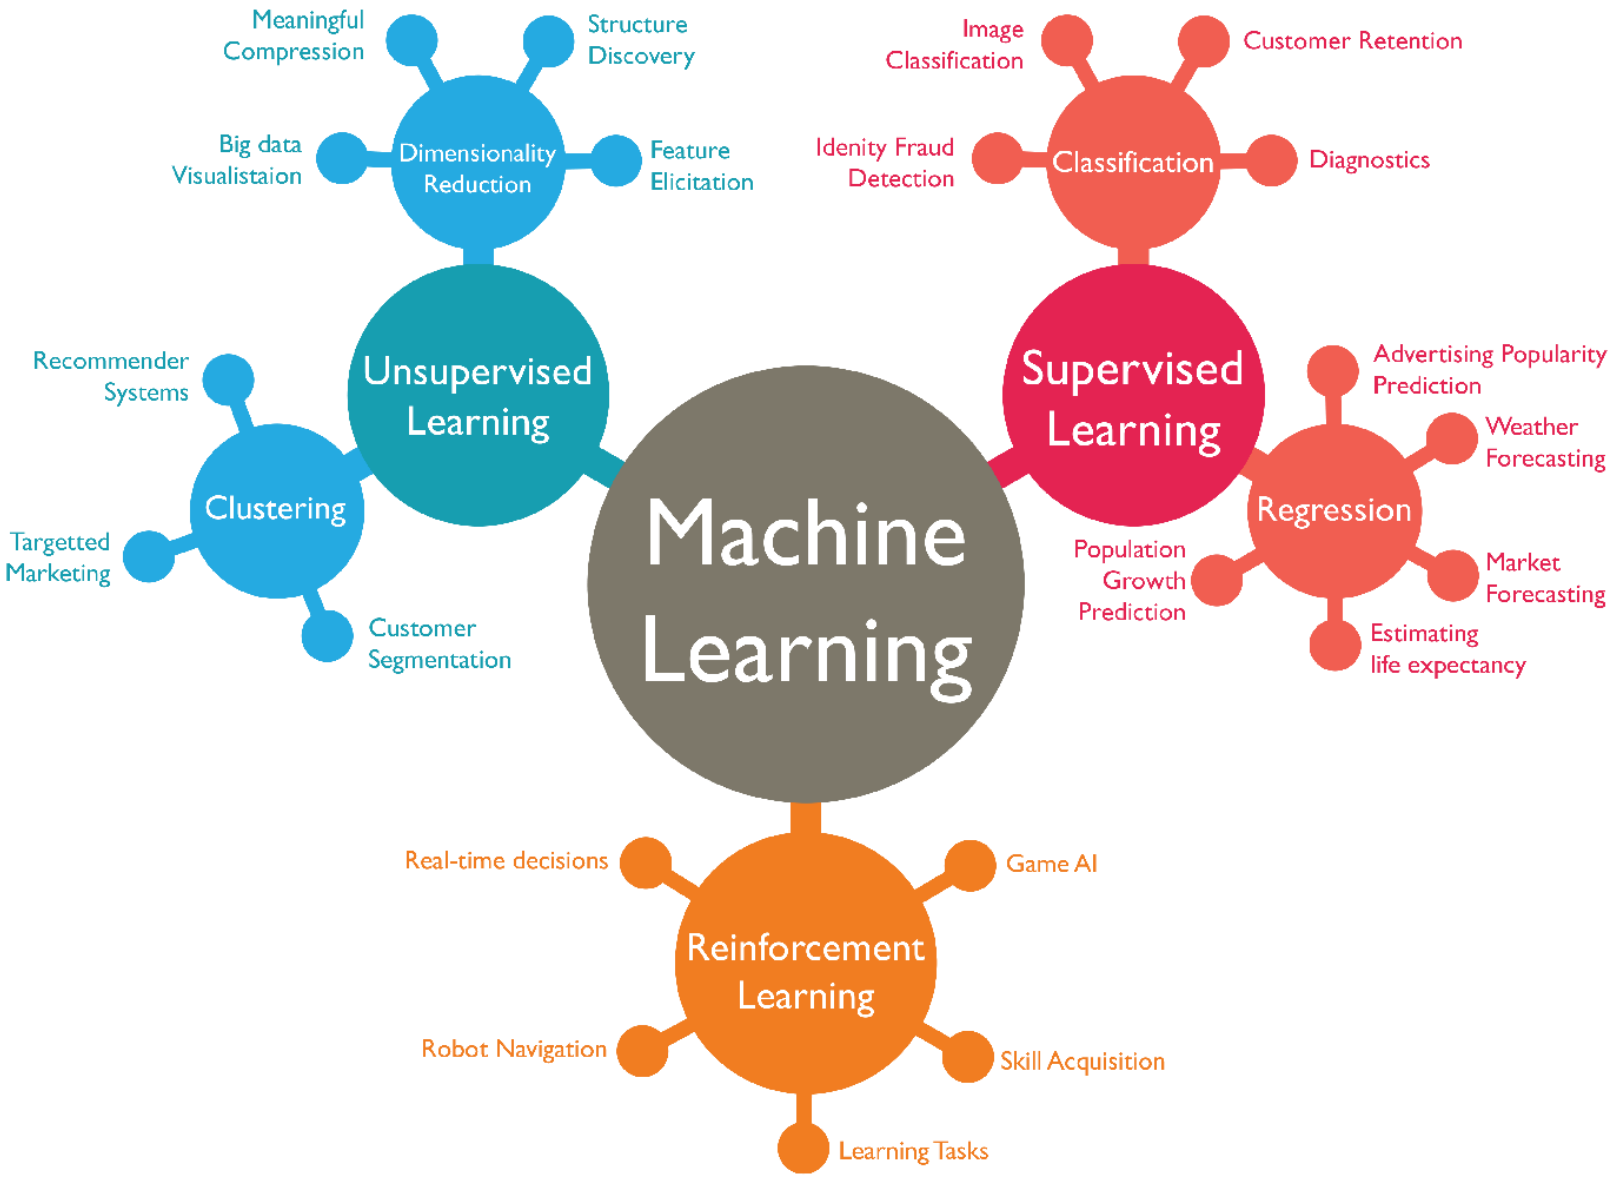

Type of Machine Learning

Supervised

Input:

- Features
- Gold labels

Tasks: task-oriented

- Classification
- Regression

Unsupervised

Input:

- Features

Tasks: data-oriented

- Clustering
- Dimension Reduction

Reinforcement

Input:

- Features
- Feed-back function

Tasks: environment

- Real-time Decision
- Learning tasks

Type of Machine Learning

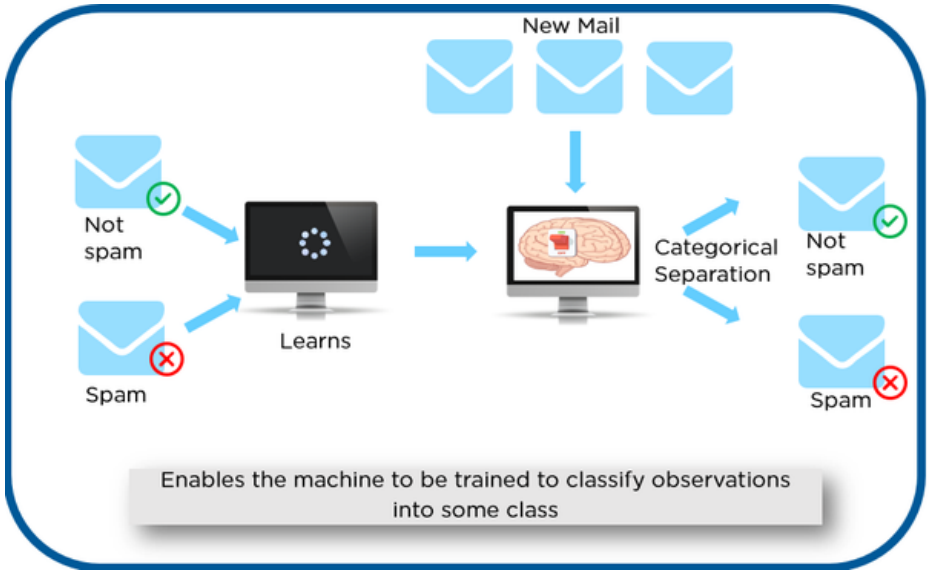

Supervised

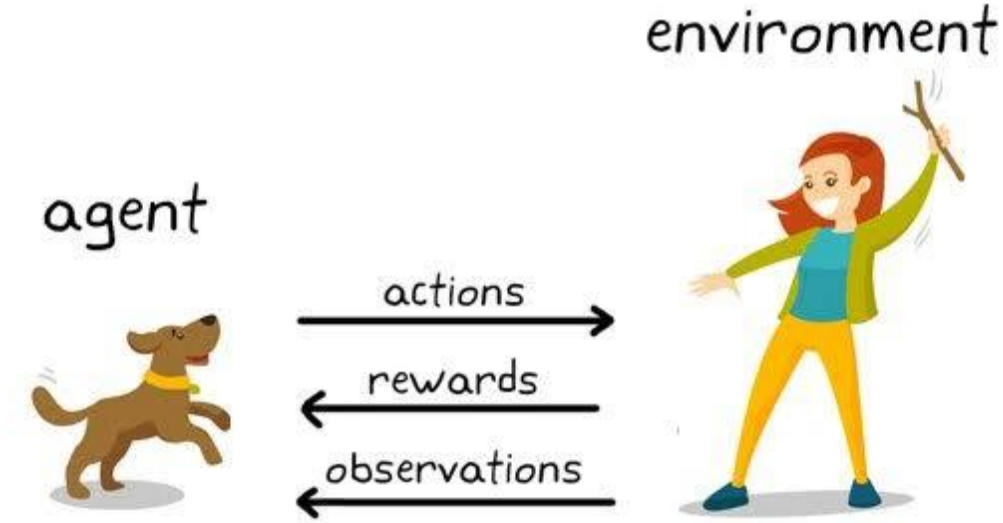

Reinforcement

Unsupervised

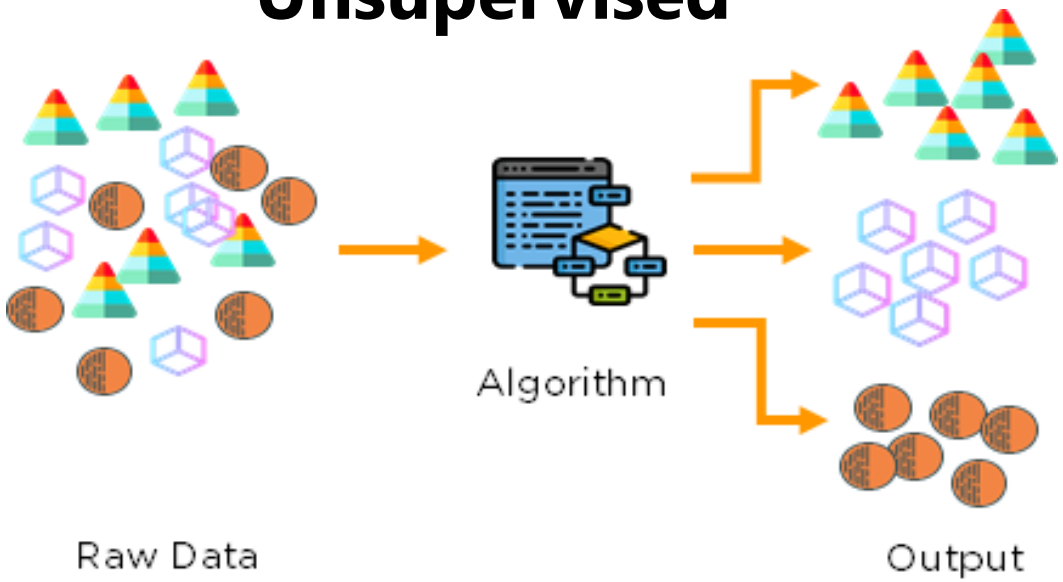

Supervised learning

- Learn from **labeled** data
  - Classification methods: predict a label
  - Regression methods: predict a quantity

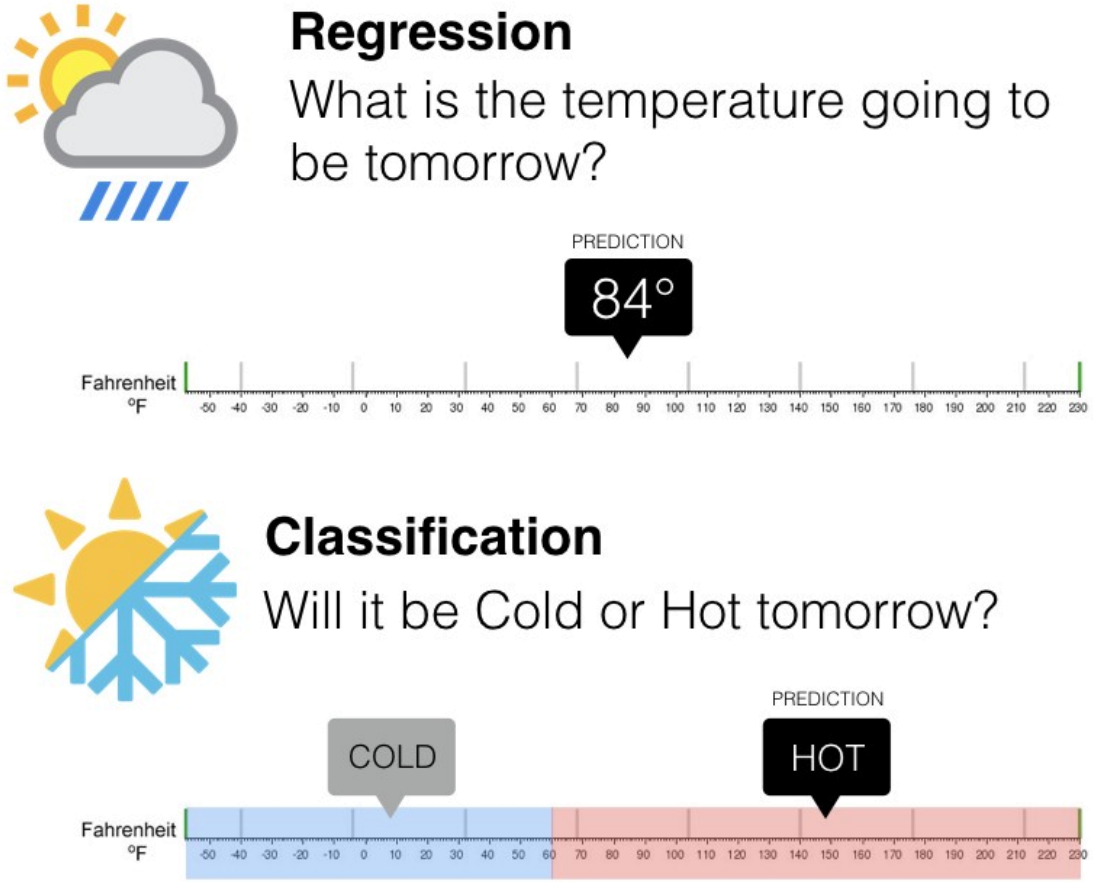

Supervised learning: Example

- Handwriting recognition:
  - Data: a set of handwriting images and their labels.
  - Goal: build a model to predict a digit for a given image.

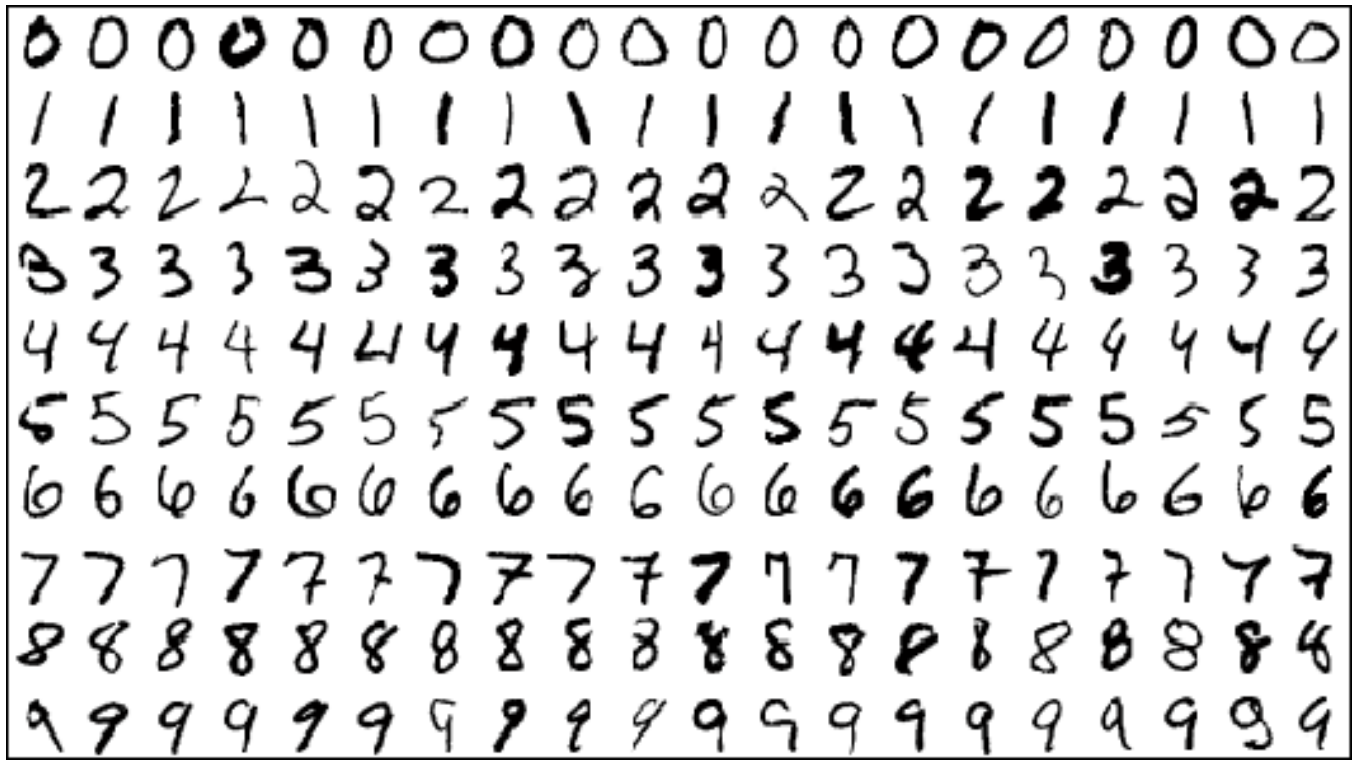

Supervised learning: Example

- Medical Costs Analysis:
  - Data: containing medical records
  - Goal: Build a model to predict medical costs for a given individual based on their demographic and health-related information

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

Unsupervised learning

- Learn from unlabeled data
  - Dimension reduction
  - Clustering
  - Association

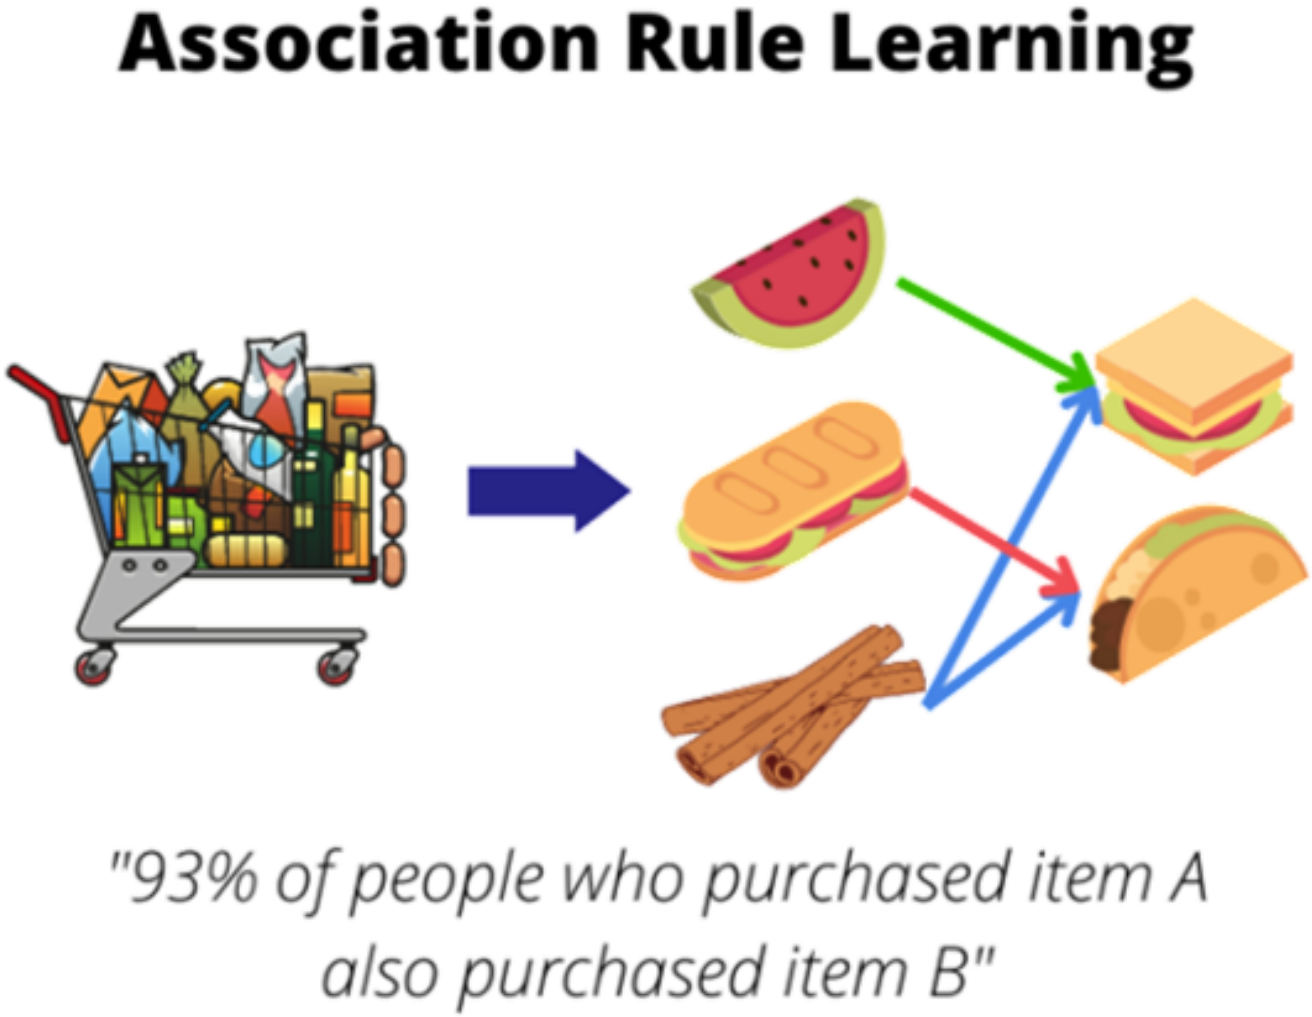

- Learn from mistakes

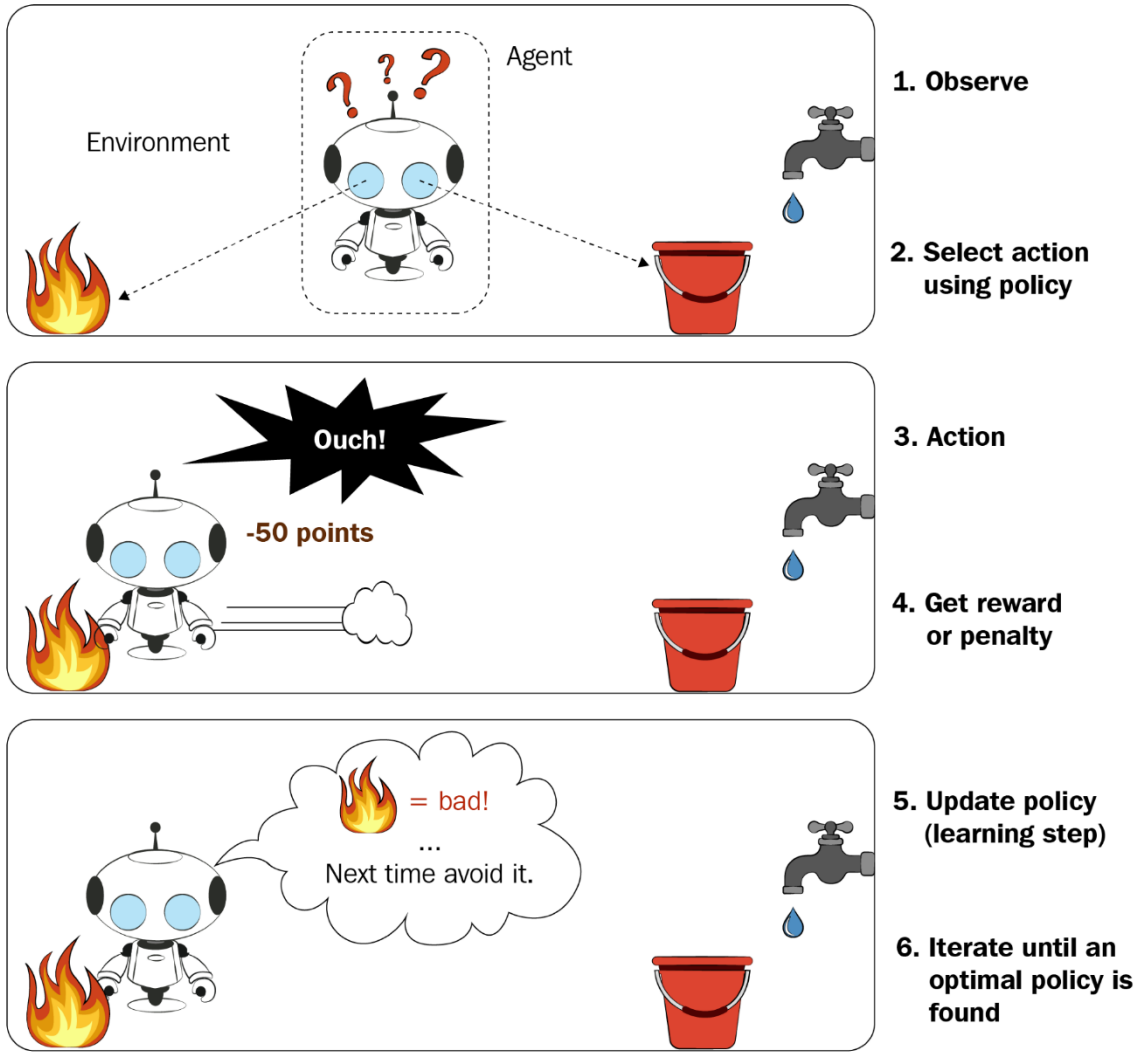

Feature Extraction

- Represent the data into the "learnable" object in computer

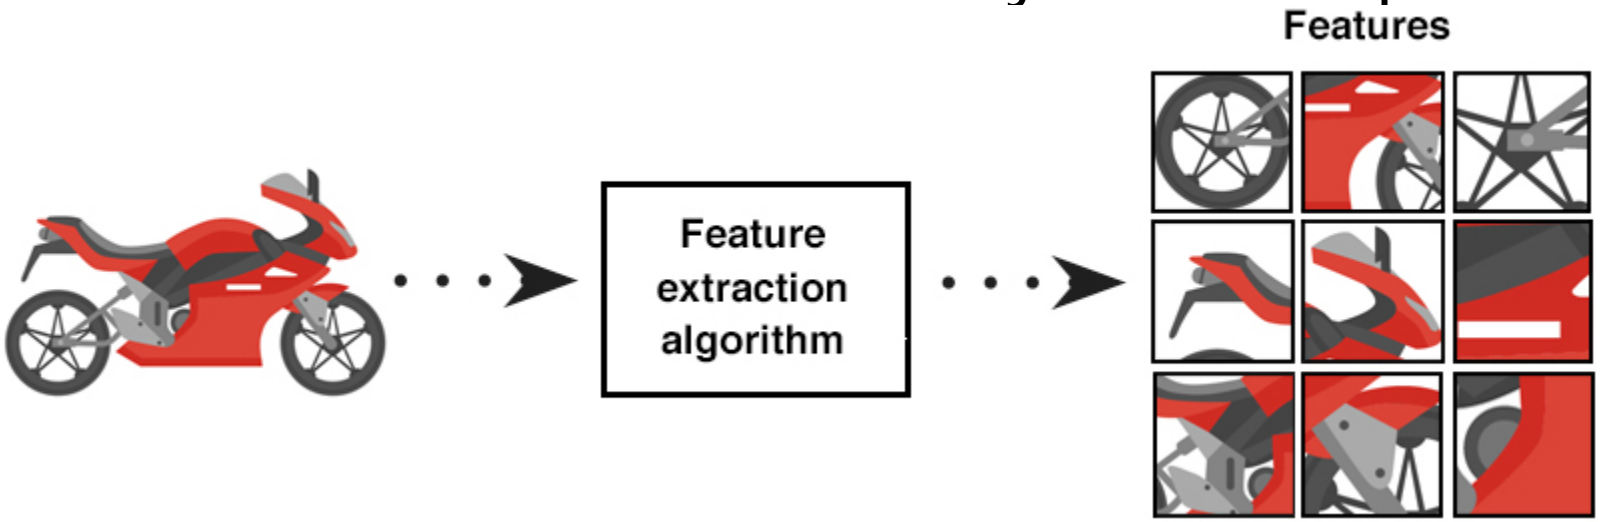

  - Image: pixels, rules, …
  - Audio: sampling, quantization
  - Text: ????

Feature Extraction

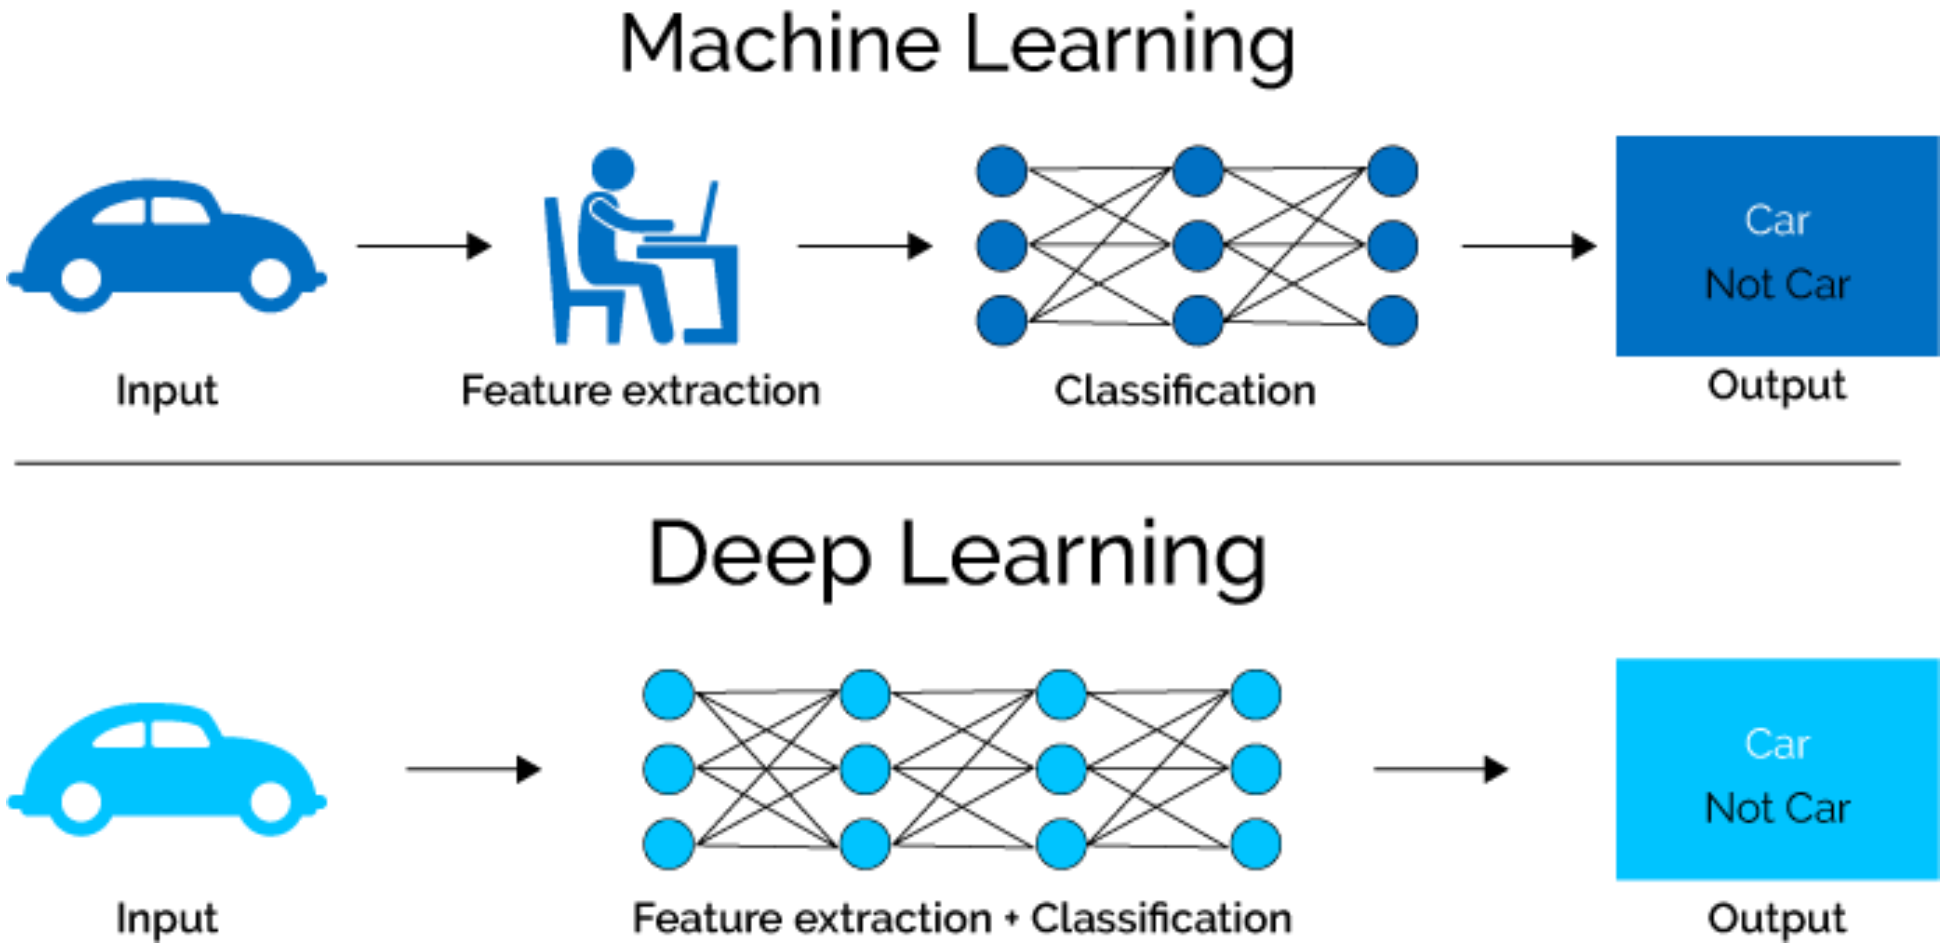

Evaluation

- How to choose the best "suitable" model

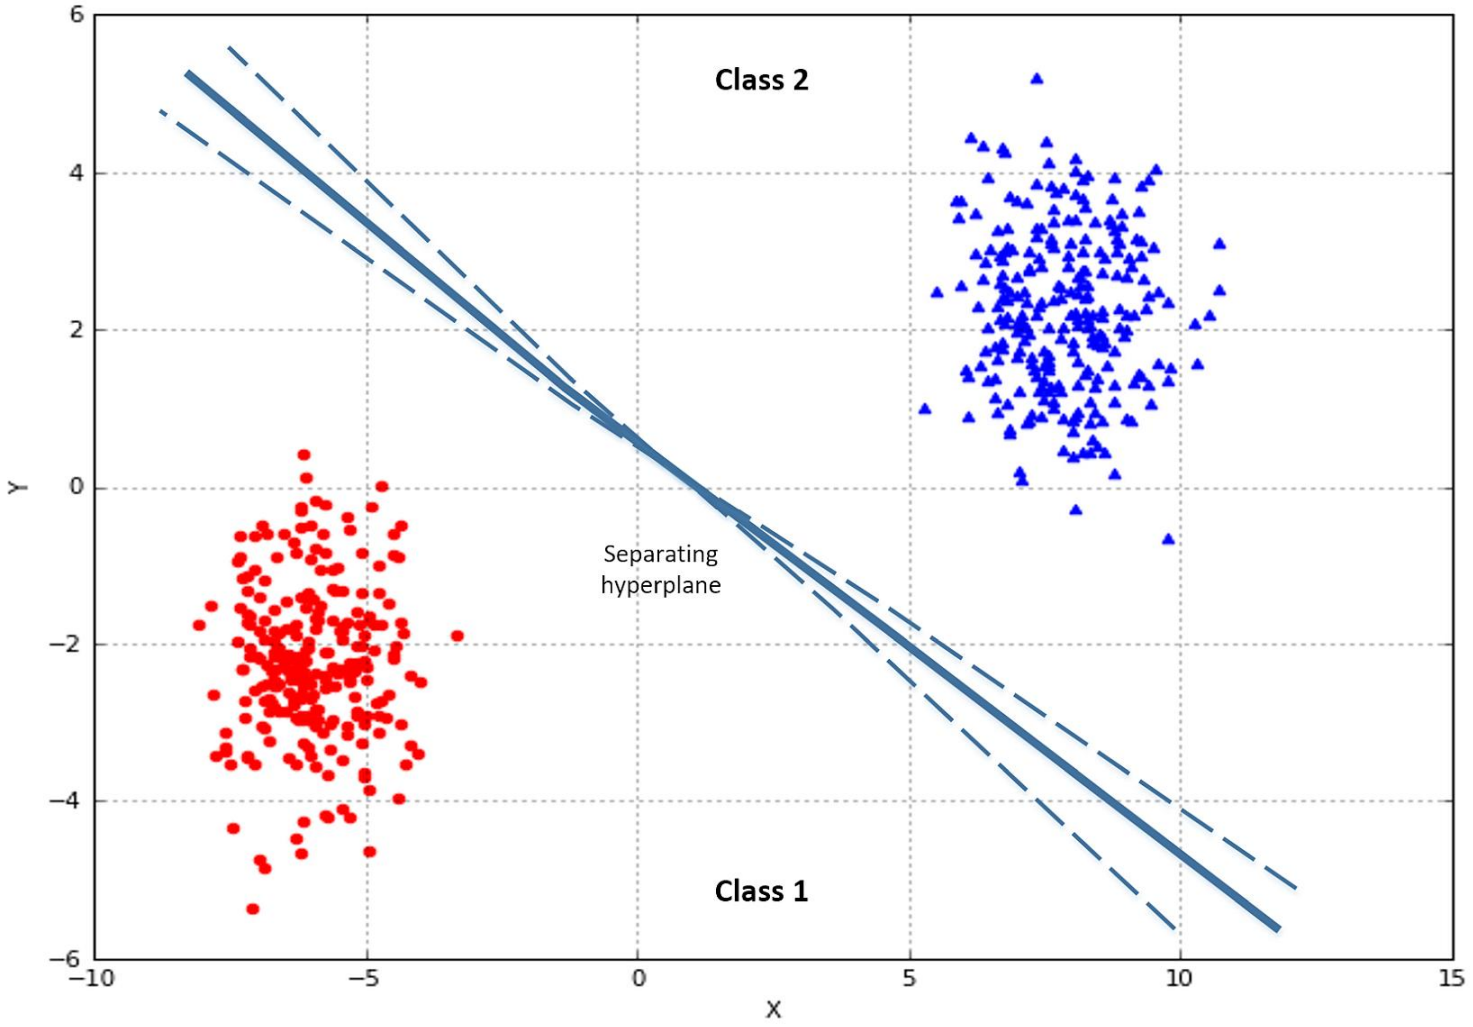

Evaluation

- Generalization refers to your model's ability to adapt properly to new, previously **unseen data**, drawn from the same distribution as the one used to create the mode

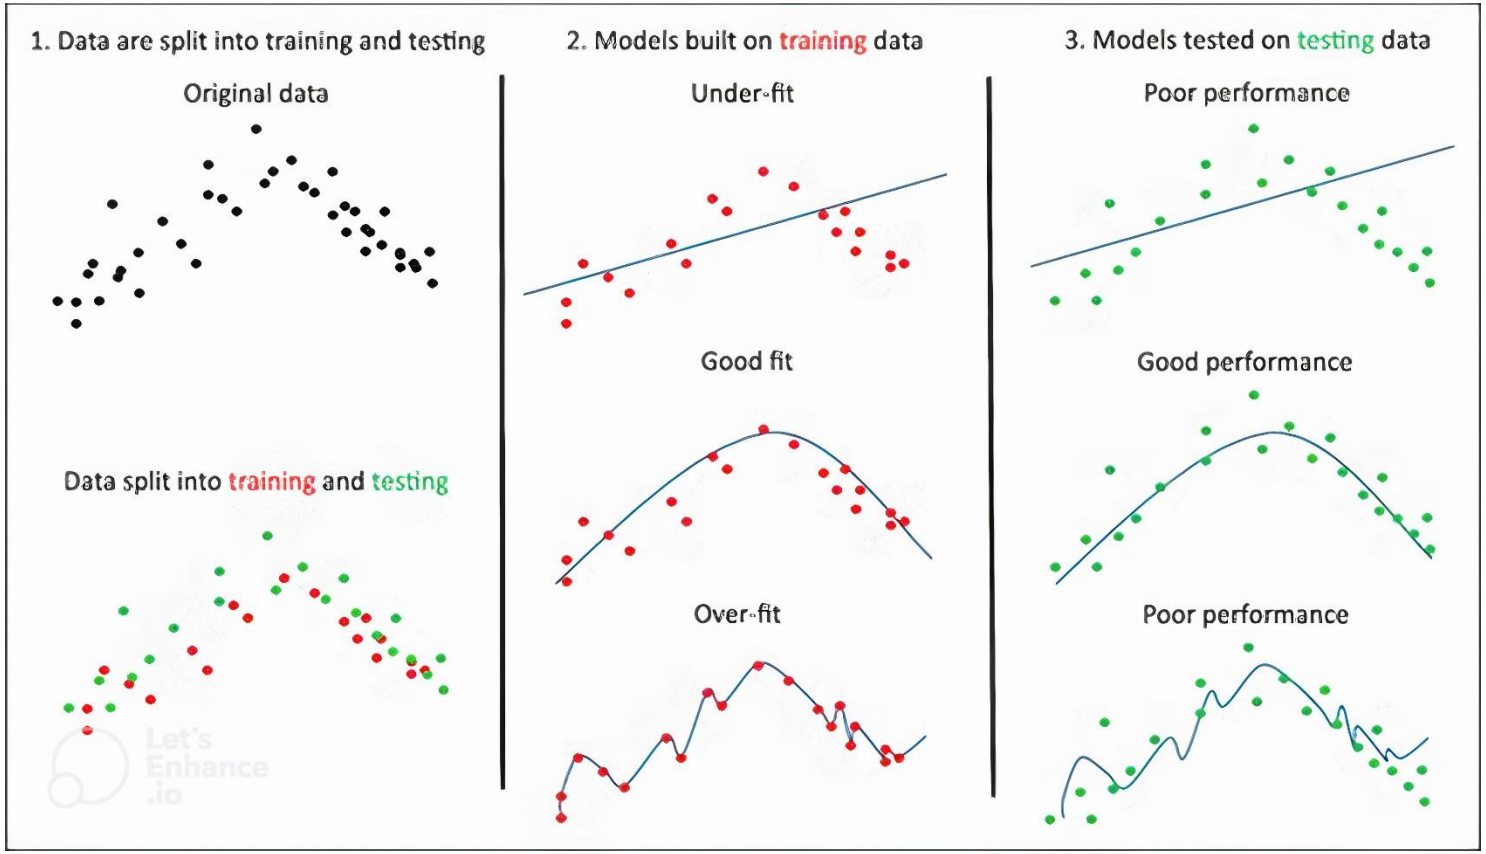

Evaluation

- Based on some metrics that compares the model's predicted labels with the known true labels

Regression	Classification	Recommender System
• Mean Absolute Error (MAE) • Root Mean Squared Error (RMSE) • R-Squared and Adjusted R-Squared	• Recall • Precision • F1-Score • Accuracy • Area Under the Curve (AUC)	• Mean Reciprocal Rank • Root Mean Squared Error (RMSE)

Evaluation: Classification

- A confusion matrix is a performance measurement technique
- Accuracy is a metric that generally describes how the model performs across all classes.

$$Accuracy = \frac{Correct\ prediction}{Total\ cases} * 100\%$$

$$Accuracy = \frac{(TP + TN)}{(TP + TN + FP + FN)} * 100\%$$

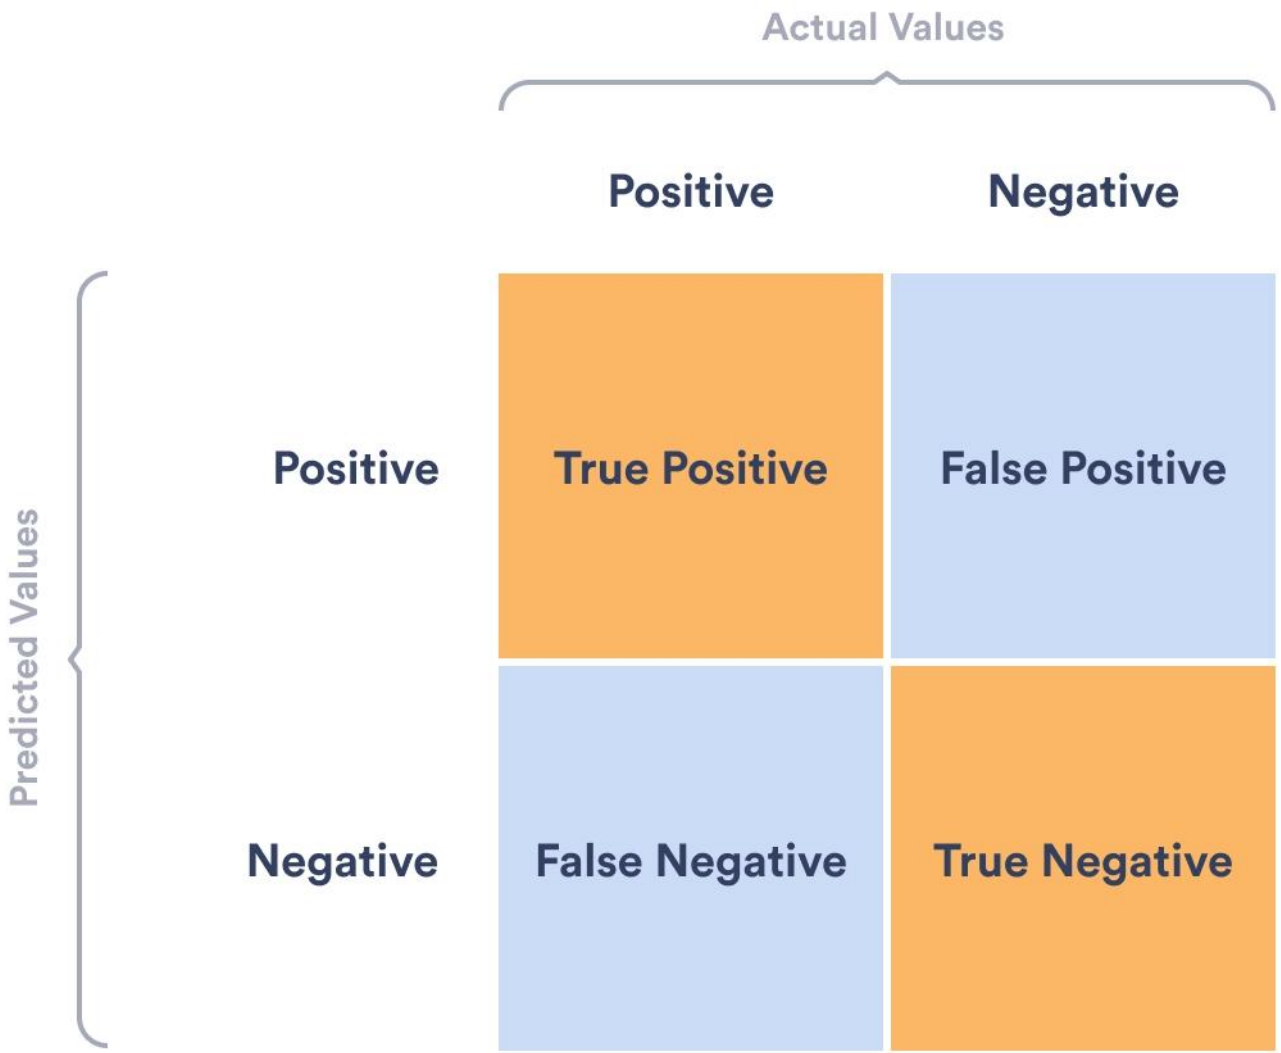

Evaluation: Classification

- Precision:

$$\text{Precision} = \frac{True\ Positive}{True\ Positive + False\ Positive}$$

$$= \frac{True\ Positive}{Total\ Predicted\ Positive}$$

- Recall:

$$\text{Recall} = \frac{True\ Positive}{True\ Positive + False\ Negative}$$

$$= \frac{True\ Positive}{Total\ Actual\ Positive}$$

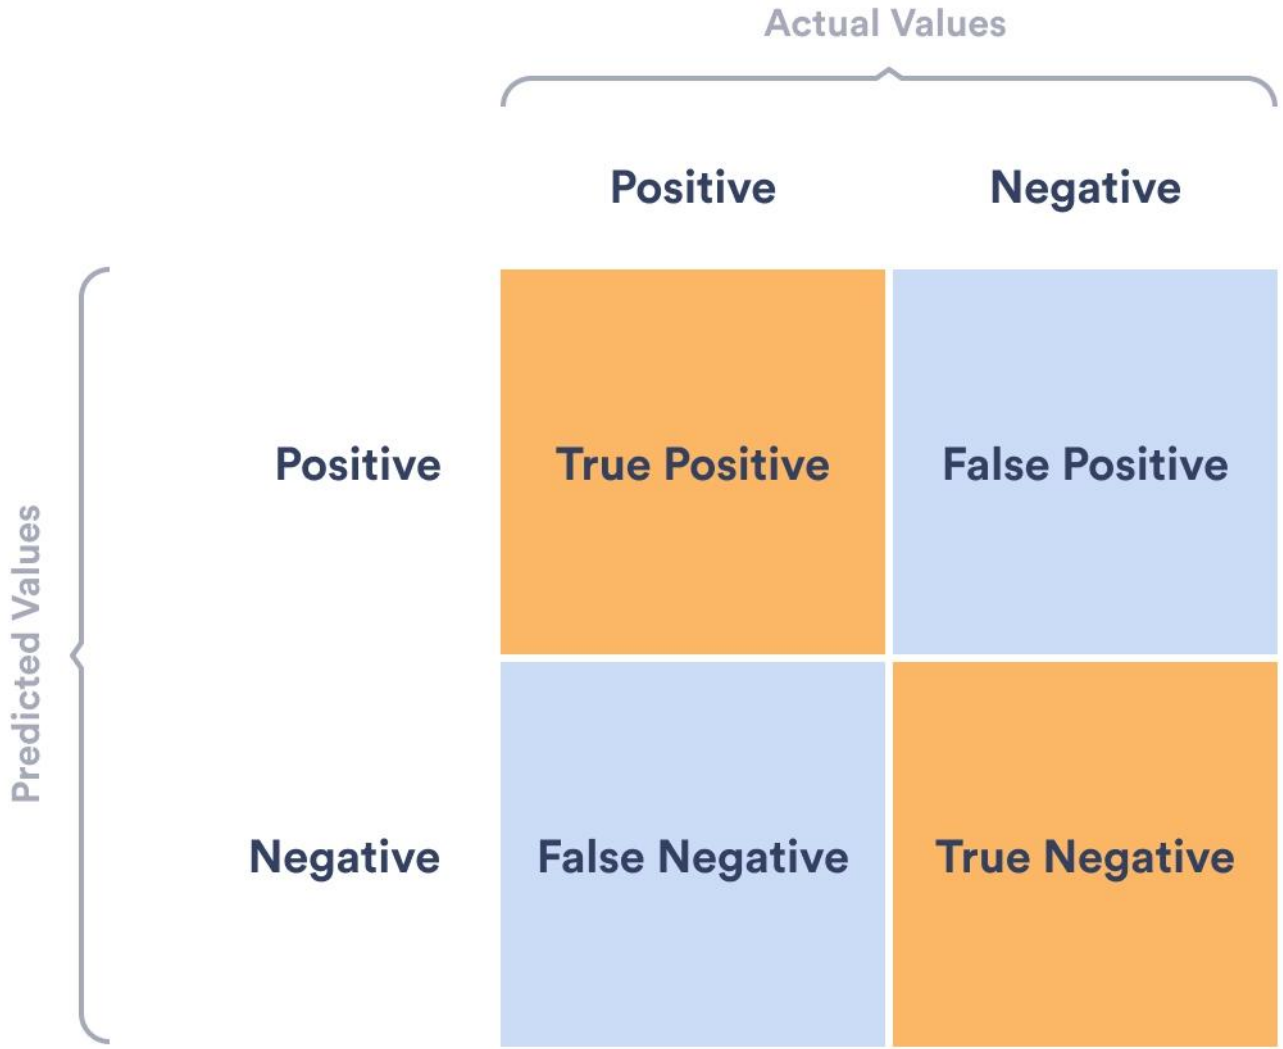

Data Splitting Techniques

- Kinds of data:
  - Train: learn from the historical patterns to reveal the hidden models
  - Validating: choose the best/good model, tune the model's hyperparameters and configurations
  - Testing: evaluate the practical performance

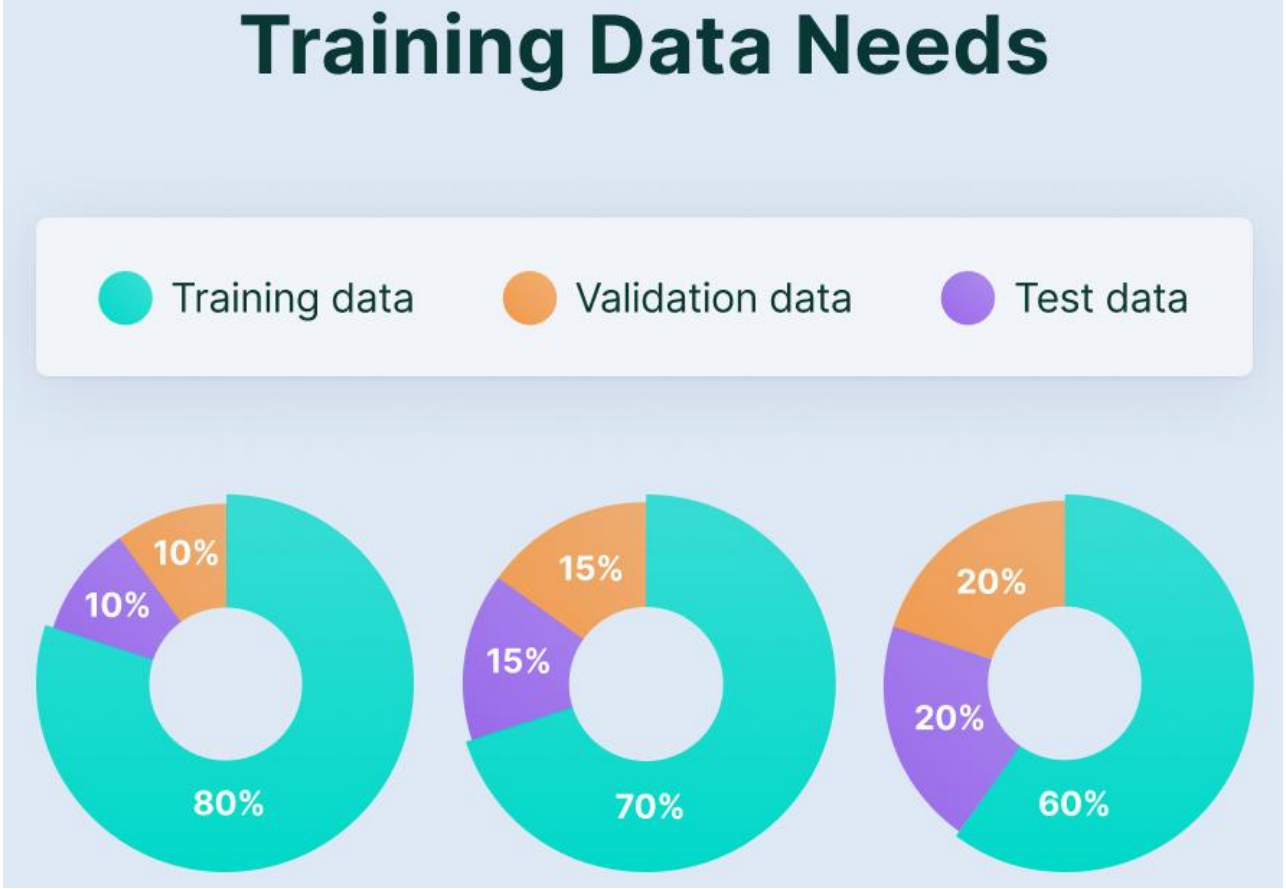

Data Splitting Techniques

- Random sampling
  - Simple
  - Stratified: random from each group (group division by label)
  - Cluster: group division by features
  - Multistage: divide by many kinds of random sampling

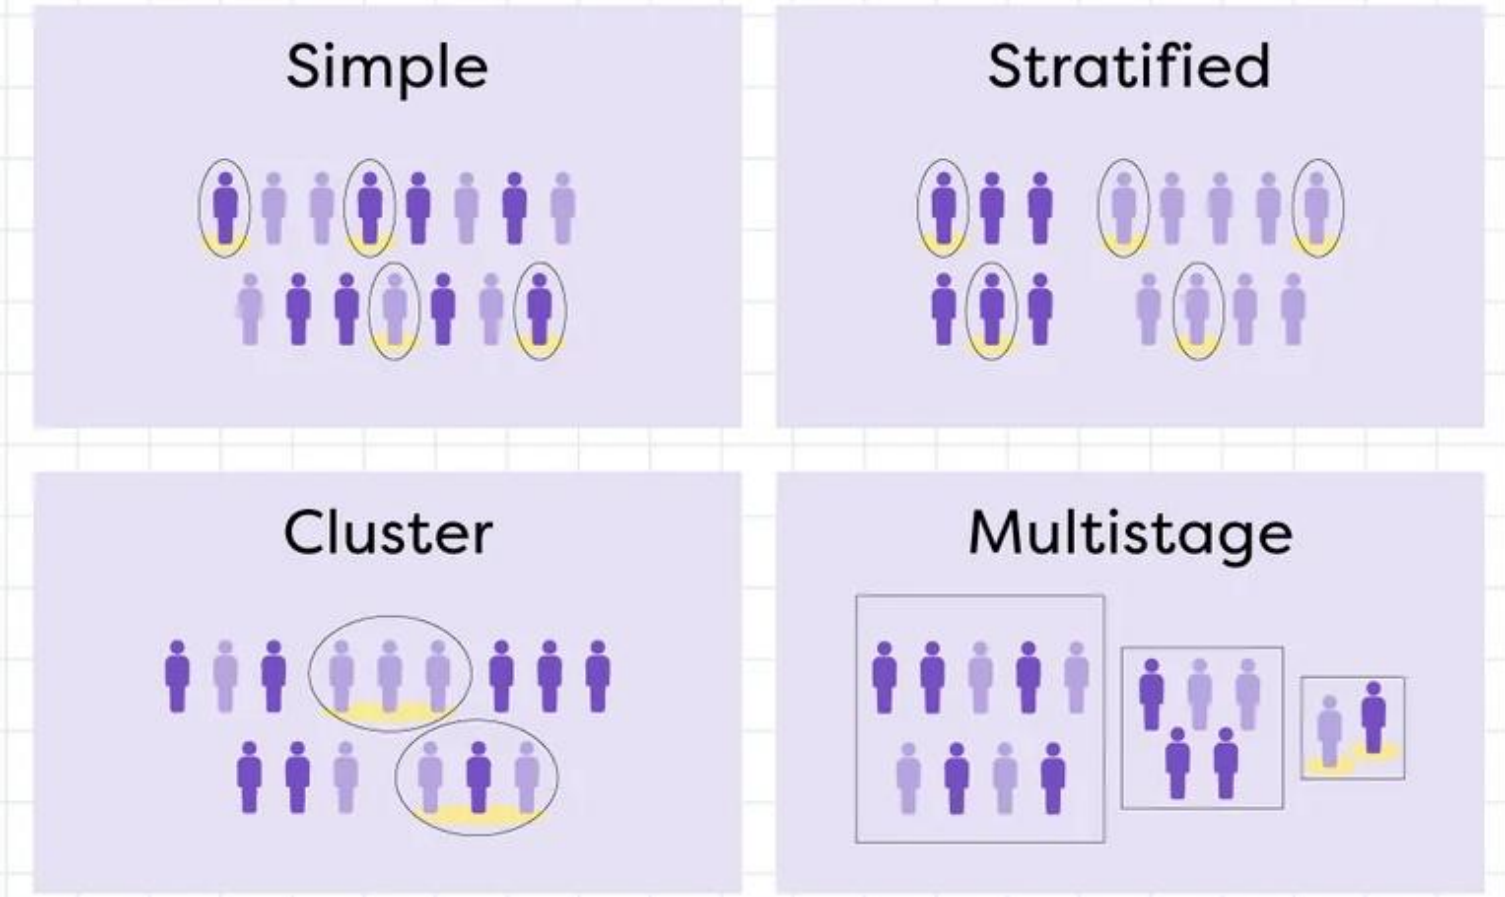

Data Splitting Techniques

- Cross-validation:
  - Cross over training data
  - Cross over all data
- This alleviates any bias occurring as selecting data in the training and validation sets.
- The goal of this technique is validating and testing, not training model.

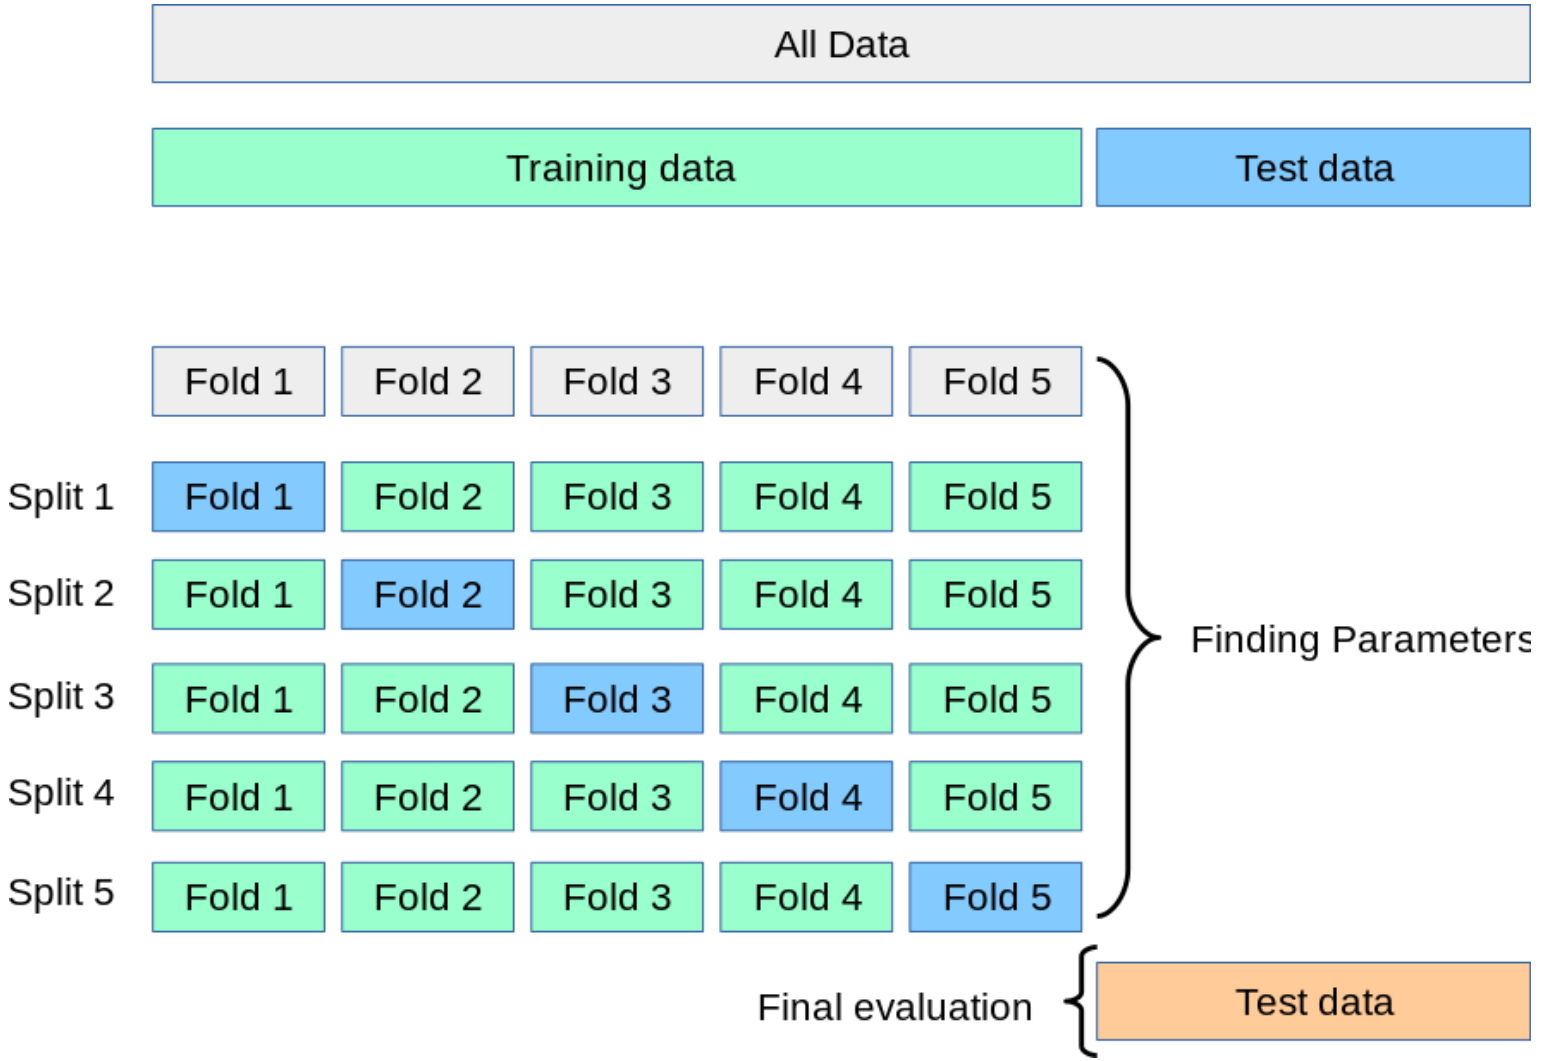

Generalization Error

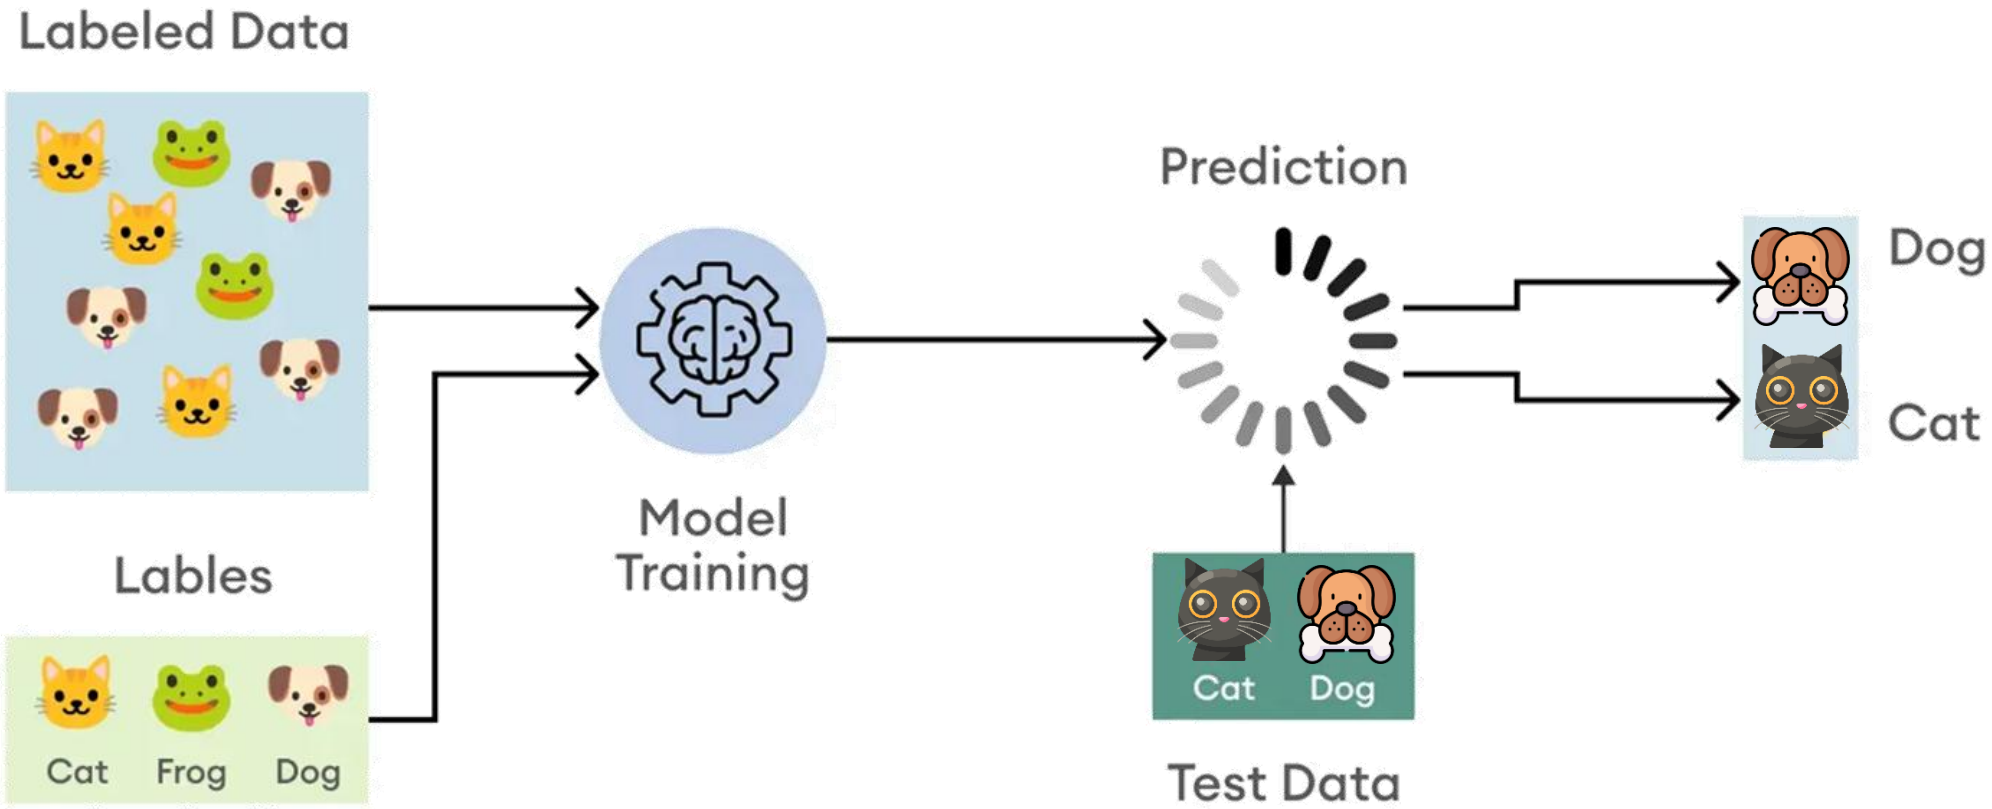

No free lunch theorem

- If an algorithm performs well on a certain class of problems, then it necessarily pays for that with degraded performance on the set of all remaining problems. (Wolpert and Macready)

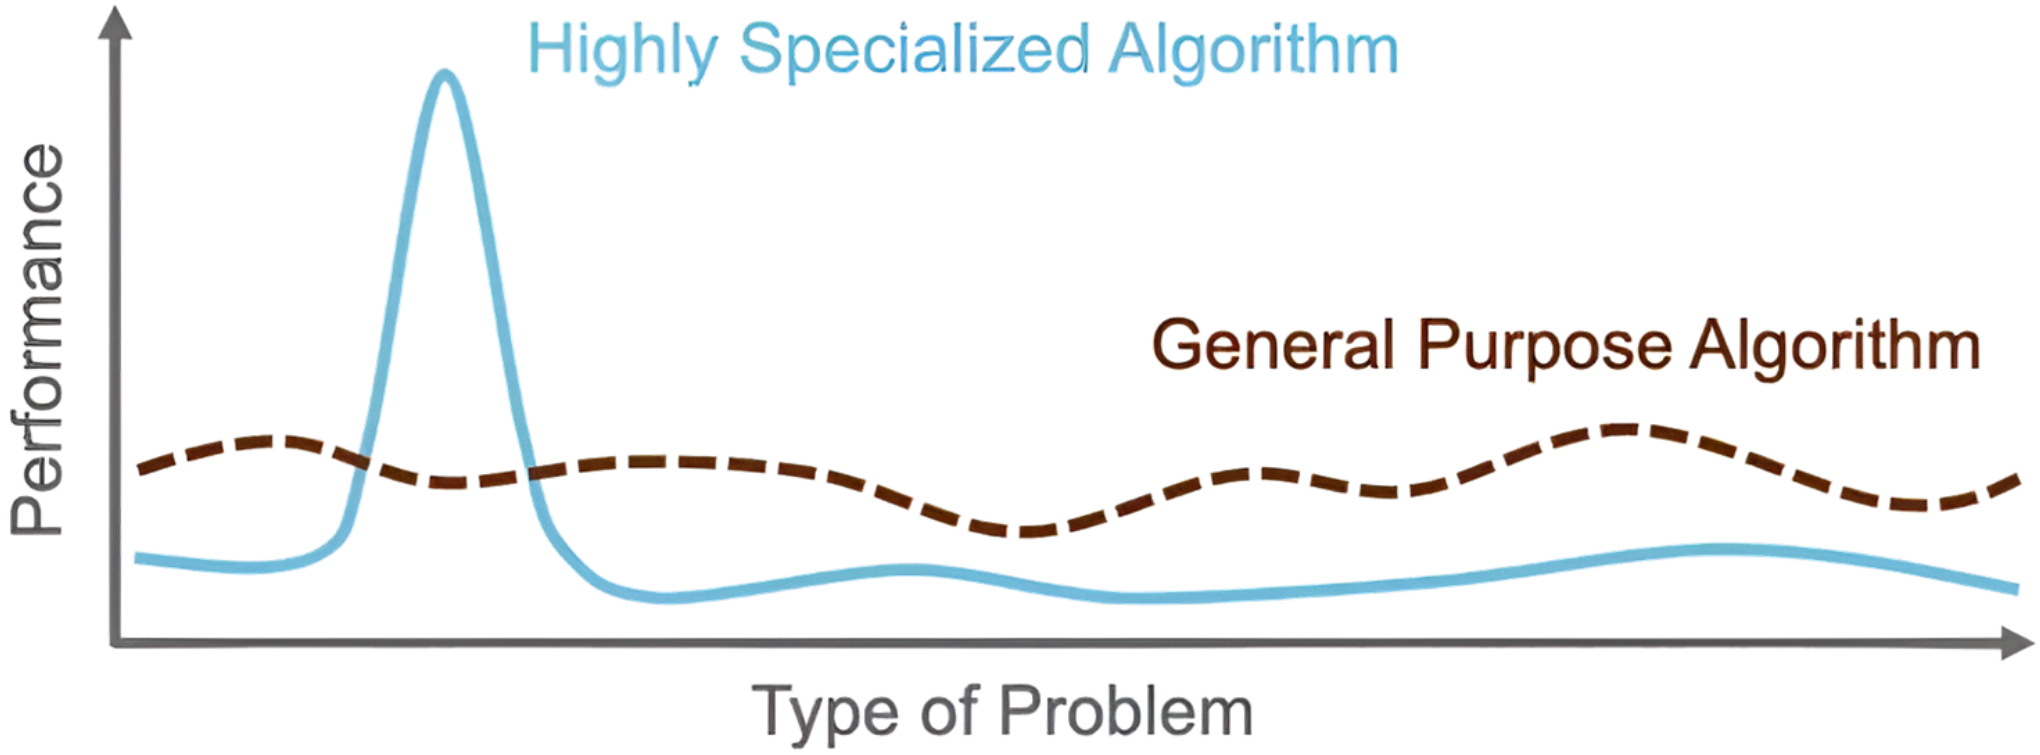

Ugly duckling theorem

- There is **no** privileged or **best** **feature** representation, and that even the notion of similarity between patterns depends implicitly on assumption that may or may not be correct. (Watanabe)

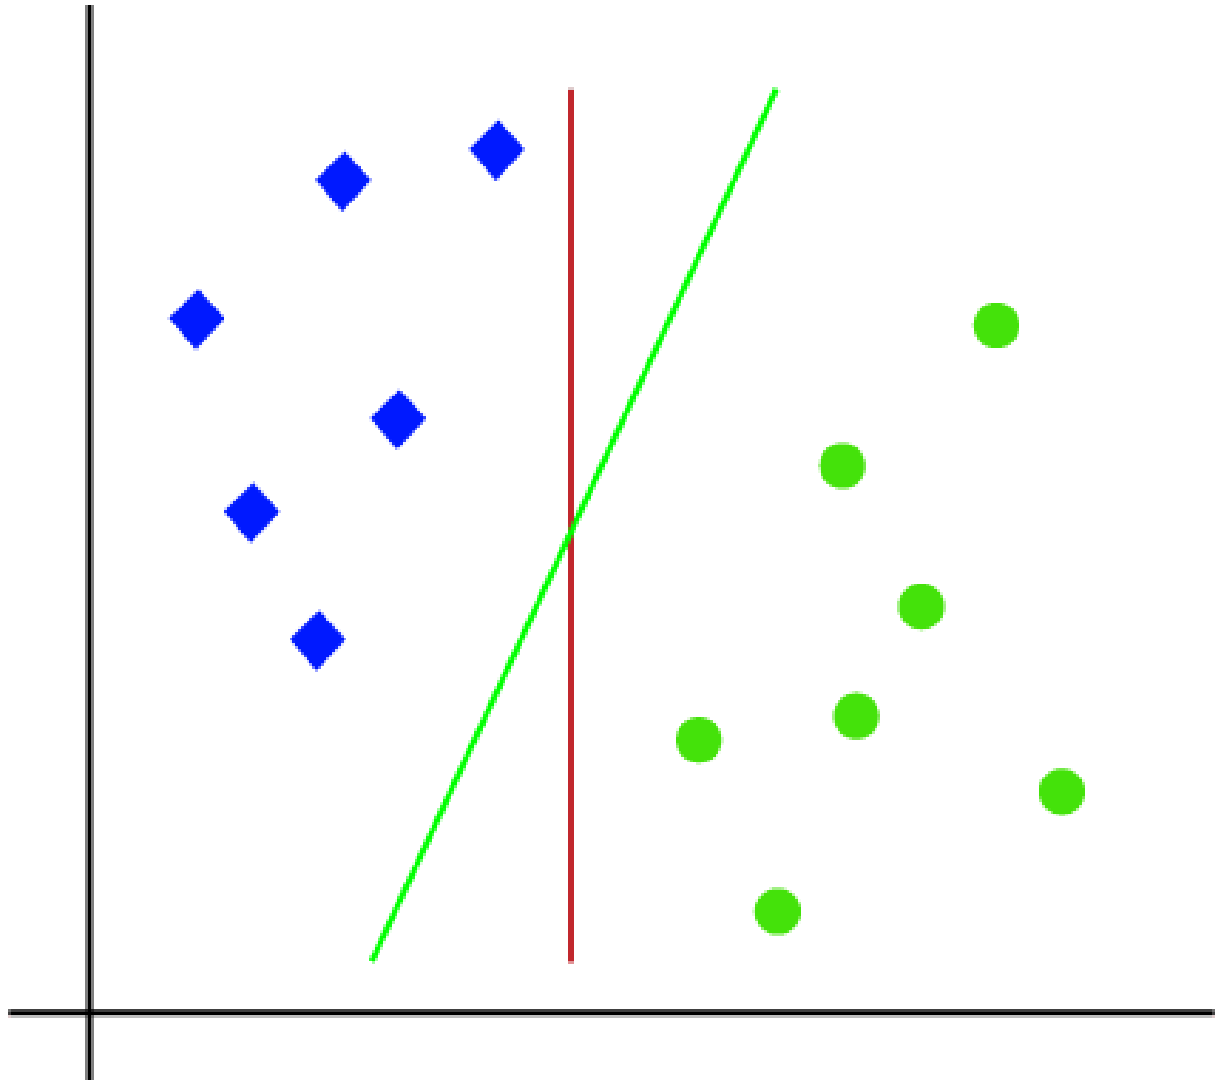

Curse of dimensionality

- Various phenomena that arise when analyzing data in high-dimensional spaces that do not occur in low-dimensional settings

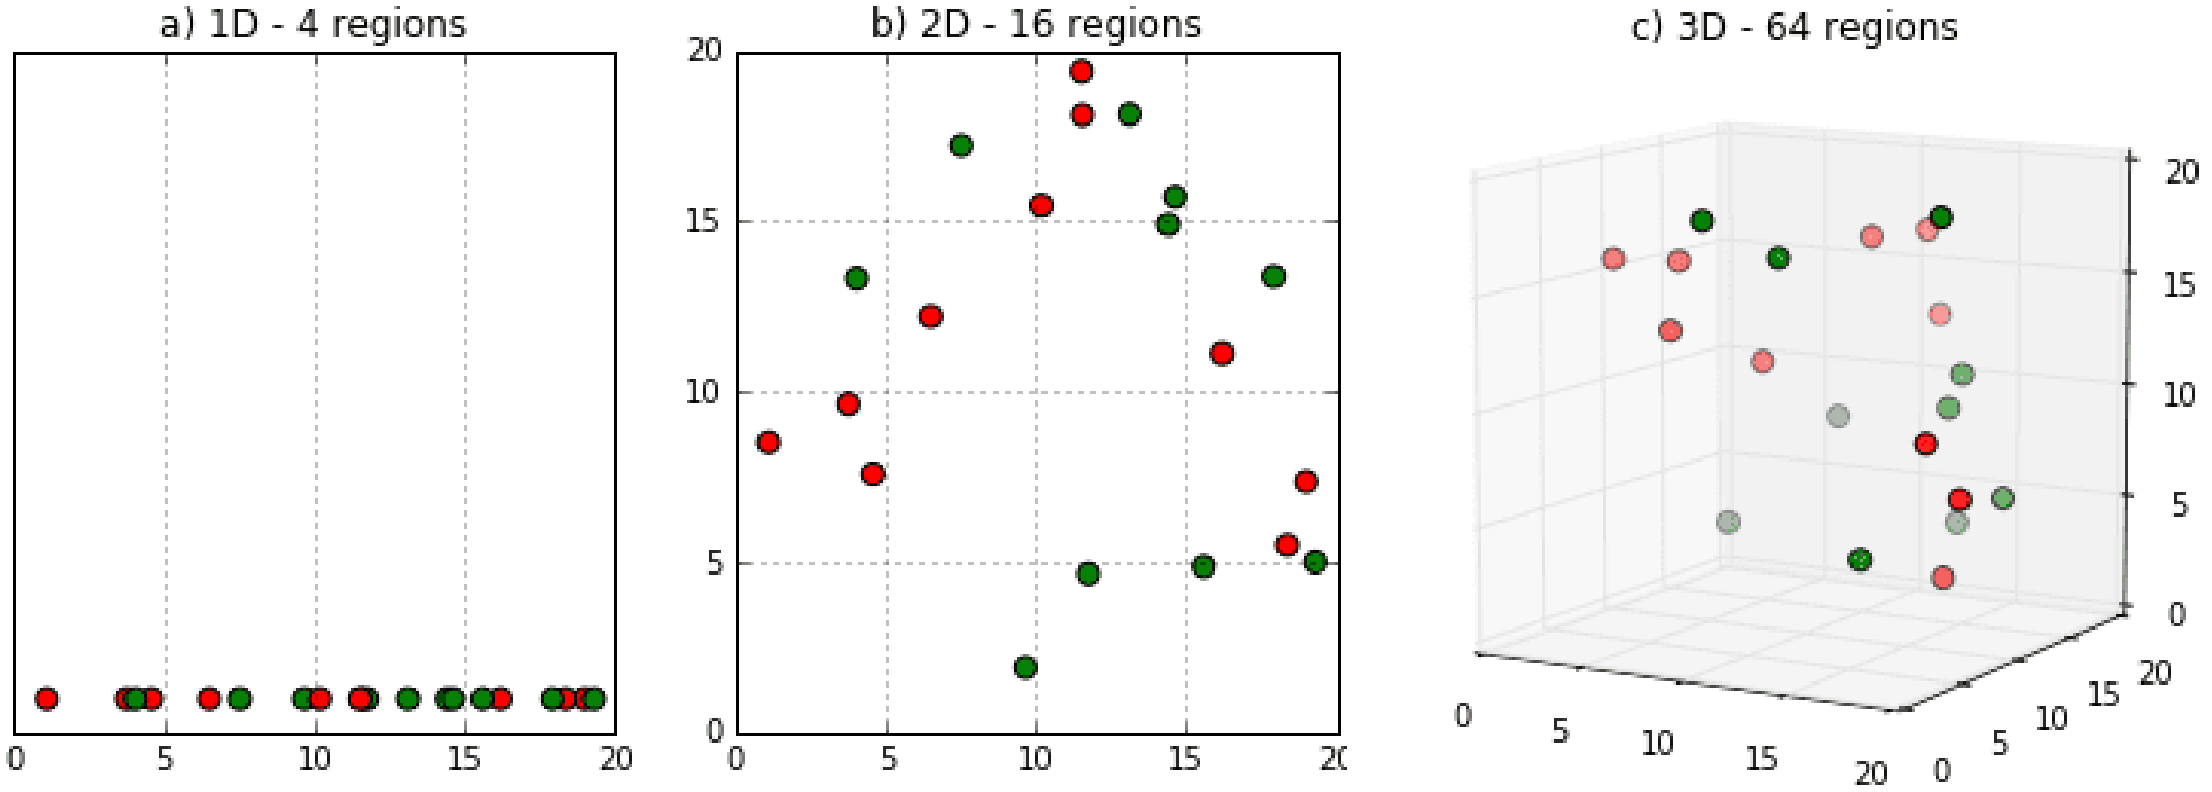

Bias & Variance

- Estimation Error

– Bias: difference between expected prediction and correct value

– Variance: variability of prediction for a given data point

– Trade-off

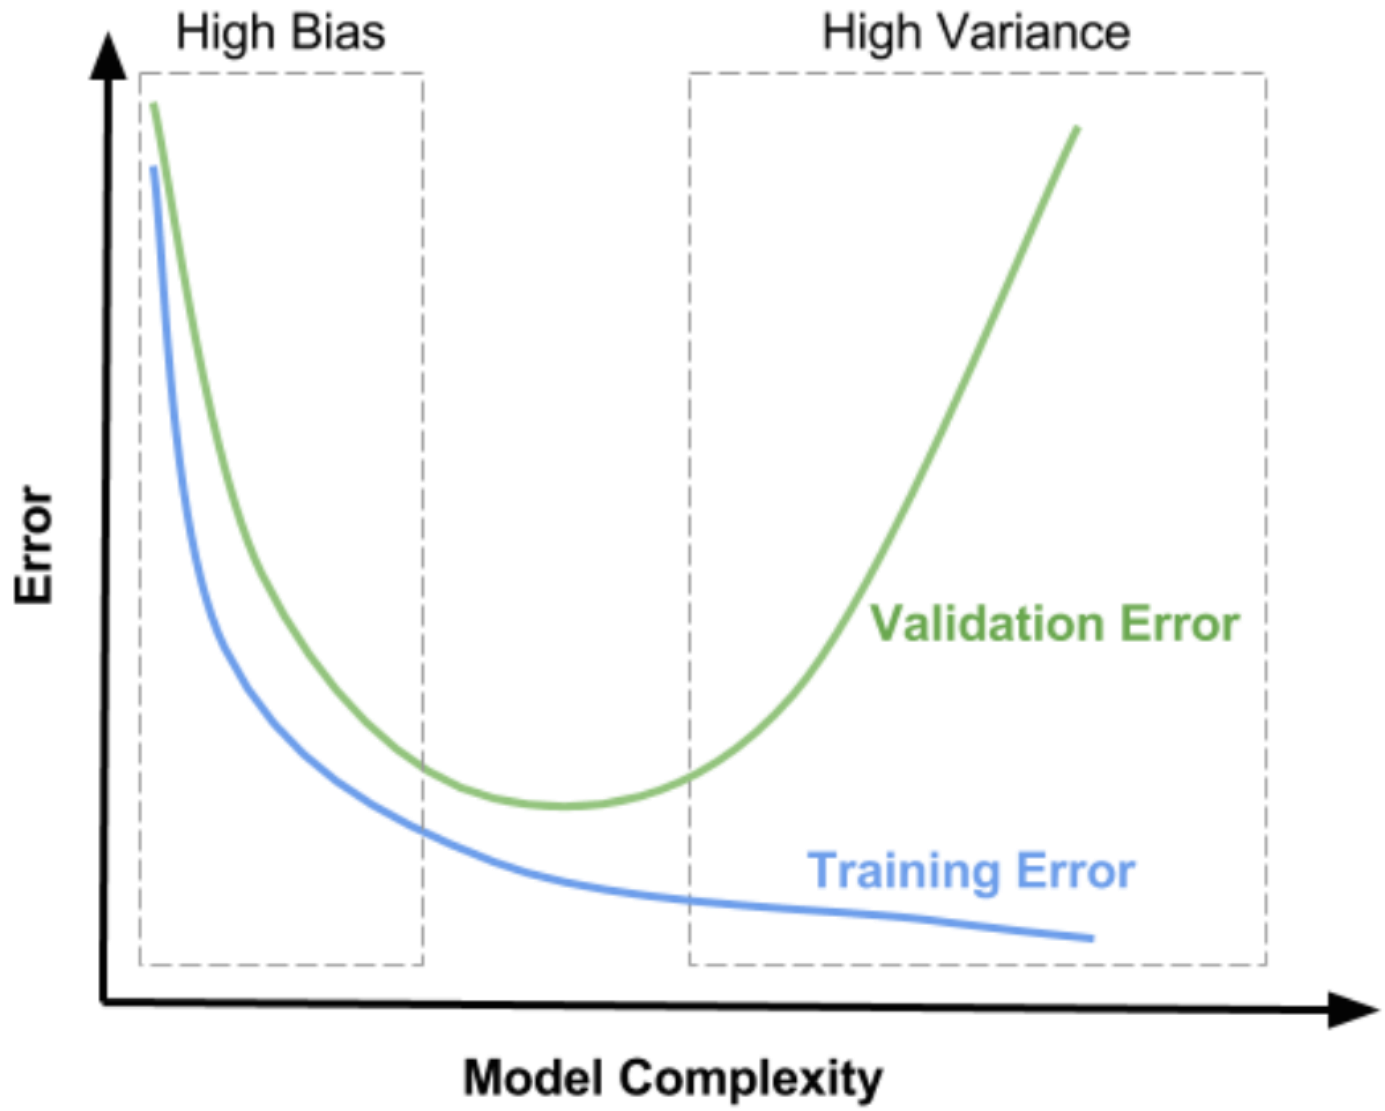

Machine Learning Algorithms

- Read more at:
  - https://www.javatpoint.com/machine-learning-algorithms
  - https://machinelearningcoban.com/

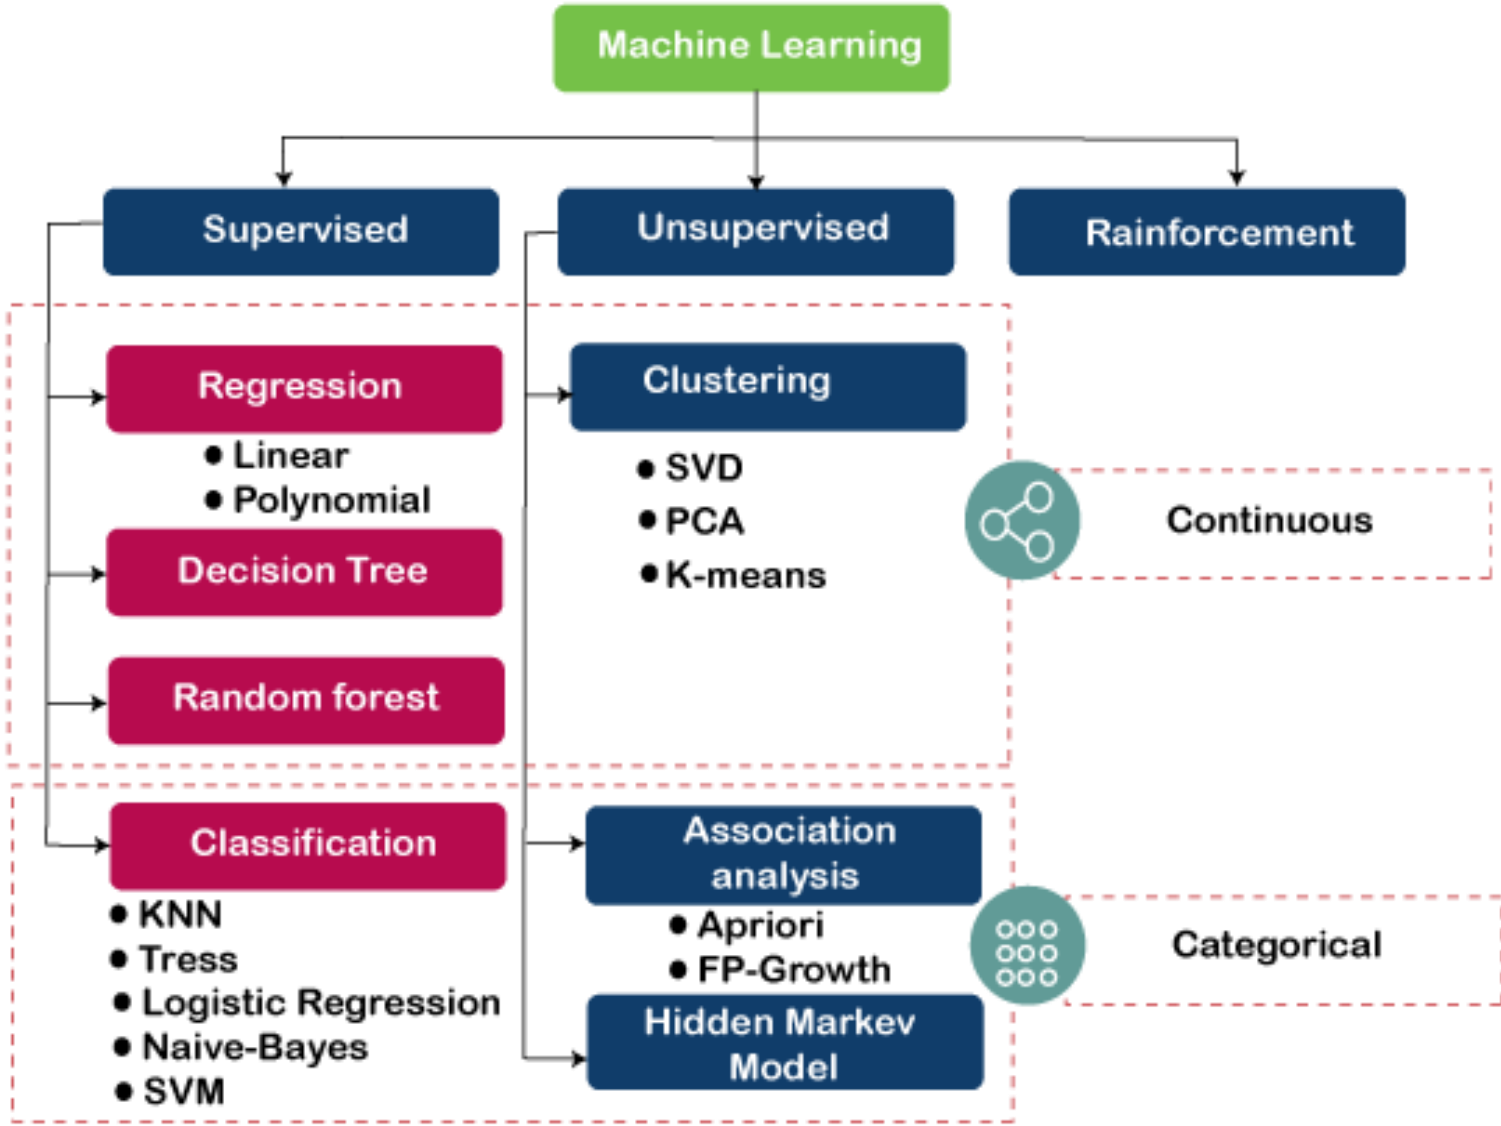

Basic Statistics

Reviews

	Population	Sample
# of subjects	N	n
Mean	$\mu = \frac{\sum_{i=1}^{N} x_i}{N}$	$\bar{x} = \frac{\sum_{i=1}^{n} x_i}{n}$
Variance	$\sigma^2 = \frac{\sum_{i=1}^{N} (x_i - \mu)^2}{N}$	$S^2 = \frac{\sum_{i=1}^{n} (x_i - \bar{x})^2}{n-1}$

Note: S^2 is the formula for unbiased sample variance, since we're dividing by $n-1$.

Standard deviation	$\sigma = \sqrt{\frac{\sum_{i=1}^{N} (x_i - \mu)^2}{N}}$	$S = \sqrt{\frac{\sum_{i=1}^{n} (x_i - \bar{x})^2}{n-1}}$

Note: Finding S by taking $\sqrt{S^2}$ reintroduces bias.

Exercise

Given a sample: X = [10, 7, -22, 4, -9]

1. Calculate the mean, variance and standard deviation
2. Normalize the data points into the range [0, 1]
3. Calculate the mean, variance and standard deviation after normalizing
4. Standardize X s.t. the mean of X is 0 and standard deviation is 1

Note: round the value up with 2 digits after the floating point (e.g., 10.23)

$$\text{Normalization: } X_{new} = \frac{X - X_{min}}{X_{max} - X_{min}}$$

$$\text{Standardization: } X_{new} = \frac{X - \bar{X}}{s}$$

Exercise: Solution

```
1.
Sample: [10, 7, -22, 4, -9]
Mean = -2
Variance = 177.5
Standard deviation = 13.32
Max value = 10
Min value = -22

2.
[10, 7, -22, 4, -9] after normalization: [1.0, 0.91, 0.0, 0.81, 0.41]

3.
Sample: [1.0, 0.90625, 0.0, 0.8125, 0.40625]
Mean = 0.62
Variance = 0.17
Standard deviation = 0.42
Max value = 1.0
Min value = 0.0

4.
[10, 7, -22, 4, -9] after standardization: [0.9, 0.68, -1.5, 0.45, -0.53]
```

Data Standardization

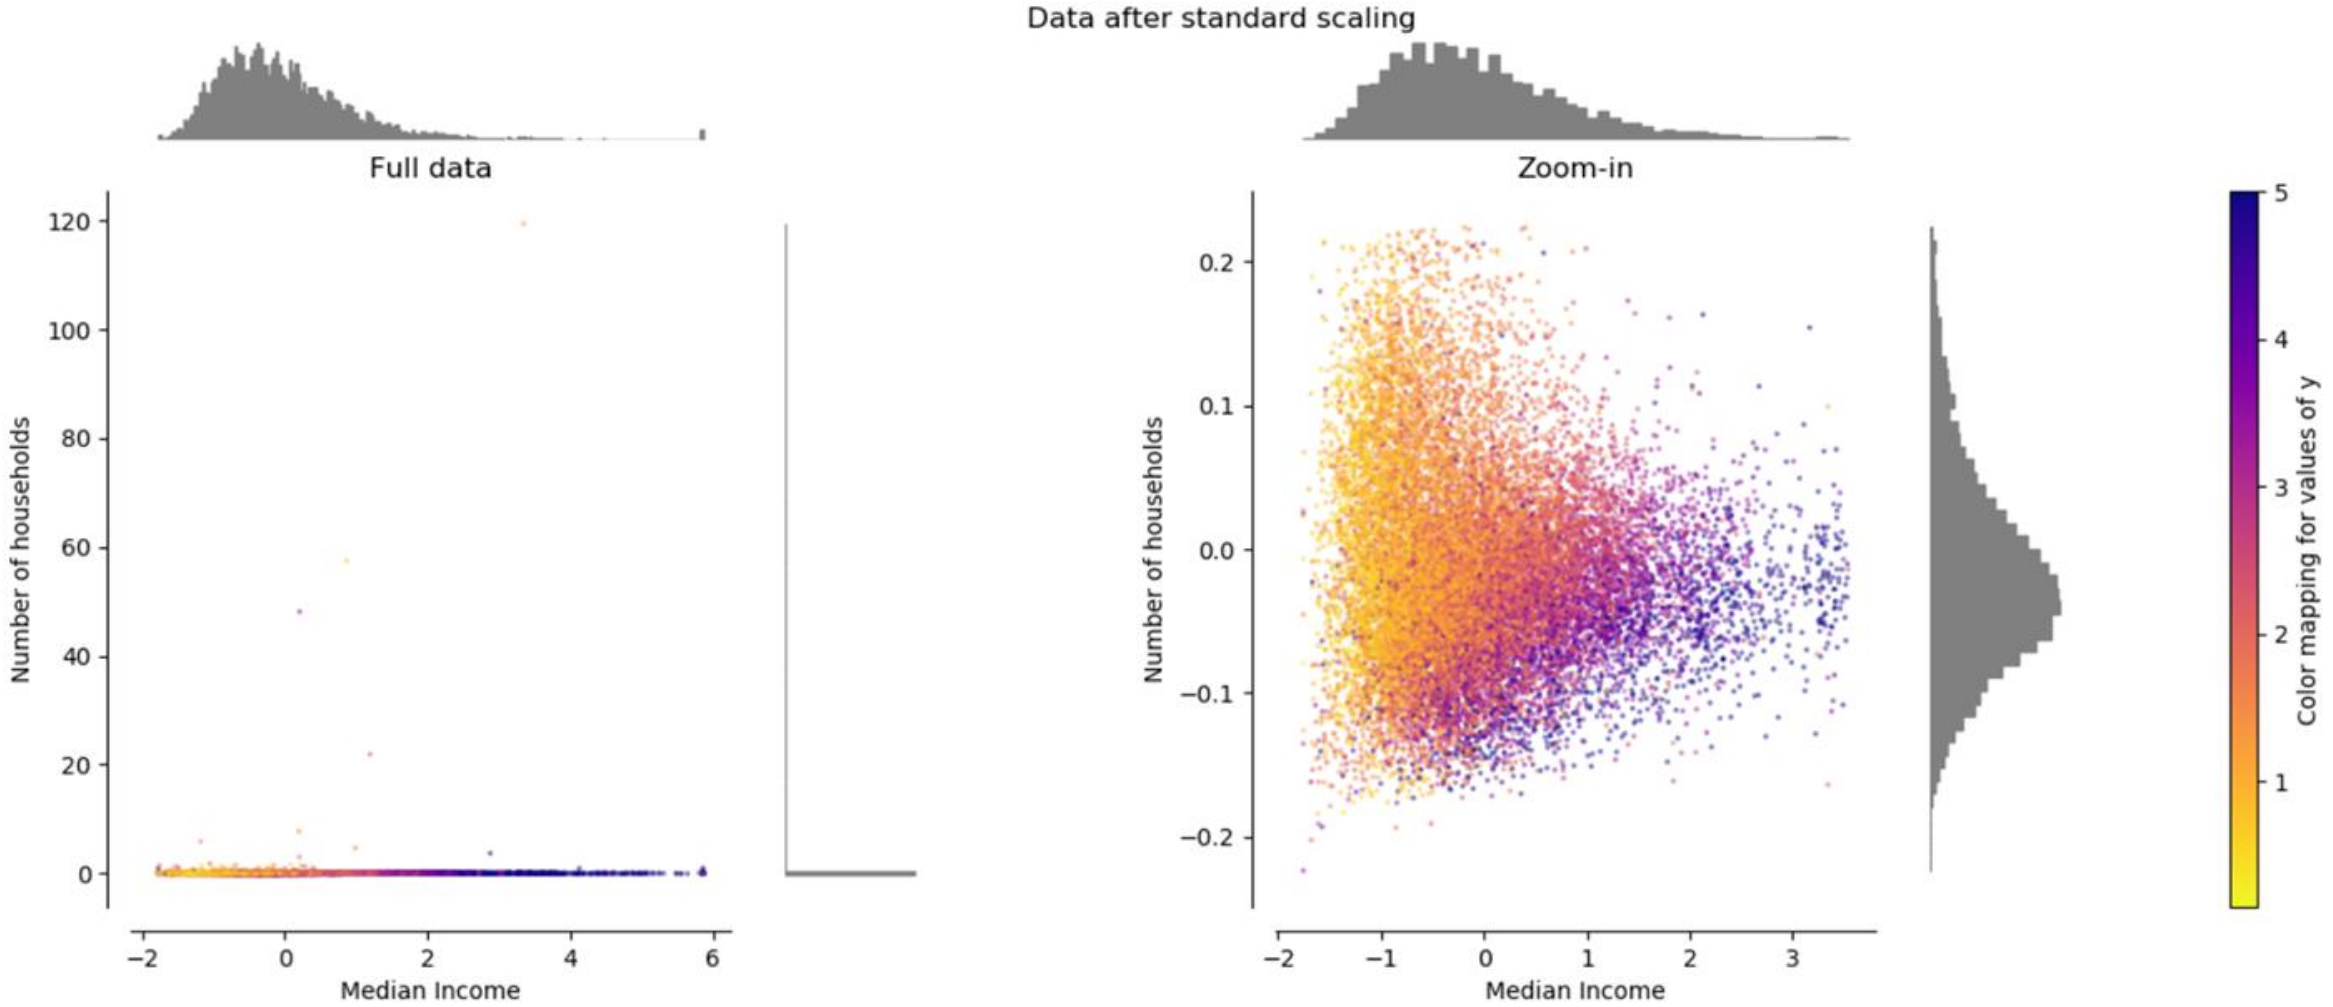

Source: https://scikit-learn.org/stable/auto_examples/preprocessing/plot_all_scaling.html

Data Standardization

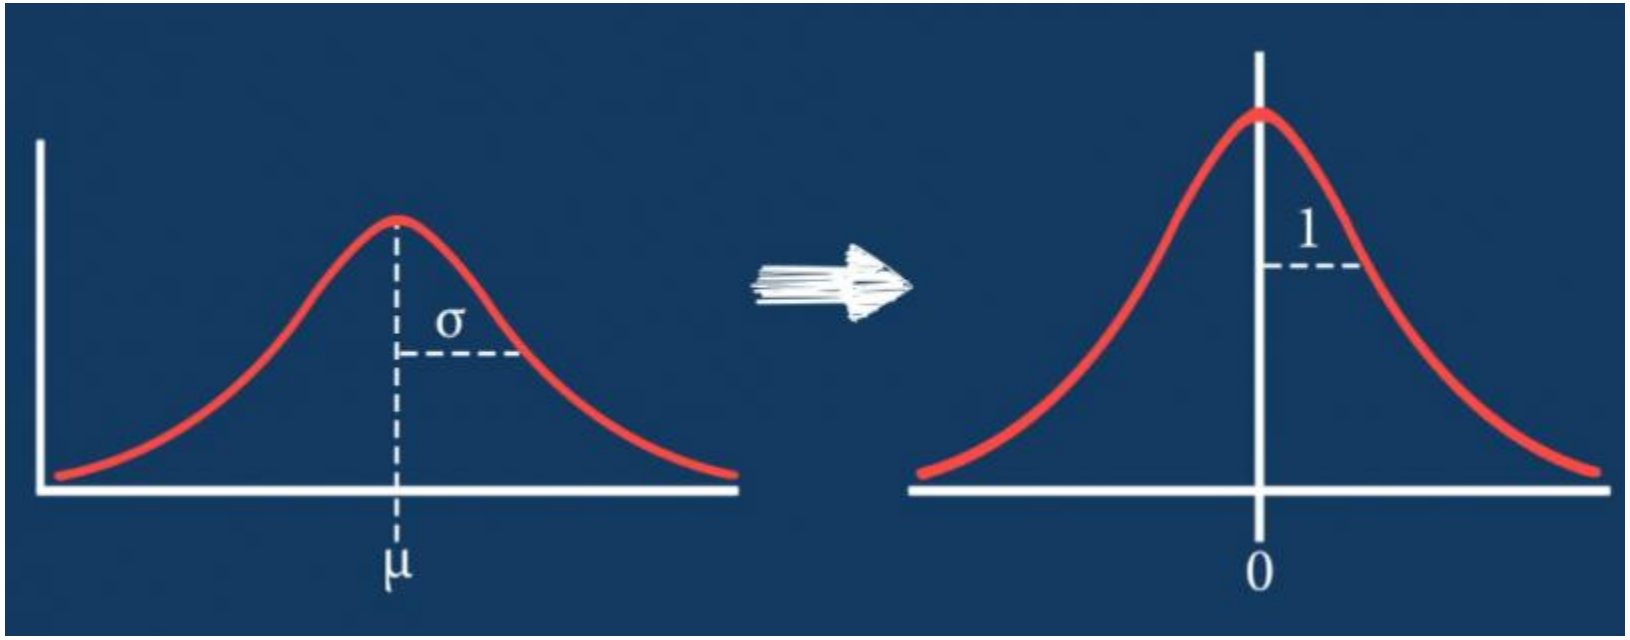

- Needed before statistic-based algorithm (e.g. clustering, PCA, SVM, ...)
- No need in regression and tree-based algorithms

Python: Introduction

Introduction

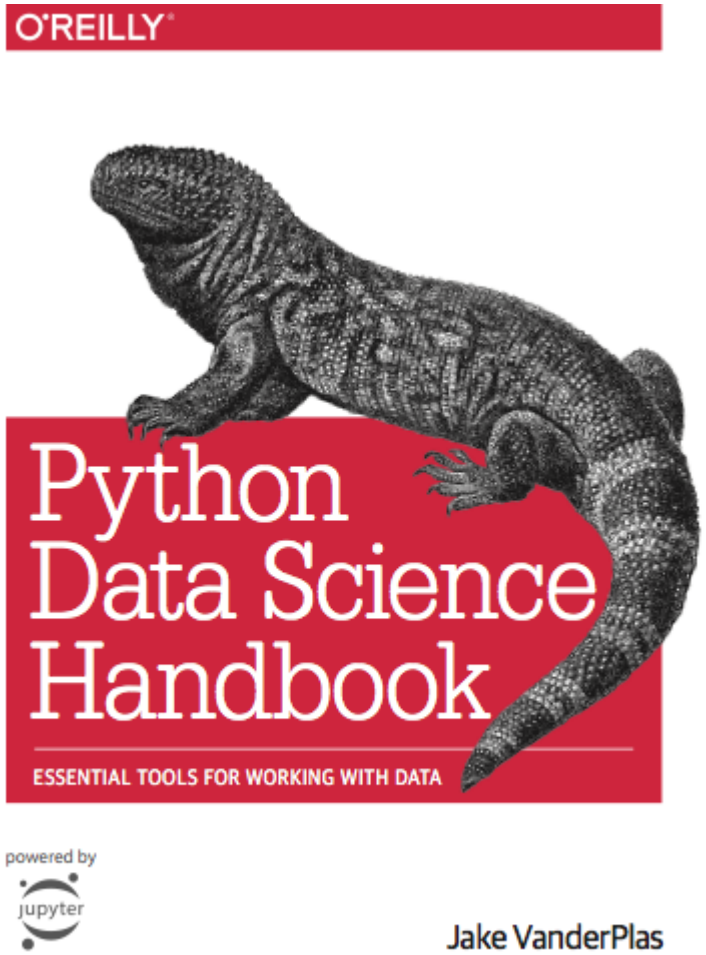

Table of Contents

Preface

1. IPython: Beyond Normal Python
   - Help and Documentation in IPython
   - Keyboard Shortcuts in the IPython Shell
   - IPython Magic Commands
   - Input and Output History
   - IPython and Shell Commands
   - Errors and Debugging
   - Profiling and Timing Code
   - More IPython Resources
2. Introduction to NumPy
   - Understanding Data Types in Python
   - The Basics of NumPy Arrays
   - Computation on NumPy Arrays: Universal Functions
   - Aggregations: Min, Max, and Everything In Between
   - Computation on Arrays: Broadcasting
   - Comparisons, Masks, and Boolean Logic
   - Fancy Indexing
   - Sorting Arrays
   - Structured Data: NumPy's Structured Arrays
3. Data Manipulation with Pandas

Link:

https://colab.research.google.com/github/jakevdp/PythonDataScienceHandbook/blob/master/notebooks/Index.ipynb

Keras

Pytorch

Google Colab

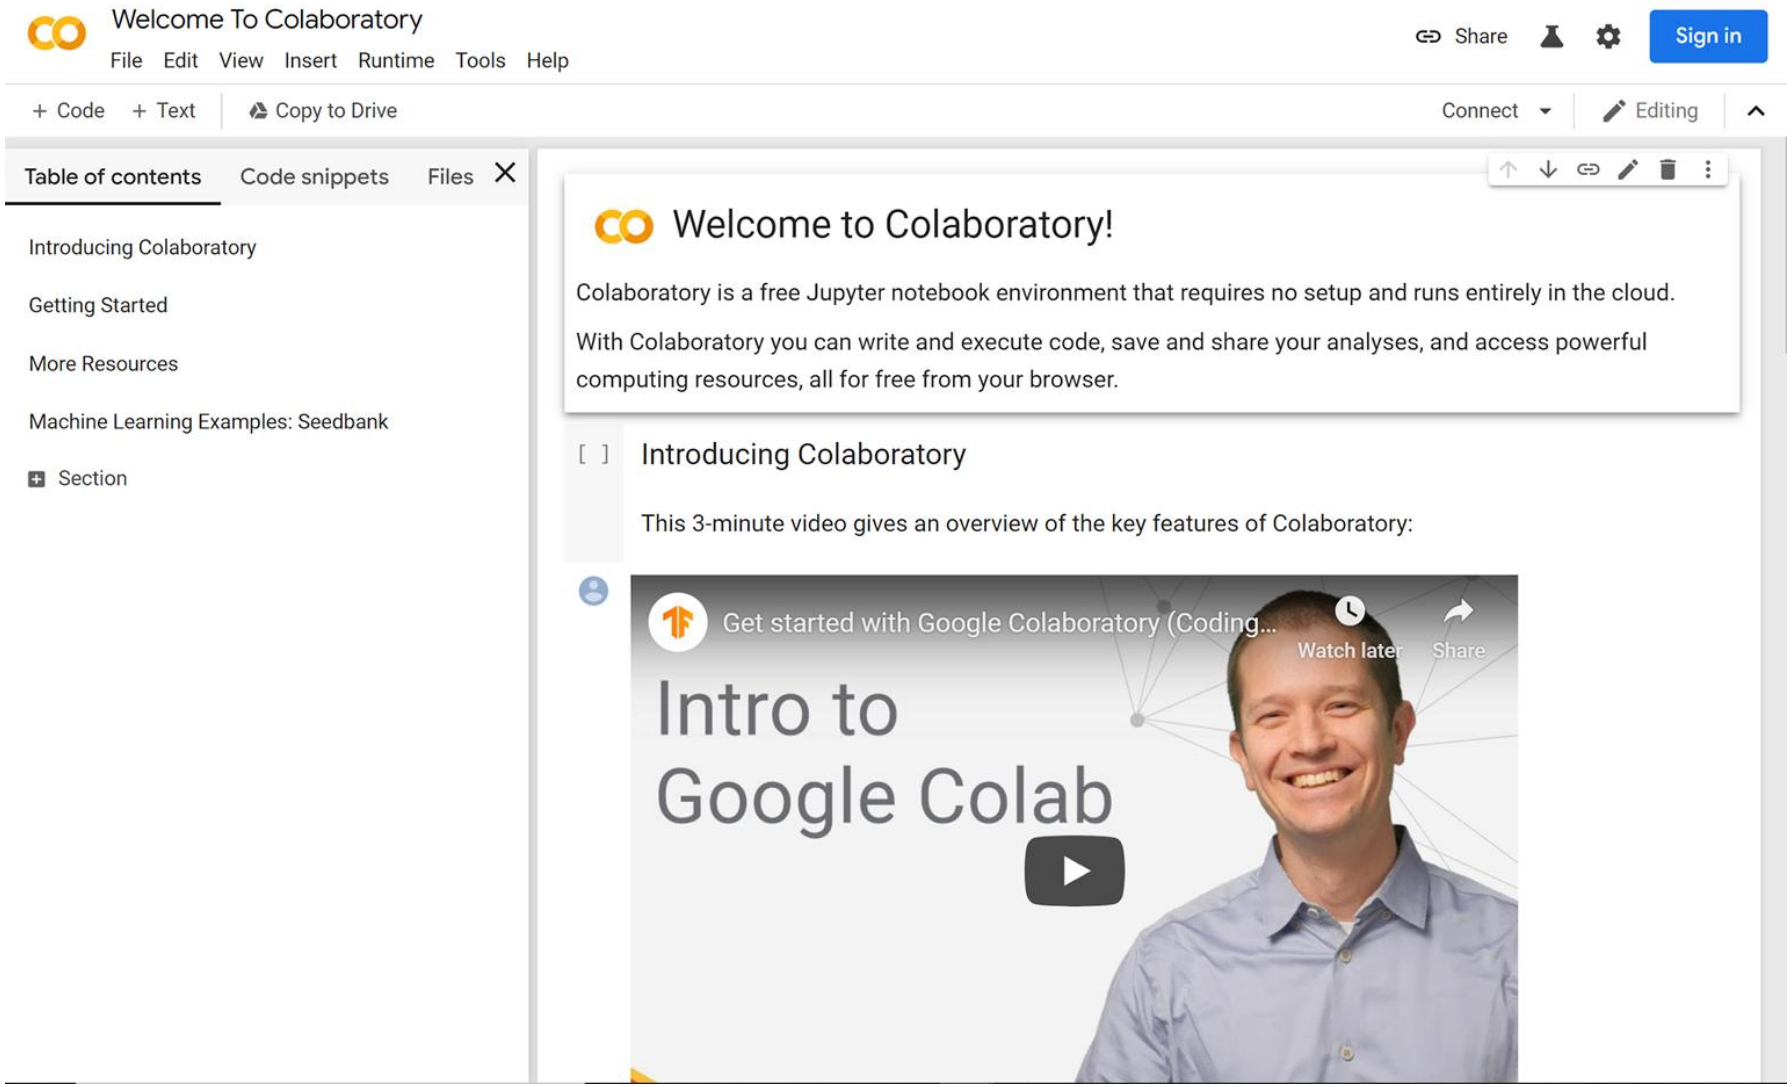

THANK YOU
for YOUR ATTENTION