Natural Language Processing Applications

Week 6: Text Similarity

fit@hcmus
KHOA CÔNG NGHỆ THÔNG TIN
TRƯỜNG ĐẠI HỌC KHOA HỌC TỰ NHIÊN

fit@hcmus

- Introduction to Similarity
- Text Similarity
- Similarity Evaluation

fit@hcmus

Applied Natural Language Processing- Text Similarity

INTRODUTION TO SIMILARITY

fit@hcmus

Introdution to Similarity

- Word similarity: finding similarities between words is a fundamental part of text similarity
- Words are considered similar if they :
  - Have the same meaning(Synonyms)
  - Have opposite meanings(Antonyms)
  - Are used in the same way(For example: red, green…)
  - Are used in the same context(For example: doctor, hospital…)
  - Are type of another word (For example: fluffy dog, dog, animal…)

fit@hcmus

Introdution to Similarity(continued)

- Text Similarity:

Hurricane Gilbert swept toward the Dominican Republic Sunday , and the Civil Defense alerted its heavily populated south coast to prepare for high **winds**, heavy **rains** and high seas.

The **storm** was approaching from the southeast with sustained **winds** of 75 mph gusting to 92 mph .

" There is no need for alarm," Civil Defense Director Eugenio Cabral said in a television alert shortly before midnight Saturday .

Cabral said residents of the province of Barahona should closely follow **Gilbert** 's movement .

An estimated 100,000 people live in the province, including 70,000 in the city of Barahona , about 125 miles west of Santo Domingo .

Tropical **Storm Gilbert** formed in the eastern Caribbean and strengthened into a **hurricane** Saturday night

The National **Hurricane** Center in Miami reported its position at 2a.m. Sunday at latitude 16.1 north , longitude 67.5 west, about 140 miles south of Ponce, Puerto Rico, and 200 miles southeast of Santo Domingo.

The National Weather Service in San Juan , Puerto Rico , said **Gilbert** was moving westward at 15 mph with a "broad area of cloudiness and heavy weather" rotating around the center of the **storm**.

The weather service issued a flash flood watch for Puerto Rico and the Virgin Islands until at least 6p.m. Sunday.

Strong **winds** associated with the **Gilbert** brought coastal flooding , strong southeast **winds** and up to 12 feet to Puerto Rico 's south coast.

fit@hcmus

Introduction to Similary(continued)

- Text Similary:
  - Document1
    - Gilbert: 3
    - Hurricane: 2
    - Rains: 1
    - Storm: 2
    - Winds: 2
  - Document2
    - Gilbert: 2
    - Hurricane: 1
    - Rains: 0
    - Storm: 1
    - Winds: 2

Cosine similarity: 0.9439

fit@hcmus

Introduction to Similary(continued)

- According to John Philip McCrae:
  - “Semantic textual similarity is the task of deciding if two sentences express a similar or identical meaning and requires a deep understanding of a sentence and its meaning in order to achieve high performance.”

fit@hcmus

Introduction to Similarity(continued)

- Applications of Text Similarity:
  - Information Retrieval
  - Text Summarization
  - Machine Translation
  - Plagiarism Detection
  - ...

fit@hcmus

Introduction to Similarity(continued)

- Information Retrieval

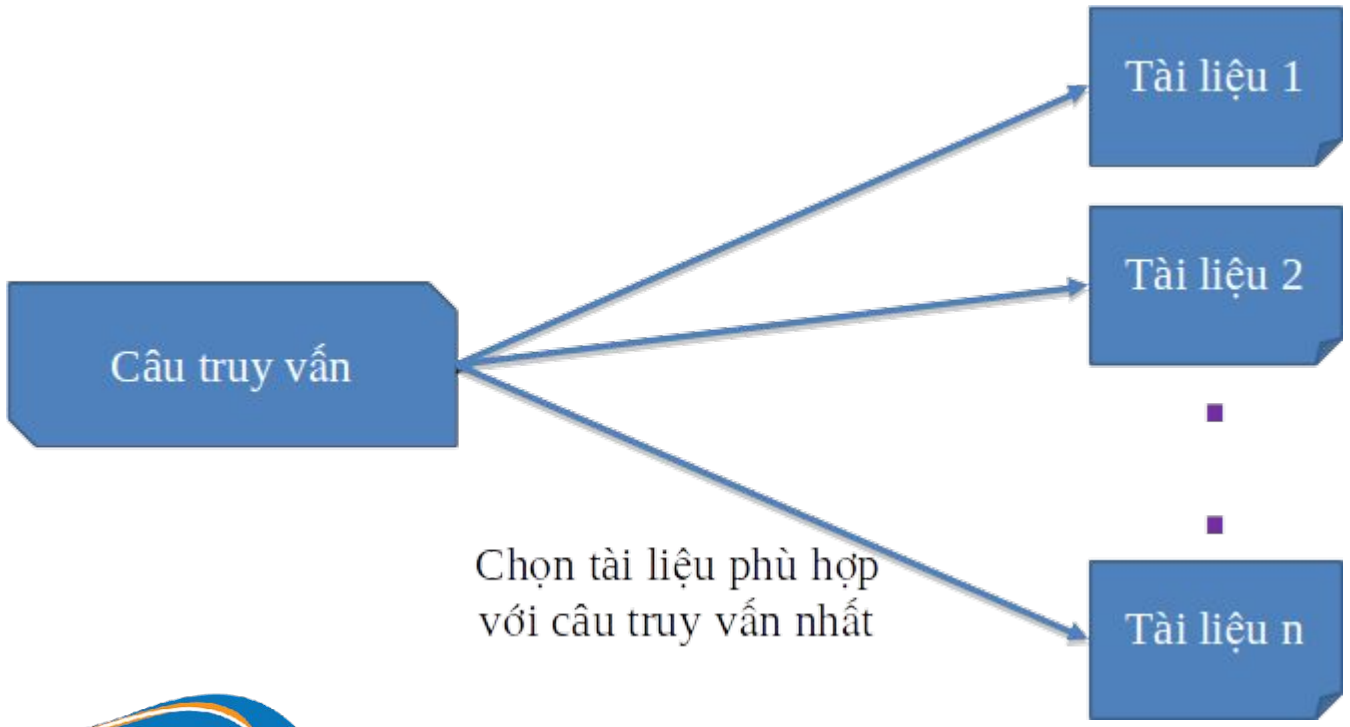

fit@hcmus

Introduction to Similary(continued)

Text Summarization

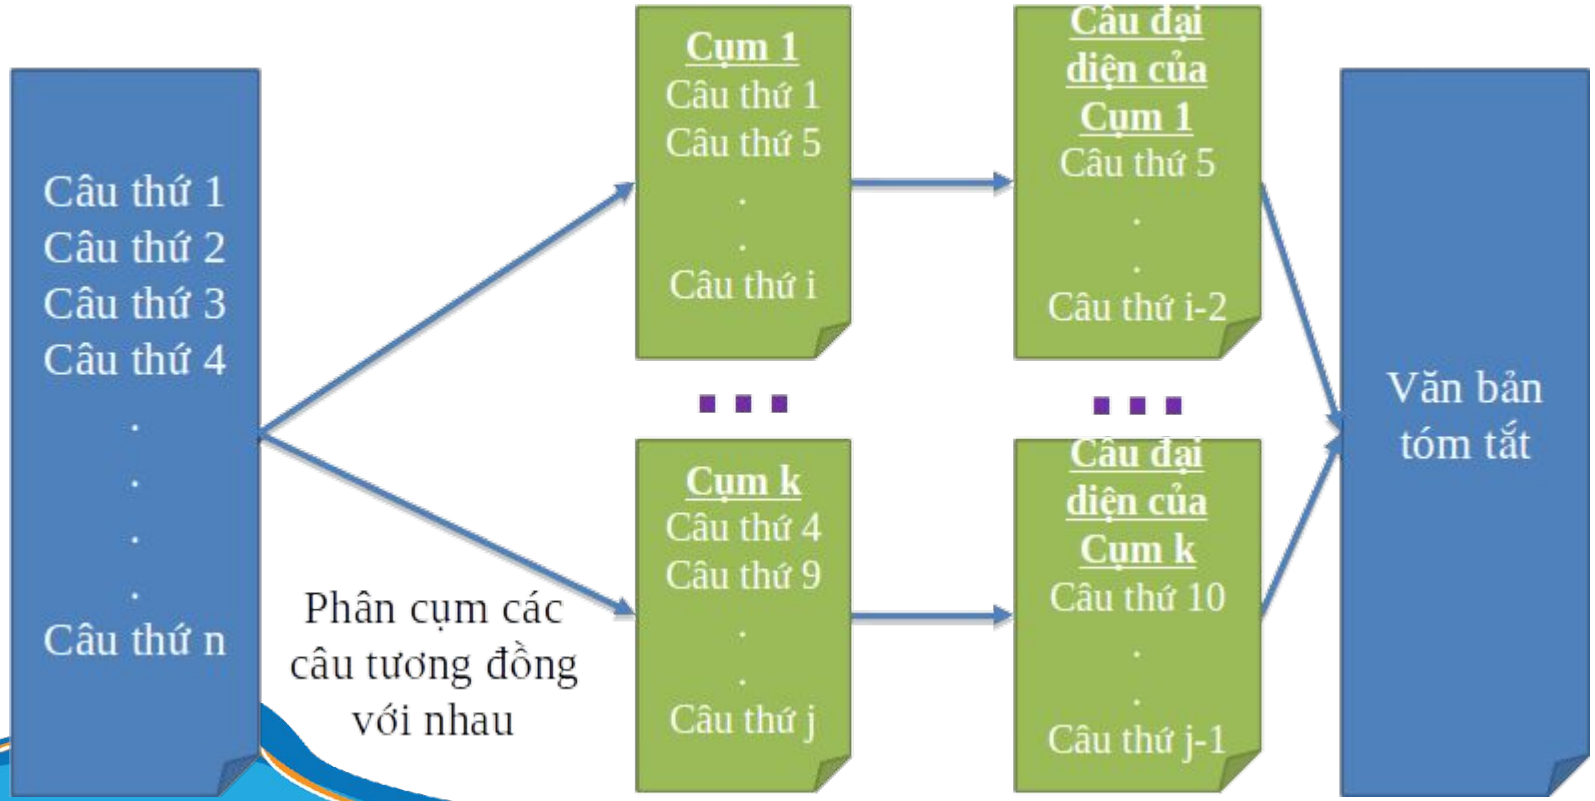

fit@hcmus

Introduction to Similarity(continued)

- Machine Translation (Quality evaluation)

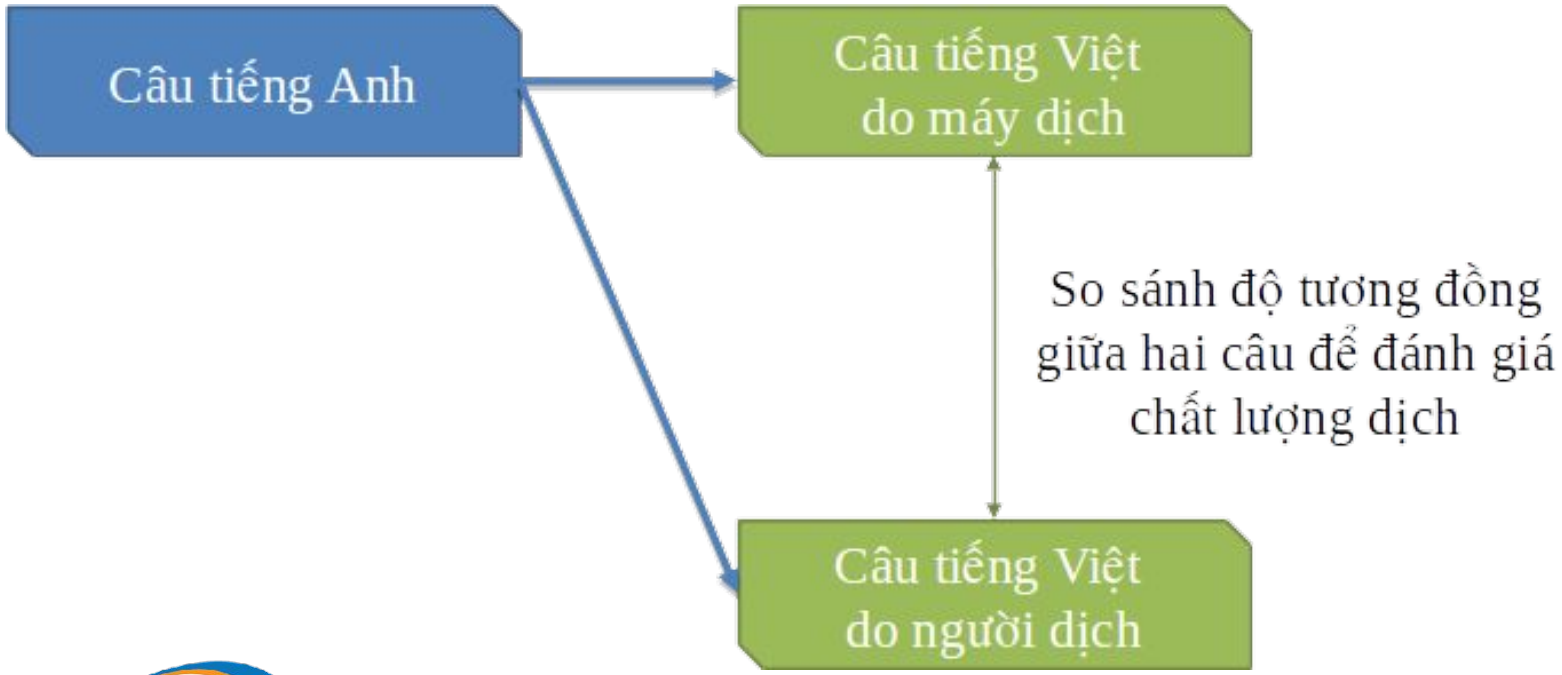

fit@hcmus

Introduction to Similarity(continued)

- Plagiarism Detection

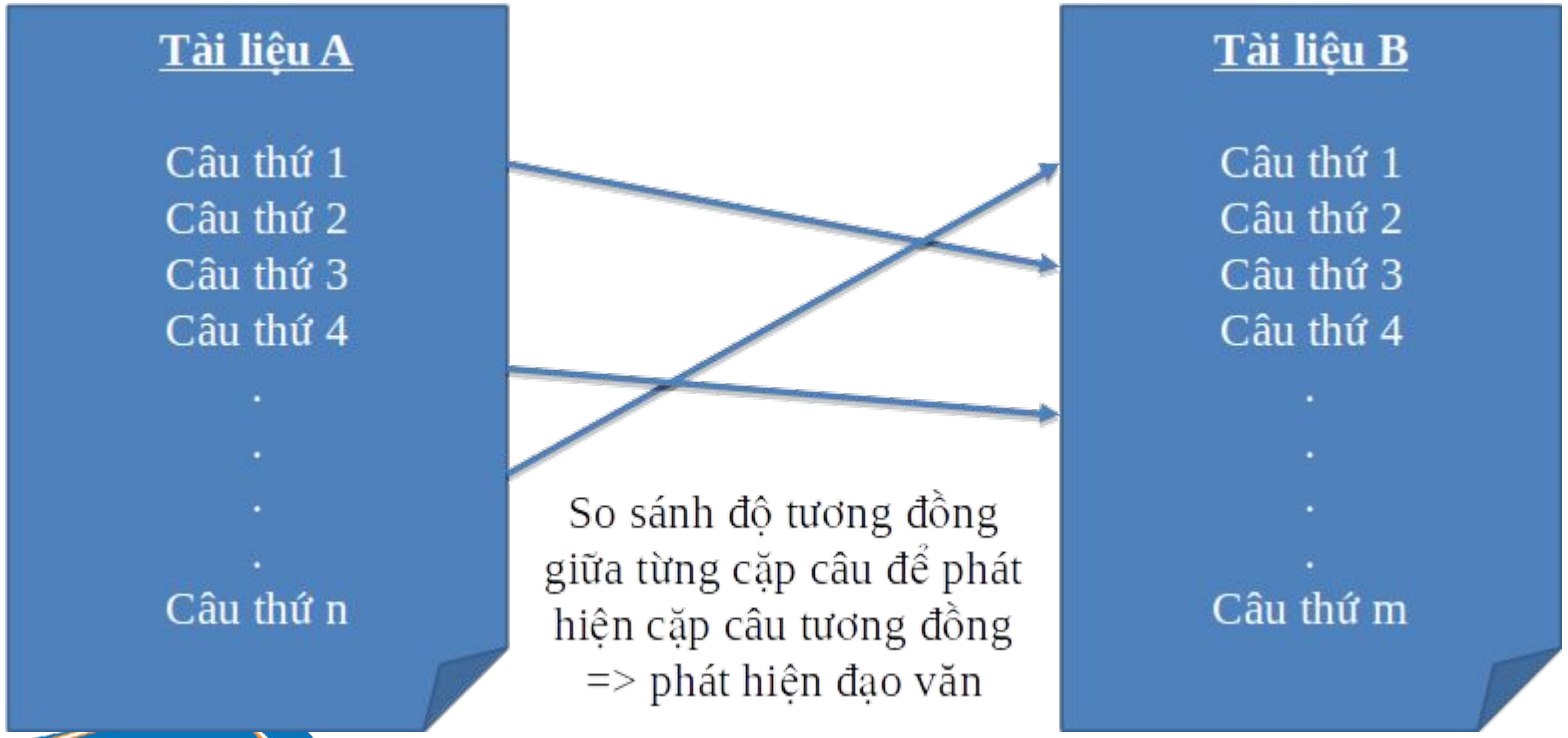

fit@hcmus

Introduction to Similary(continued)

- There are 2 popular scales :
  - Scale 0 to 1 is used in detecting rewrites text while retaining full meaning(paraphrase identification)
  - Scale 0 to 5:

fit@hcmus

Introduction to Similary(continued)

Scale 0 to 1:

- The following pair of sentences is labeled as 1
  - Sentance 1: Customers will have to use a decoder card from the cable TV provider to plug in the set.
  - Sentance 2: To watch pay TV, customers will plug in the television set a decoder card provided by the cable TV provider.
- And the following pair of sentences is labeled as 0
  - Sentance 1: With an interpreter like you everything will be fine.
  - Sentance 2: Today's interpreter is Mr. Nam.

fit@hcmus

Introduction to Similary(continued)

- Scale 0 to 5:

5	**Hai câu tương đồng hoàn toàn**
	Con chim đang tắm trong bồn rửa. Con chim non đang tắm trong bồn nước.
4	**Hai câu tương đồng phần lớn, nhưng khác nhau vài chi tiết không quan trọng.**
	Hai chàng trai đang chơi trò chơi điện tử trên một chiếc ghế dài. Hai chàng trai đang chơi trò chơi điện tử.
3	**Hai câu gần tương đồng, nhưng khác nhau hoặc thiếu một vài thông tin quan trọng.**
	John cho biết anh ấy được xem là một nhân chứng chứ không phải là một nghi phạm. “Anh ấy không phải là kẻ tình nghi nữa.” John nói.
2	**Hai câu không tương đồng, nhưng có chung một vài thông tin.**
	Chúng bay ra khỏi tổ theo từng nhóm. Chúng cùng bay vào tổ.
1	**Hai câu không tương đồng, nhưng có cùng chung chủ đề.**
	Người phụ nữ đang chơi đàn vĩ cầm. Người phụ nữ trẻ thích nghe đàn ghita.
0	**Hai câu hoàn toàn khác nhau.**
	Con chó đang chạy trên tuyết. Một người lái xe đua đang lái xe của mình qua bãi bùn.

fit@hcmus

Applied Natural Language Processing- Text Similarity

TEXT SIMILARITY METHODS

fit@hcmus

Text Similarity methods

- Text Distance
- Text Representation

fit@hcmus

Text Similarity methods

- Text distance: describes the semantic proximity of two text words from the perspective of distance

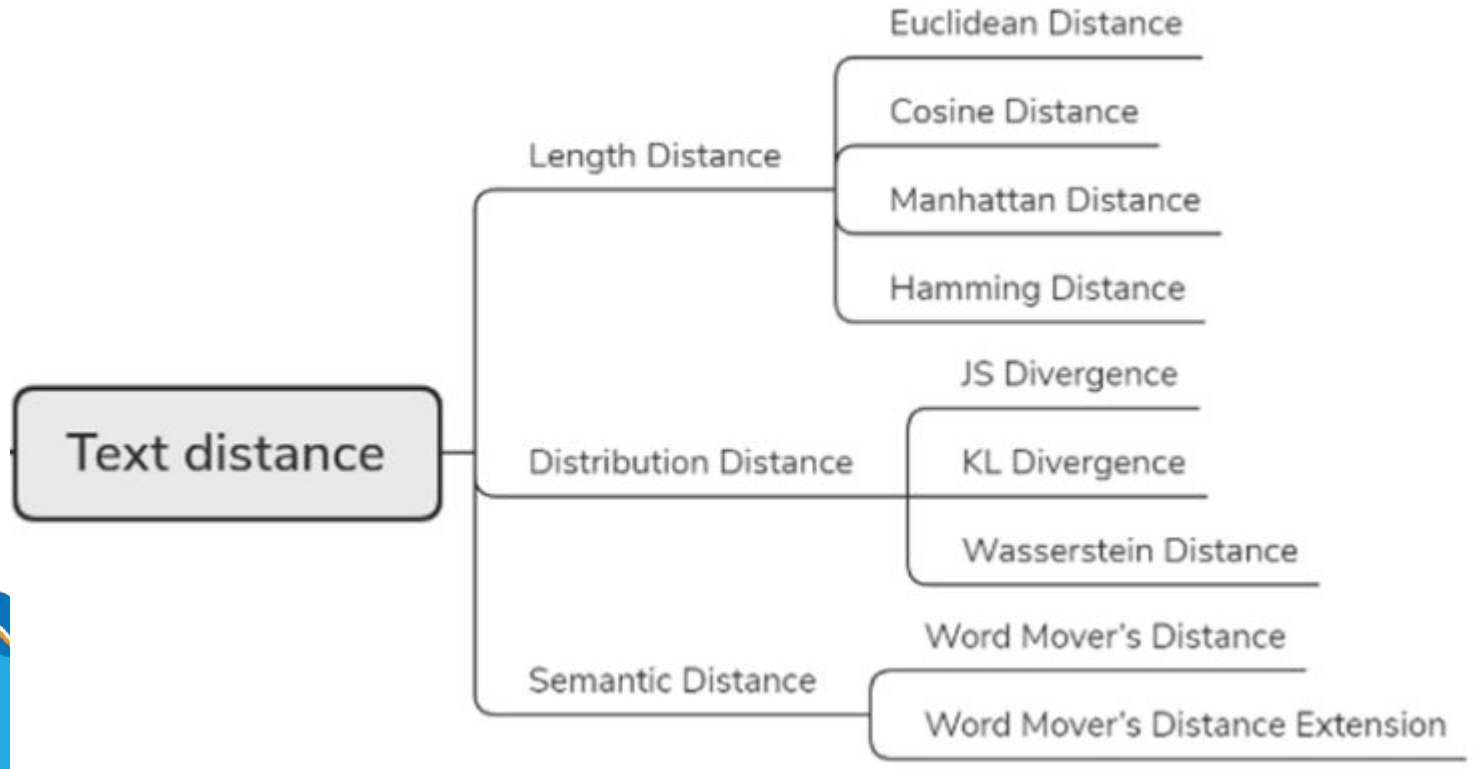

fit@hcmus

Text Similarity methods

- Text representation:

  numerically represent the unstructured text documents to make them mathematically computable.

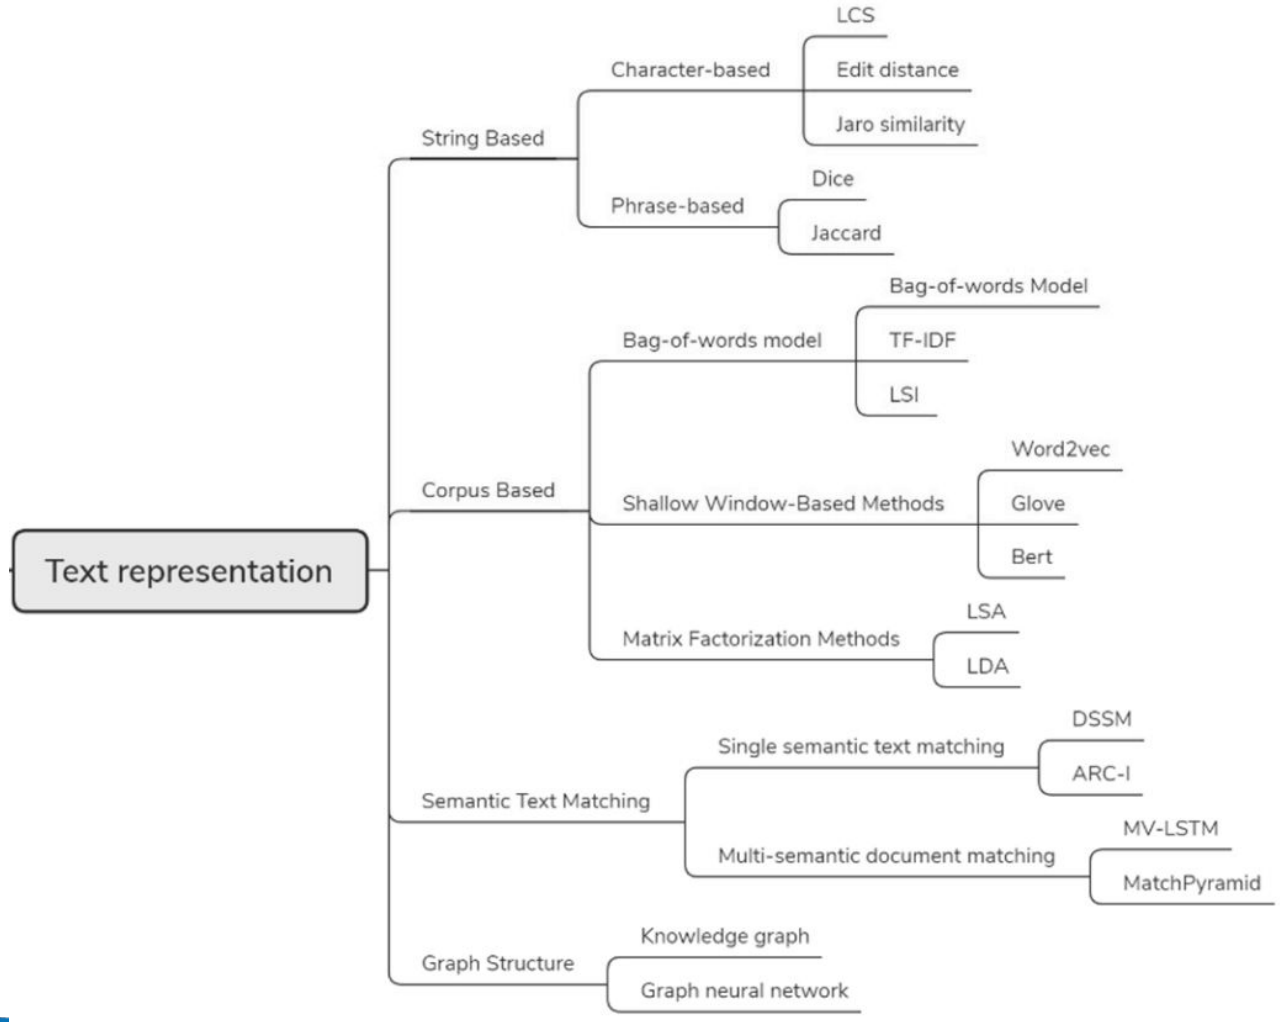

fit@hcmus

Text Distance

- Main methods:
  - Length Distance
  - Distribution Distance
  - Semantic Distance

fit@hcmus

Text Distance

- Length Distance:
  - Euclidean Distance

$$d(S_a, S_b) = \sqrt{\sum_{i=1}^{n} \left(S_a^{(i)} - S_b^{(i)}\right)^2}$$

  - Cosine Distance

$$\text{Sim}(S_a, S_b) = \cos\Theta = \frac{\vec{S_a} \cdot \vec{S_b}}{\|S_a\| \cdot \|S_b\|}$$

  - Manhattan Distance

$$\text{Sim}(x, y) = |x_1 - x_2| + |y_1 - y_2|$$

fit@hcmus

Text Distance

- Length Distance:
  - Suitable for symmetrical problems
    - Sim(A, B) = Sim(B, A)
    - => But for question Q to retrieve answer A, the corresponding similarity is not symmetrical.
  - Lack of statistical characteristics of the data

fit@hcmus

Text Distance

- Distribution Distance:
  - Kullback–Leibler Divergence

$$\mathrm{d}(\mathrm{p}\|\mathrm{q}) = \sum_{i=1}^{n} p(x) log \frac{p(x)}{q(x)}$$

  - Jensen–Shannon Divergence

$$\mathrm{JS}(P_1\|P_2) = \frac{1}{2} KL(P_1\| \frac{P_1 + P_2}{2}) + \frac{1}{2} KL(P_2\| \frac{P_1 + P_2}{2})$$

  - Wasserstein Distance

$$W(p_r, p_g) = \inf_{\gamma \sim \prod (p_r, p_g)} E_{(x,y)\sim\gamma}[\|x - y\|]$$

fit@hcmus

Text Distance

- Semantic Distance :
  - Word Mover's Distance

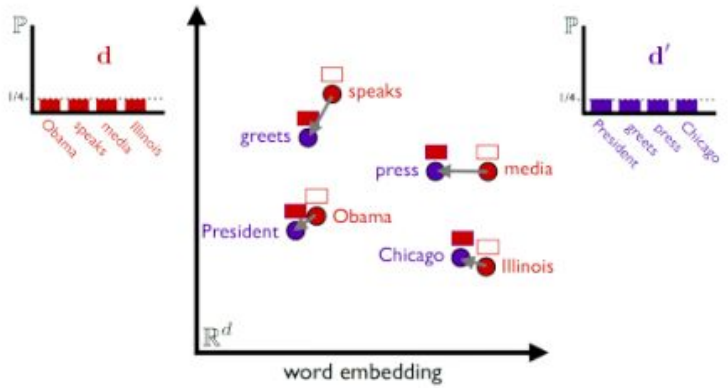

  - Word Mover's Distance Extension
    - Use the Mahalanobis distance instead of the Euclidean distance

Text representation

- String-based:
  - Operate on string sequences and character composition
  - Includes 2 methods:
    - Character-based
    - Phrase-based

fit@hcmus

Text representation

- String based:
  - Longest common substring:

$$\mathrm{LCS}(S_a, S_b) = \begin{cases} 0, \ if \ S_a = 0 \ or \ S_a = 0 \\ 1 + LCS(S_a - 1, S_b - 1), if \ x[S_a] == y[S_b] \\ max\begin{cases} LCS(S_a, S_b - 1) \\ LCS(S_a - 1, S_b) \end{cases} \ if \ x[S_a] \neq y[S_b] \end{cases}$$

  - Jaro Similarity:

$$\mathrm{Sim} = \begin{cases} 0, if \ m = 0 \\ \frac{1}{3}\left(\frac{m}{|S_a|} + \frac{m}{|S_b|} + \frac{m-t}{m}\right) \end{cases}$$

fit@hcmus

Text representation

- String based:
  - Edit Distance:
    - Levenshtein distance (L distance)
      - The minimum number of single-character edits (insertions, deletions or substitutions) required to change S1 into S2.
    - Damerau–Levenshtein Distance (D distance)
      - Like L distance, but with the addition of the transposition operation
    - Optimal String Alignment
      - Like D distance
      - No substring/subsequence is edited more than once.

fit@hcmus

Text representation

- Corpus based:
  - Use data from corpus
    - Textual feature
    - Co-occurrence probability
  - Includes 3 methods
    - Bag-of-words
    - Distributed representation
    - Matrix factorization

fit@hcmus

Text representation

It was the best of times,
it was the worst of times,
it was the age of wisdom,
it was the age of foolishnes

- “it”
- “was”
- “the”
- “best”
- “of”
- “times”
- “worst”
- “age”
- “wisdom”
- “foolishness”

- BOW:
  - Count the number of times each word appears

```
"it was the best of times"      = [1, 1, 1, 1, 1, 1, 0, 0, 0, 0]
"it was the worst of times"     = [1, 1, 1, 0, 1, 1, 1, 0, 0, 0]
"it was the age of wisdom"      = [1, 1, 1, 0, 1, 0, 0, 1, 1, 0]
"it was the age of foolishness" = [1, 1, 1, 0, 1, 0, 0, 1, 0, 1]
```

Text representation

- Term frequency–inverse document frequency (TF - IDF):
  - To measure how important a word is to a document in a collection (or corpus) of documents
    - TF: the ratio of a word's occurrence in a document

$$tf\ (w, d) = Freq\ (w, d)$$

    - IDF: indicates the amount of information provided

$$idf\ (w, D\ ) = log\frac{|D|}{N(w)}$$

$$tf\text{-}idf\ (w, d, D) = tf\ (w, d) \times idf\ (w, D)$$

fit@hcmus

Text representation

- word2vec:
  - Continuous Bag-of-words (BOW)
  - Word-skip grams (skip-gram)

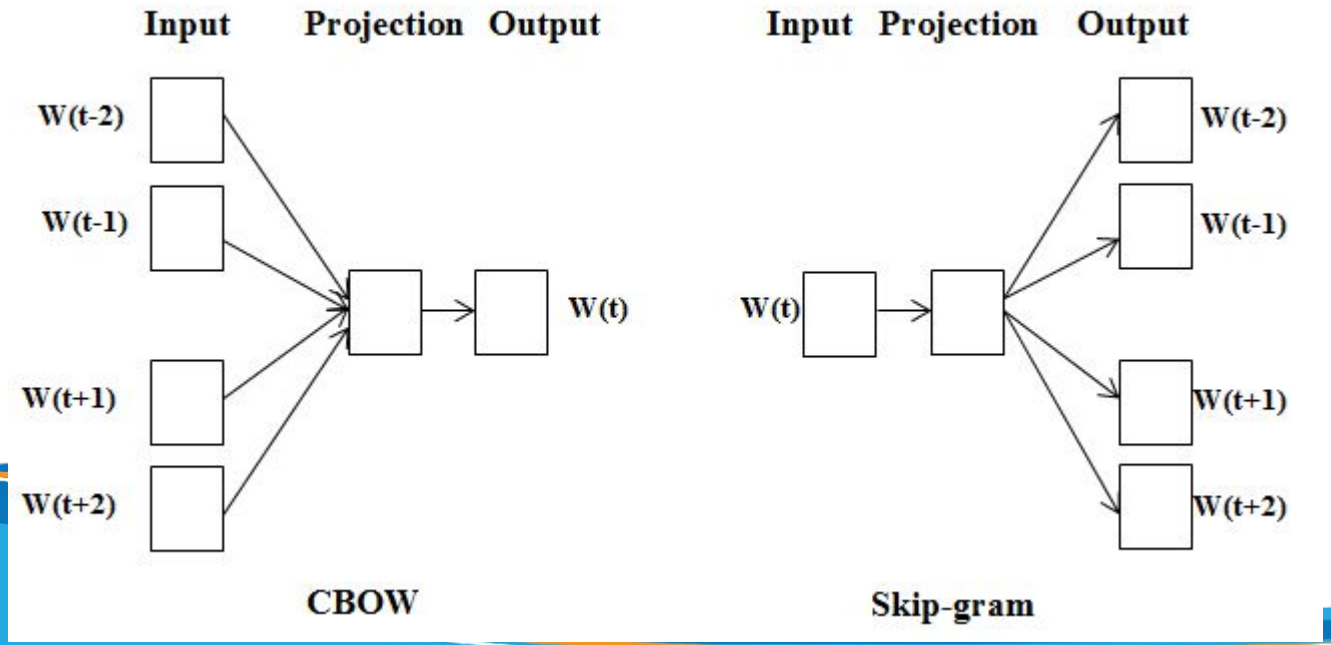

fit@hcmus

Text representation

- Glove:
  - Words with similar meanings tend to appear in similar contexts
  - Encode the ratios of co-occurrence probabilities with vector differences
- Bert: Bidirectional Encoder Representations from Transformers
  - Already pre-trained on massive datasets

fit@hcmus

Corpus

- STS Benchmark (STSb)
  - Includes 8,628 sentence pairs :
    - Train: 5,749
    - Develop: 1,500
    - Test: 1,379
  - Three categories : captions, news, and forums

fit@hcmus

Corpus

Dataset Name	Sentence pairs	Similarity score range	Year
LiSent	65	0 - 4	2007
SRS	30	0 - 4	2007
STS2012	5250	0 - 5	2012
STS2013	2250	0 - 5	2013
STS2014	3750	0 - 5	2014
SICK	10000	1 - 5	2014
STS2015	3000	0 - 5	2015
STS2016	1186	0 - 5	2016
STS2017	1750	0 - 5	2017

Applied Natural Language Processing- Text Similarity

EVALUATING

fit@hcmus

Similarity Evaluation

- Pearson correlation coefficient:

$$r = \frac{\sum_{i=1}^{n}(x_i - \bar{x})(y_i - \bar{y})}{\sqrt{\sum_{i=1}^{n}(x_i - \bar{x})^2 \sum_{i=1}^{n}(y_i - \bar{y})^2}}$$