


Artificial Intelligence

ADVERSARIAL SEARCH

Nguyễn Tiến Huy
ntienhuy@fit.hcmus.edu.vn

Outline

- The concept of games in AI
- Optimal decisions in games
- α - β Pruning
- Imperfect, real-time decisions
- Stochastic games



The concept of games in AI

Search in multiagent environments

- Each agent needs to **consider the actions of other agents** and how they affect its own welfare.
- The **unpredictability of other agents** introduce contingencies into the agent's problem-solving process



Game theory

- **Game theory** views any multiagent environment as a game.
 - The impact of each agent on the others is “significant,” regardless of whether the agents are cooperative or competitive.
- **Types of games**

	Deterministic	Chance
Perfect information	Chess, Checkers, Go, Othello	Backgammon, Monopoly
Imperfect information		Bridge, poker, scrabble nuclear war

Types of Games



Adversarial search

- **Adversarial search** (known as **games**) covers **competitive environments** in which the agents' goals are in conflict.
- **Zero-sum games of perfect information**
 - Deterministic, fully observable environments, turn-taking, two-player
 - The utility values at the end are always equal and opposite.

O		X
X	X	O
O		



Games vs. Search problems

- **Complexity**: games are too hard to be solved
 - Chess: $b \approx 35$, $d \approx 100$ (50 moves/player) \rightarrow graph of 10^{40} nodes, search tree of 35^{100} or 10^{154} nodes
 - Go: $b \approx 1000$ (!)
- **Time limits**: make some decision even when calculating the optimal decision is infeasible
- **Efficiency**: penalize inefficiency severely
 - Several interesting ideas on how to make the best possible use of time are spawn in game-playing research.

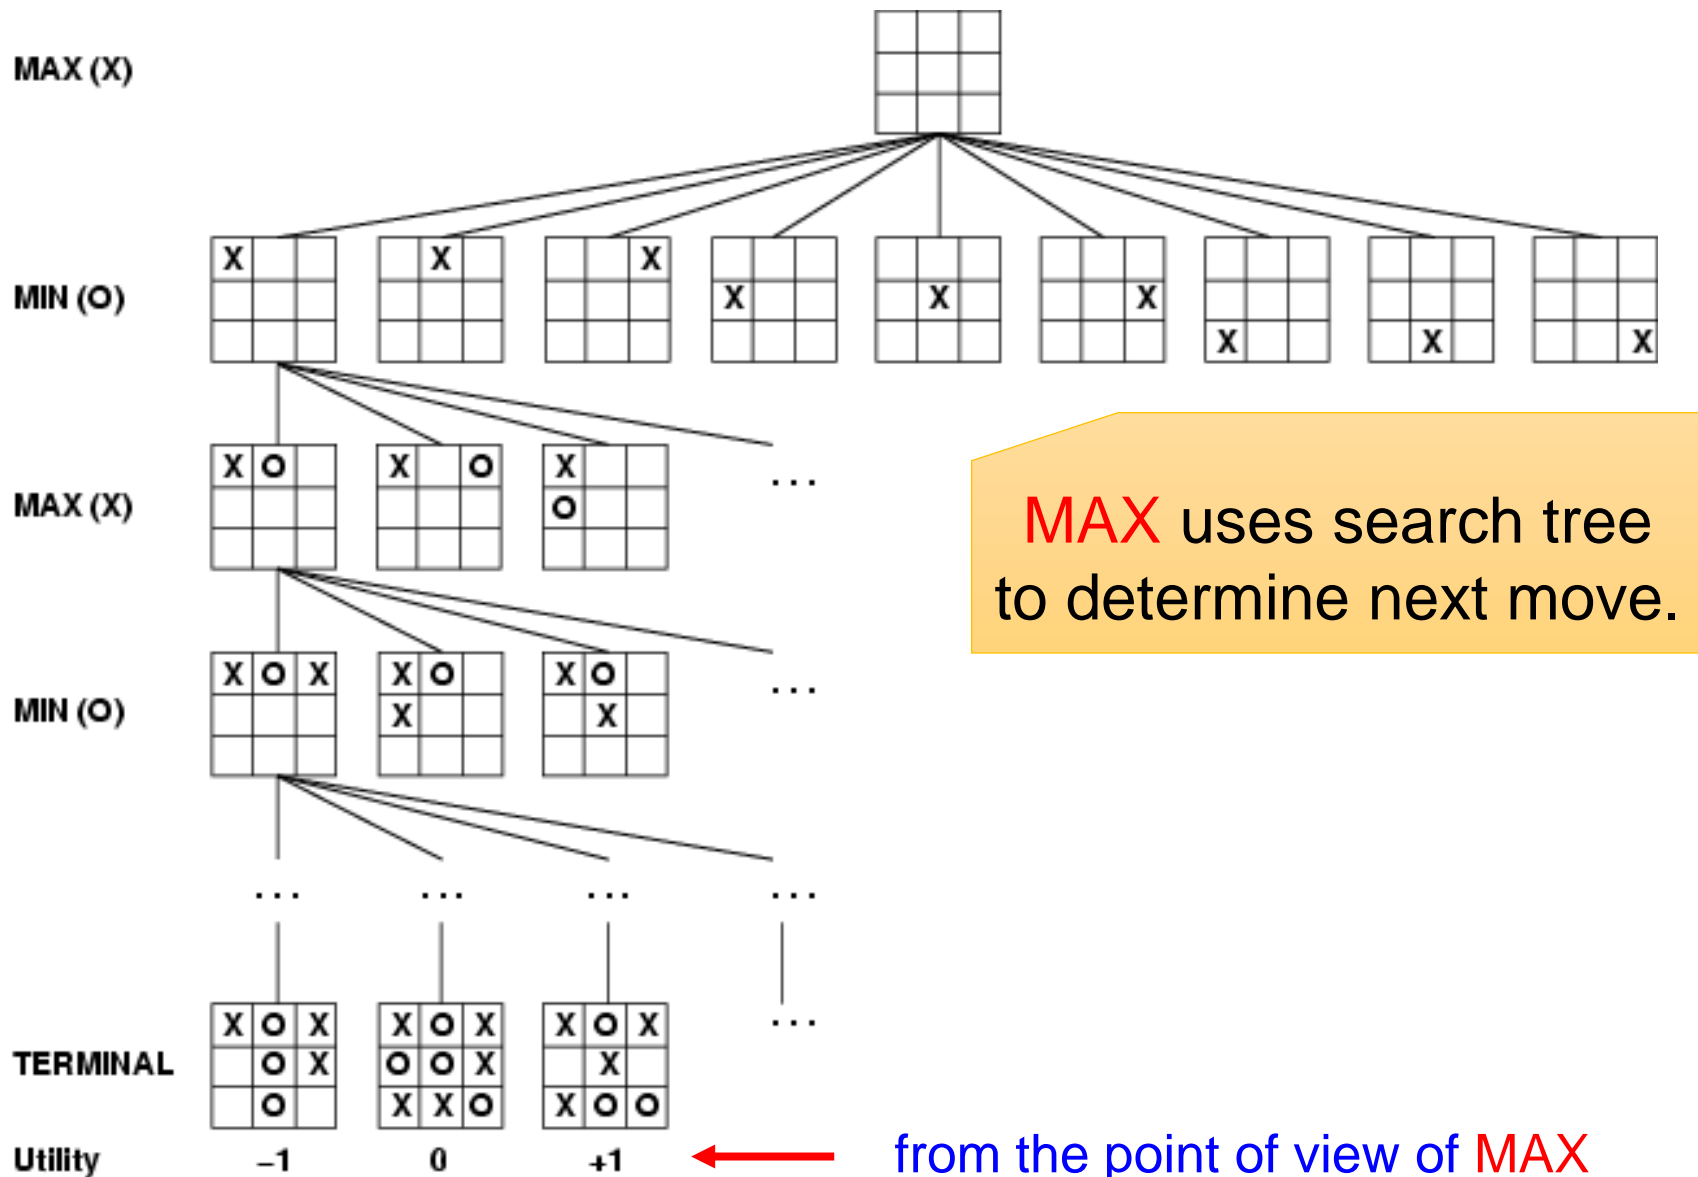
Primary assumptions

- Two players only, called MAX and MIN.
 - MAX moves first, and then they take turns moving until the game ends
 - Winner gets reward, loser gets penalty.
- Both players have complete knowledge of the game's state
 - E.g., chess, checkers and Go, etc. Counter examples: poker
- No element of chance
 - No dice thrown, no cards drawn, etc.
- Zero-sum games
 - The total payoff to all players is the same for every game instance.
- Rational players
 - Each player always tries to maximize his/her utility

Games as search

- S_0 – **Initial state**: How the game is set up at the start
 - E.g., board configuration of chess
- $PLAYER(s)$: Which player has the move in a state, **MAX/MIN?**
- $ACTIONS(s)$ – **Successor function**: A list of (move, state) pairs specifying legal moves.
- $RESULT(s, a)$ – **Transition model**: Result of move a on state s
- $TERMINAL - TEST(s)$: Is the game finished?
 - States where the game has ended are called **terminal states**
- $UTILITY(s, p)$ – **Utility function**: A numerical value of a terminal state s for a player p
 - E.g., chess: win (+1), lose (-1) and draw (0), backgammon: [0, 192]

The game tree of Tic-Tac-Toe



Examples of game: Checkers



- Complexity

- $\sim 10^{18}$ nodes, which may require 100k years with 106 positions/sec

- **Chinook** (1989-2007)

- The first computer program that won the world champion title in a competition against humans
 - 1990: won 2 games in competition with world champion Tinsley (final score: 2-4, 33 draws). 1994: 6 draws

- Chinook's search

- Ran on regular PCs, played perfectly by using [alpha-beta search](#) combining with [a database of 39 trillion endgame positions](#)

Examples of game: Chess

- Complexity

- $b \approx 35$, $d \approx 100$, 10^{154} nodes (!!)
- Completely impractical to search this

- **Deep Blue** (May 11, 1997)

- Kasparov lost a 6-game match against IBM's Deep Blue (1 win Kasp – 2 wins DB) and 3 ties.
- In the future, focus will be to allow computers to **LEARN** to play chess rather than being **TOLD** how it should play



Deep Blue

- Ran on a parallel computer with **30** IBM RS/6000 **processors** doing alpha–beta search
- Searched up to **30 billion positions**/move, average depth **14** (be able to reach to **40** plies)
- Evaluation function: **8000** features
 - highly specific patterns of pieces (~4000 positions)
 - 700,000 grandmaster games in database
- Working at **200 million positions**/sec, even Deep Blue would require **10^{100}** years to evaluate all possible games.
 - (The universe is only 10^{10} years old.)
- Now: algorithmic improvements have allowed programs running on standard PCs to win World Computer Chess Championships.
 - Pruning heuristics reduce the effective branching factor to less than 3




GO

1 million trillion trillion trillion
trillion more configurations
than chess!

- **Complexity**

- Board of 19x19, $b \approx 361$, average depth ≈ 200
- 10^{174} possible board configuration.
- Control of territory is unpredictable until the endgame




- **AlphaGo** (2016) by Google

- Beat 9-dan professional Lee Sedol (4-1)
- Machine learning + Monte Carlo search guided by a “value network” and a “policy network” (implemented using *deep neural network* technology)
- Learn from human + Learn by itself (self-play games)


An overview of AlphaGo

AlphaGo Overview

based on: Silver, D. et al. Nature Vol 529, 2016
copyright: Bob van den Hoek, 2016



Optimal decisions in games



- *Minimax algorithm*
- *Optimal decisions in multiplayer games*

Optimal decision in games

- Normal search problem

- The optimal solution is a sequence of action leading to a goal state.

- Games

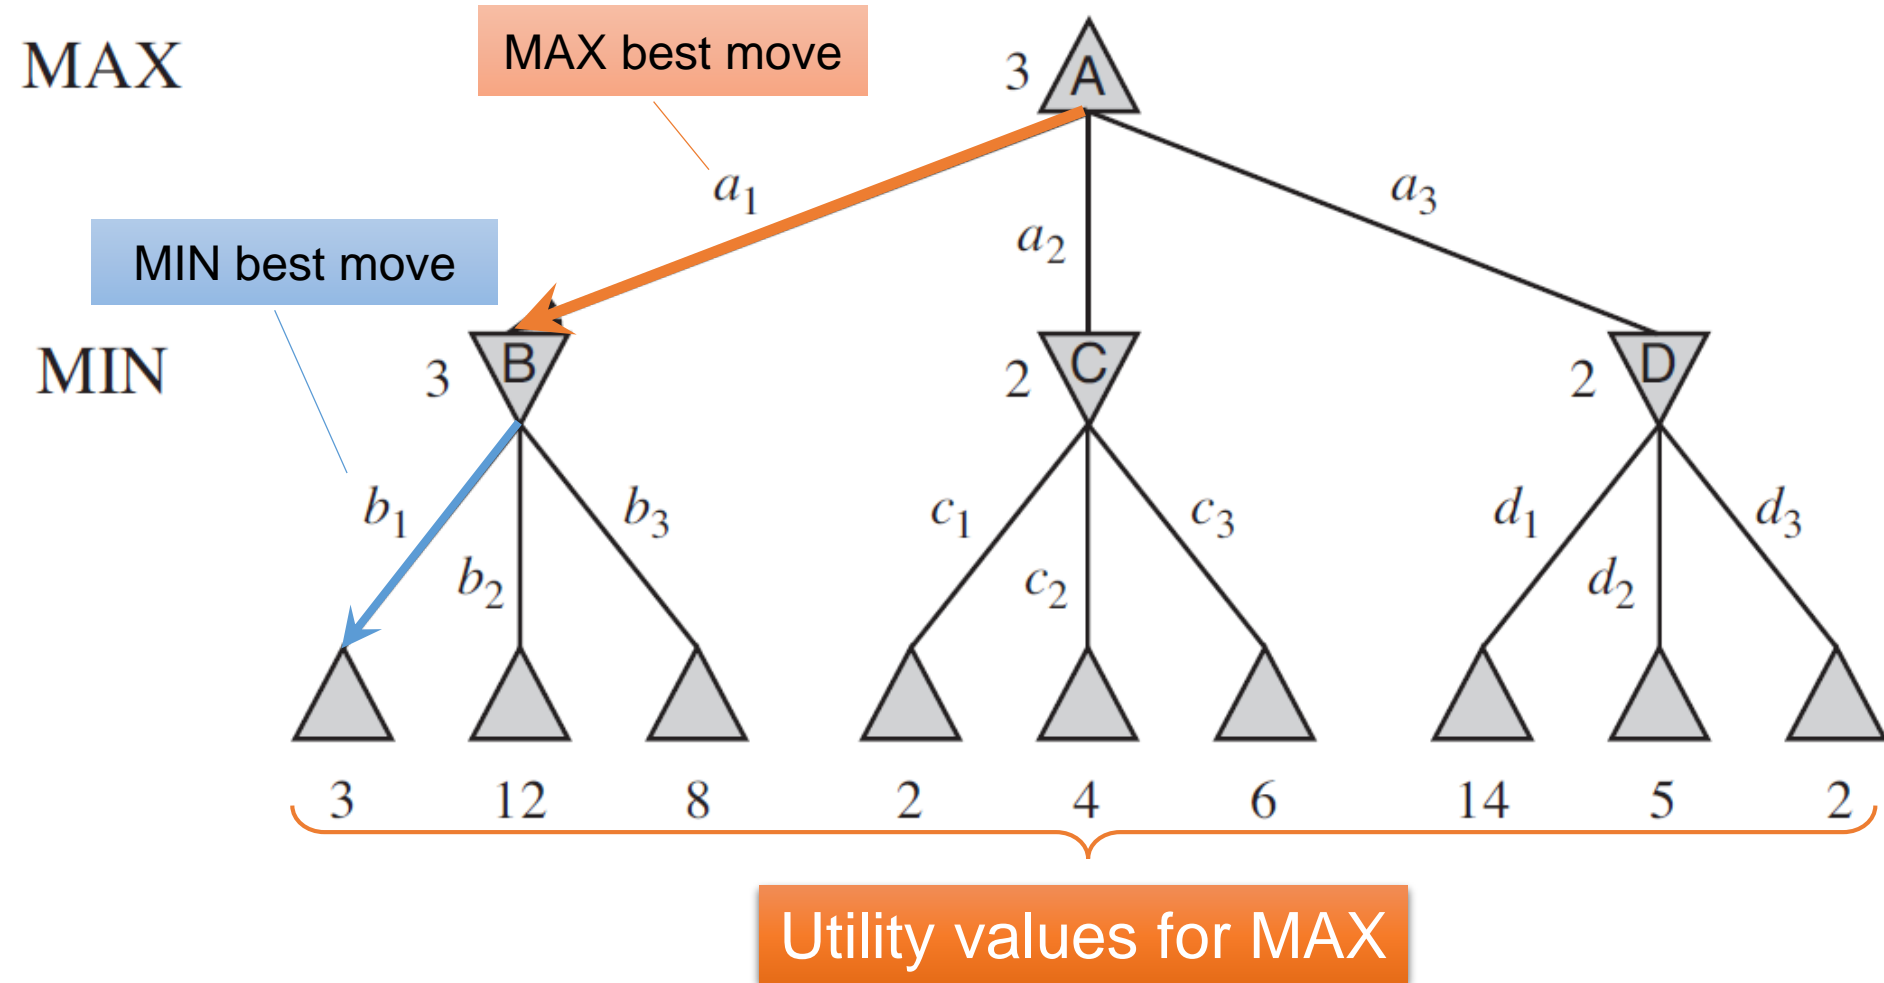
- The optimal strategy is a search path that guarantee win for a player
 - This can be determined from the **minimax value** of each node.

$$\text{MINIMAX}(s) = \begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

For MAX


Assume that both players play optimally from there to the end of the game

An example of two-ply game tree



Minimax algorithm

- Make a **minimax decision** from the current state, using a recursive computation of minimax values at each successor
 - The recursion proceeds all the way down to the leaves, and then back up the minimax values through the tree as it unwinds.



Minimax algorithm

function MINIMAX-DECISION(*state*) **returns** an action
return $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(state, a))$

function MAX-VALUE(*state*) **returns** a utility value
if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow -\infty$
for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$
return *v*

function MIN-VALUE(*state*) **returns** a utility value
if TERMINAL-TEST(*state*) **then return** UTILITY(*state*)
 $v \leftarrow \infty$
for each *a* **in** ACTIONS(*state*) **do**
 $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$
return *v*

Properties of Minimax algorithm

- A complete **depth-first exploration** of the game tree

- **Completeness**

- Yes (if tree is finite)

- **Optimality**

- Yes (against an optimal opponent)

Note:

m: the maximum depth of the tree

b: the legal moves at each point

- **Time complexity**

- $O(b^m)$

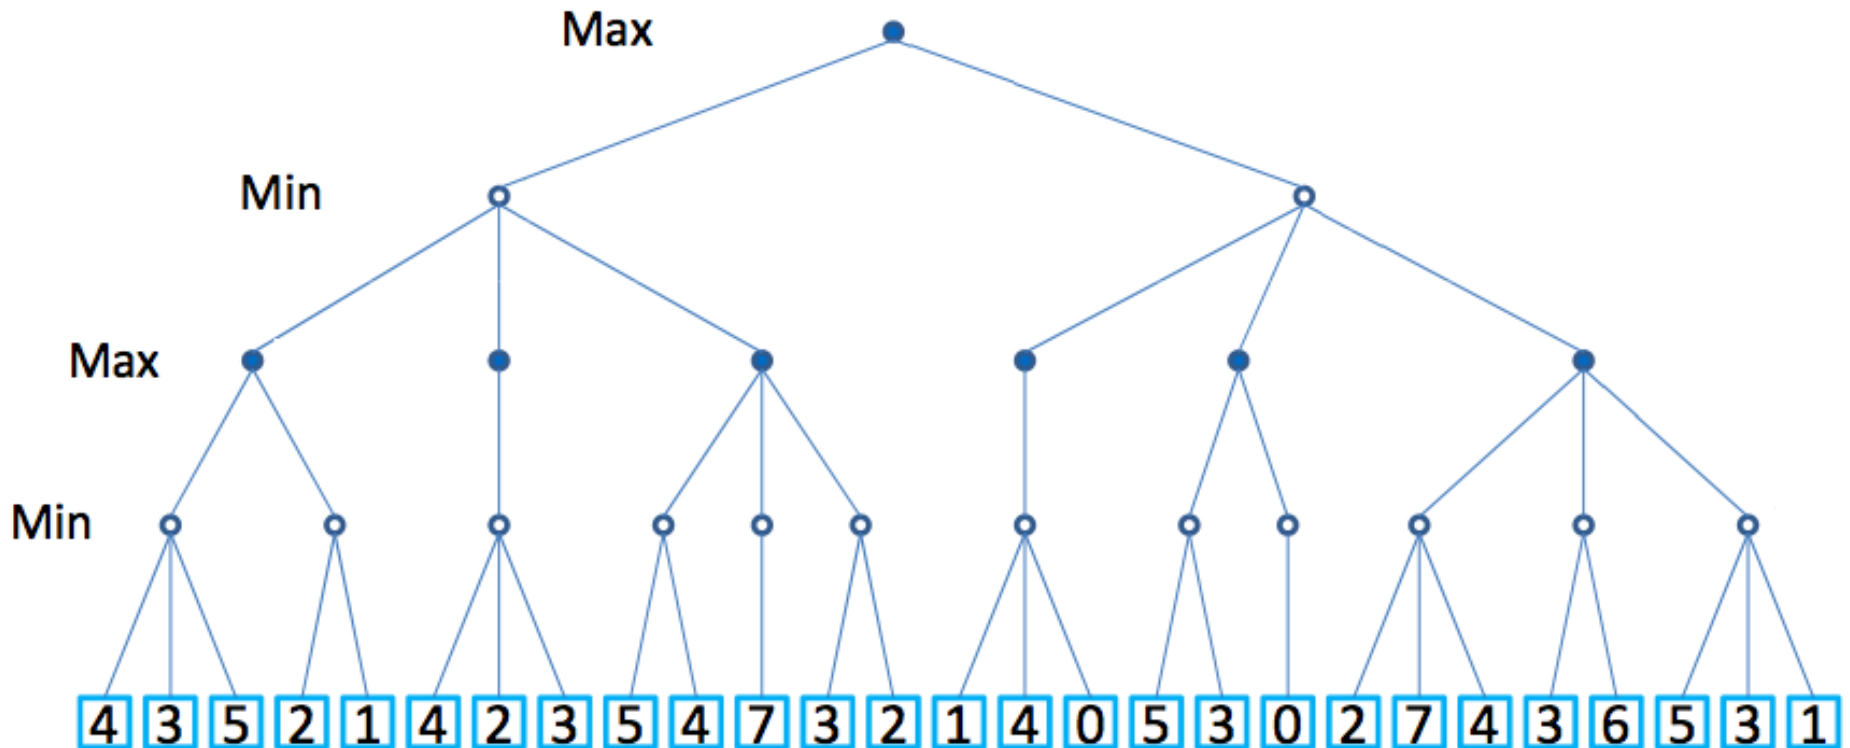
- **Space complexity**

- $O(bm)$ (depth-first exploration)

For chess, $b \approx 35, m \approx 100$ for "reasonable" games
→ exact solution completely infeasible

Quiz 01: Minimax algorithm

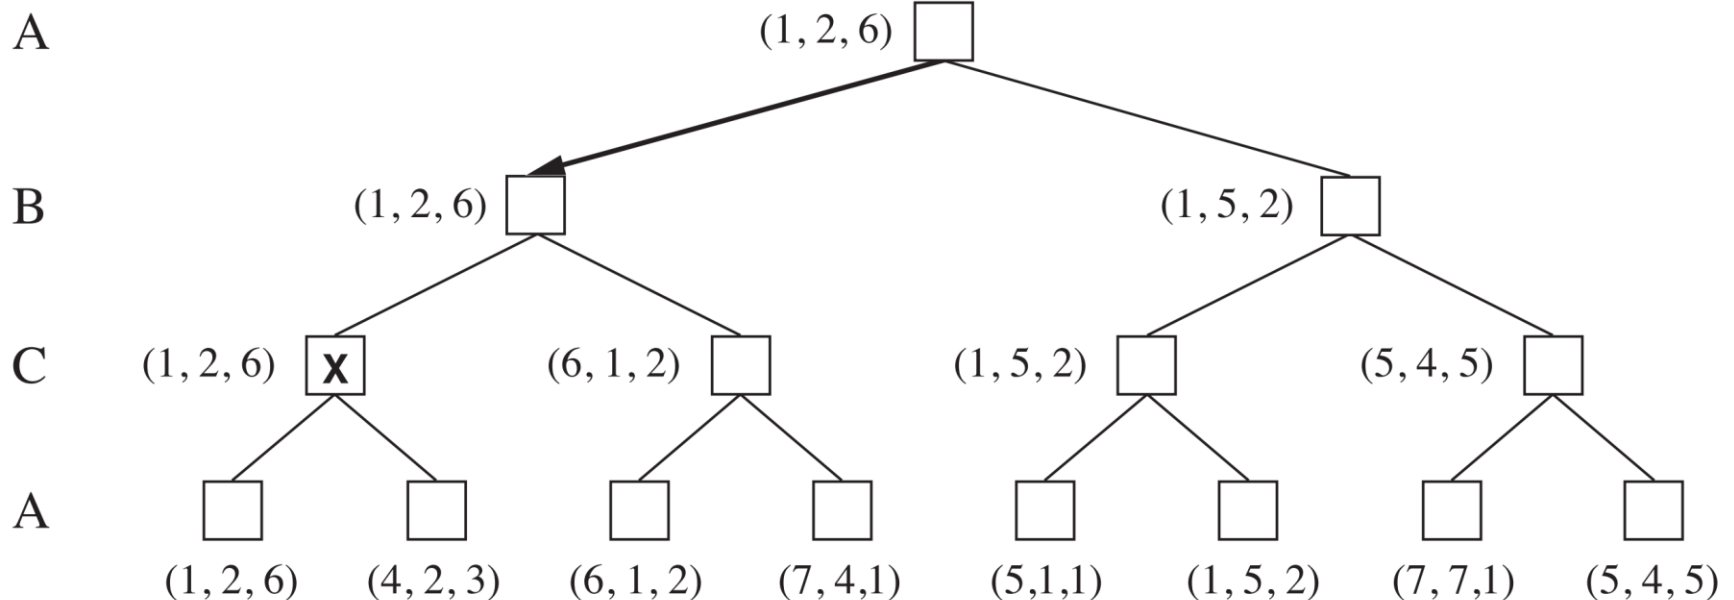
- Calculate the utility value for the remaining nodes
- Which node should MAX and MIN choose?



Optimality in multiplayer games


- A single value is replaced with a vector of values.
 - the UTILITY function **returns a vector of utilities**
- For terminal states, this vector gives the utility of the state from each player's viewpoint.

to move




Optimality in multiplayer games

- Multiplayer games usually involve **alliances**, which are made and broken as the game proceeds.



A and B are weak while C is strong.
A forms an alliance with B.



C becomes weak.
A or B could violate the agreement

- If the game is not zero-sum, then collaboration can also occur with just two players.












Alpha-beta pruning

Problem with minimax search

- The number of game states is **exponential** in the tree's depth
→ Do not examine every node
- **Alpha-beta pruning**: Prune away branches that cannot possibly influence the final decision
- **Bounded lookahead**
 - Limit depth for each search
 - This is what chess players do: look ahead for a few moves and see what looks best





Another way to look at this is as a simplification of the formula for MINIMAX.


Let the two unevaluated successors of node C have values x and y .

Then the value of the root node is given by

$$\begin{aligned}
 \text{MINIMAX}(\text{root}) &= \max(\min(3, 12, 8), \min(2, x, y), \min(14, 5, 2)) \\
 &= \max(3, \min(2, x, y), 2) \\
 &= \max(3, z, 2) \quad \text{where } z = \min(2, x, y) \leq 2 \\
 &= 3.
 \end{aligned}$$

Alpha-beta pruning

- If a move n is determined to be worse than move m that has already been examined and discarded, then examining move n once again is pointless.



α = the value of the **best** (i.e., highest-value) choice we have found so far at any choice point **along the path for MAX**.

β = the value of the **best** (i.e., lowest-value) choice we have found so far at any choice point **along the path for MIN**.

Alpha-beta search algorithm

```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in  $\text{ACTIONS}(\text{state})$  with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if  $\text{TERMINAL-TEST}(\text{state})$  then return  $\text{UTILITY}(\text{state})$   
   $v \leftarrow -\infty$   
  for each  $a$  in  $\text{ACTIONS}(\text{state})$  do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$   
     $\alpha \leftarrow \text{MAX}(\alpha, v)$   
return  $v$ 
```

Alpha-beta search algorithm


```
function MIN-VALUE( $state, \alpha, \beta$ ) returns a utility value
  if TERMINAL-TEST( $state$ ) then return UTILITY( $state$ )
   $v \leftarrow +\infty$ 
  for each  $a$  in ACTIONS( $state$ ) do
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$ 
    if  $v \leq \alpha$  then return  $v$ 
     $\beta \leftarrow \text{MIN}(\beta, v)$ 
  return  $v$ 
```

Properties of alpha-beta pruning

- Pruning **does not affect** the result
 - Its worst case is as good as the minimax algorithm
- **Good move ordering** improves effectiveness of pruning
 - With "perfect ordering": time complexity $O(b^{m/2}) \rightarrow \times 2$ search depth
 - The effective branching factor becomes \sqrt{b} instead of b .
 - E.g., for chess, about 6 instead of 35.

Quiz 02: Alpha-beta pruning

- Calculate the utility value for the remaining nodes.
- Which nodes should be pruning?



Imperfect real-time decisions



- *Evaluation functions*
- *Cutting off search*
- *Forward pruning*
- *Search versus Lookup*

Heuristic minimax

- Both minimax and alpha-beta pruning search all the way to terminal states.
 - This depth is usually **impractical** because moves must be made in a reasonable amount of time (~ minutes).
- Cut off the search earlier with some depth limit
- Use an evaluation function
 - An estimation for the desirability of position (win, lose, tie?)

$$\text{H-MINIMAX}(s, d) = \begin{cases} \text{EVAL}(s) & \text{if CUTOFF-TEST}(s, d) \\ \max_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{H-MINIMAX}(\text{RESULT}(s, a), d + 1) & \text{if PLAYER}(s) = \text{MIN}. \end{cases}$$

Evaluation functions

- These evaluation function should order the terminal states in the **same way as the true utility function** does
 - States that are wins must evaluate better than draws, which in turn must be better than losses.
- **The computation must not take too long!**
- For nonterminal states, their orders should be strongly **correlated with the actual chances of winning.**

Evaluation functions

- For chess, typically linear weighted sum of features

$$Eval(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

- where f_i could be the numbers of each kind of piece on the board, and w_i could be the values of the pieces
- E.g., $Eval(s) = 9q + 5r + 3b + 3n + p$
- Implicit strong assumption: the contribution of each feature is independent of the values of the other features.
 - E.g., assign the value 3 to a bishop ignores the fact that bishops are more powerful in the endgame → **Nonlinear combination**

Cutting off search


- *Minimax Cutoff* is identical to *Minimax Value* except
 1. *Terminal?* is replaced by *Cutoff?*
 2. *Utility* is replaced by *Eval*

if CUTOFF-TEST(state, depth) then return EVAL(state)

- Does it work in practice?
 - $b^m = 10^6, b = 35 \rightarrow m = 4$
 - 4-ply lookahead is a hopeless chess player!
 - 4-ply \approx human novice, 8-ply \approx typical PC, human master, 12-ply \approx Deep Blue, Kasparov

A more sophisticated cutoff test




- **Quiescent positions** are those unlikely to exhibit wild swings in value in the near future.
 - E.g., in chess, positions in which favorable captures can be made are not quiescent for an evaluation function counting material only
- **Quiescence search:** expand nonquiescent positions until quiescent positions are reached.




Stochastic games

Stochastic behaviors

- Uncertain outcomes controlled by chance, not an adversary!
- *Why wouldn't we know what the result of an action will be?*
 - Explicit randomness: rolling dice
 - Unpredictable opponents: the ghosts respond randomly
 - Actions can fail: when a robot is moving, wheels might slip

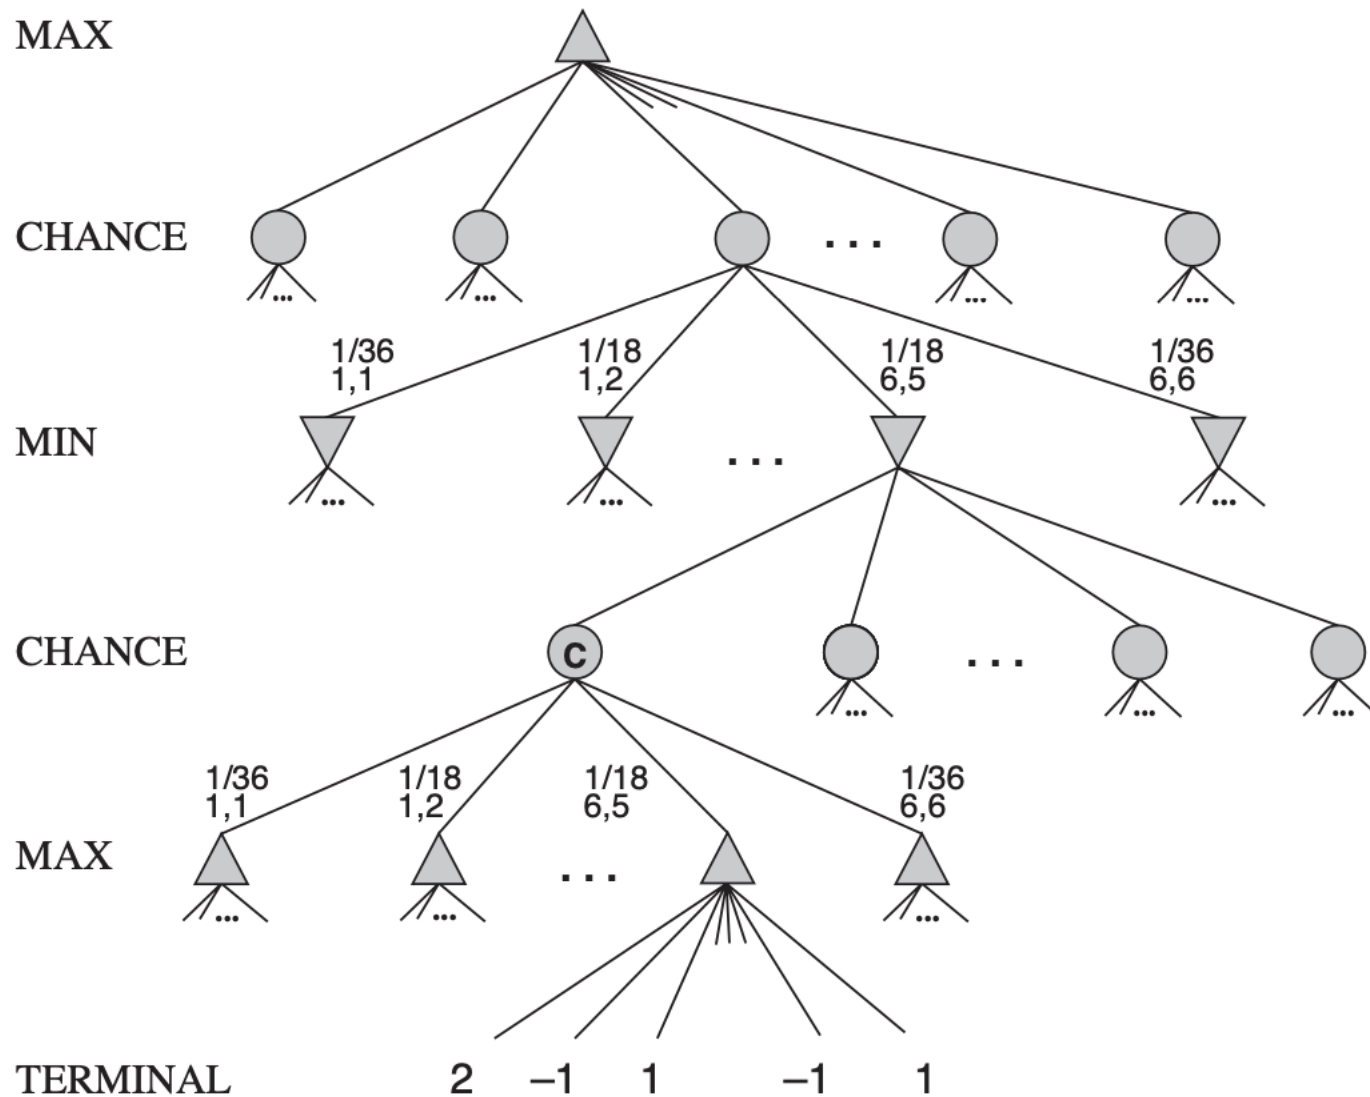


A stochastic game - Backgammon



A typical backgammon position. The goal of the game is to move all one's pieces off the board. White moves clockwise toward 25, and Black moves counterclockwise toward 0. A piece can move to any position unless multiple opponent pieces are there; if there is one opponent, it is captured and must start over. In the position shown, White has rolled 6–5 and must choose among four legal moves: (5–10,5–11), (5–11,19–24), (5–10,10–16), and (5–11,11–16), where the notation (5–11,11–16) means move one piece from position 5 to 11, and then move a piece from 11 to 16.

A game tree for a backgammon position



Expectiminimax

EXPECTIMINIMAX(s) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_a \text{EXPECTIMINIMAX}(\text{RESULT}(s, a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \\ \sum_r P(r) \text{EXPECTIMINIMAX}(\text{RESULT}(s, r)) & \text{if } \text{PLAYER}(s) = \text{CHANCE} \end{cases}$$



THE END