

Balanced Binary Search Tree

Bùi Tiến Lên

2022



KHOA CÔNG NGHỆ THÔNG TIN
TRƯỜNG ĐẠI HỌC KHOA HỌC TỰ NHIÊN

Contents



1. **Balanced Tree**

2. **AVL Tree**

3. **Red-Black Tree**

4. **Optimal binary search trees**

5. **Workshop**



Balanced Tree

- Rotations
- Strategies in Balancing Tree

Introduction



Balance may be defined by:

- Comparing the numbers of nodes of the two subtrees
- Height balancing: comparing the heights of the two sub trees
- Null-path-length balancing: comparing the null-path-length of each of the two sub-trees
- Weight balancing: comparing the number of null sub-trees in each of the two sub trees



Introduction (cont.)

Concept 1

A binary tree is **balanced** if the difference in the numbers of nodes of both subtrees of any node in the tree either **zero** or **one**.

Concept 2

A binary tree is **height-balanced** if the difference in height of both subtrees of any node in the tree either **zero** or **one**.

- A complete binary tree is is height-balanced

Rotations



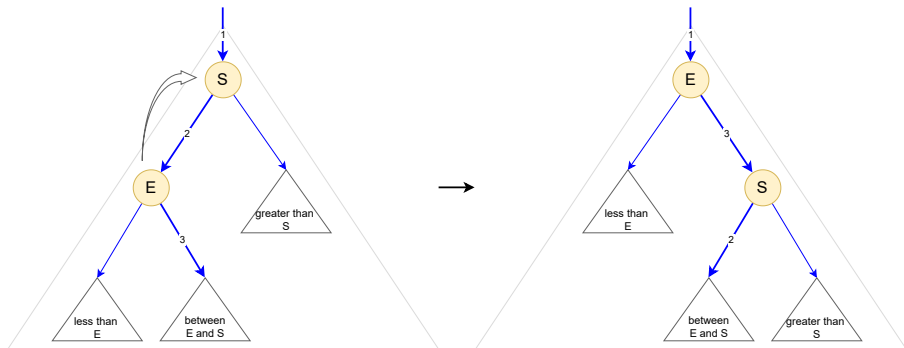
- A **rotation** allows us to interchange the role of the root and one of the root's children in a tree while still **preserving the BST ordering** among the keys in the nodes.
- There are two kinds of rotations: right rotation and left rotation




Right rotation

A right rotation involves the root and the left child. The rotation puts the root on the right, essentially reversing the direction of the left link of the root:


- Before the rotation, it points from the root to the left child
- After the rotation, it points from the old left child (the new root) to the old root (the right child of the new root)



Example



- Make right rotation at 15



Implementation

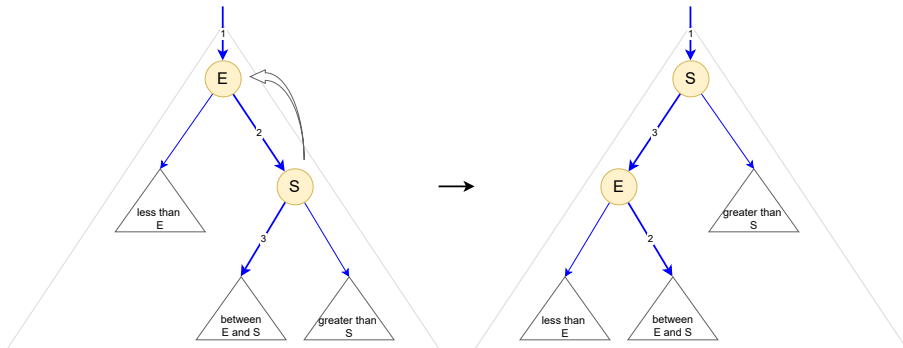


```
void rightRotate(link& h) {  
    link x = h->left;  
    h->left = x->right;  
    x->right = h;  
    h = x;  
}
```



Left rotation

A left rotation involves the root and the right child.



Implementation



```
void leftRotate(link& h) {  
    link x = h->right;  
    h->right = x->left;  
    x->left = h;  
    h = x;  
}
```

Strategies in Balancing Tree



- **Global rebalancing:** an approach to producing better balance in BSTs is periodically to rebalance them explicitly.
 - costs at least linear time in the size of the tree
- **Local rebalancing:** balancing BSTs after each operation (insert, delete)

DSW algorithm



The algorithm was proposed by Colin Day and later improved by Quentin F. Stout and Bette L. Warren.

Idea of algorithm:

1. Transform an arbitrary BST into a *linked-list-like-tree* called **backbone** or **vine** by rotations
2. Transform this tree into a **perfectly balanced tree** by rotations

DSW algorithm (cont.)



```
CREATEBACKBONE(root)
```

```
  p := root
```

```
  while p ≠ null?
```

```
    if p has a left child?
```

```
      make right rotation at p
```

```
    else
```

```
      p := p → right
```



DSW algorithm (cont.)

```
CREATECOMPLETETREE(root)
```

```
   $n \leftarrow$  the number of nodes
```

```
   $m \leftarrow 2^{\lfloor \log_2(n+1) \rfloor} - 1$ 
```

```
  make  $n - m$  left rotations starting from the top of backbone
```

```
  while ( $m > 1$ )
```

```
     $m \leftarrow m/2$ 
```

```
    make  $m$  left rotations starting from the top of backbone
```

Illustration



Figure 1: BST




Figure 2: Backbone





Figure 3: Perfect





AVL Tree

- Insertion
- Deletion



AVL Tree

Concept 3

AVL tree

- proposed by two Soviet scientists G. M. Adelson-Velskii and E. M. Landis
- is BST tree which is **height-balanced**

$$\forall p : |\mathbf{height}(\mathbf{LeftSubtree}(p)) - \mathbf{height}(\mathbf{RightSubtree}(p))| \leq 1 \quad (1)$$



The Height of an AVL Tree

Consider the worst case,

- To determine the maximum height that an AVL tree with N nodes can have, we can instead ask what is the minimum number of nodes that an AVL tree of height h can have (called AVL tree F_h).
- We have the recurrence relation

$$|F_h| = |F_{h-1}| + |F_{h-2}| + 1 \quad (2)$$

where $|F_0| = 1$ and $|F_1| = 2$

- Solve the equation, we have

$$|F_h| + 1 \approx \frac{1}{\sqrt{5}} \left[\frac{1 + \sqrt{5}}{2} \right]^{h+2} \quad (3)$$





- The height of an AVL tree in the worst case

$$h \approx 1.44 \log_2 |F_h| = 1.44 \log_2 N \quad (4)$$



Rebalancing technique

- After an insertion/deletion, we may find a node whose new balance violates the AVL condition.
- Four cases: LL imbalance, LR imbalance, RR imbalance, RL imbalance





Rebalancing technique (cont.)

Balanced Tree
Rotations
Strategies in Balancing Tree

AVL Tree

Insertion
Deletion

Red-Black Tree


Insertion
Deletion

Optimal binary search trees

Static
Dynamic
Splaying

Workshop


- Case **LL imbalance** is corrected by executing a **single right rotation** at the node with the imbalance.



Rebalancing technique (cont.)



- Case **LR imbalance** is corrected by executing a **double rotations**





Insertion

```
INSERT(root,key)  
  if root = null  
    root := new NODE(key)  
    return  
  if root → key = key  
    return  
  if root → key < key  
    INSERT(root → right,key)  
  if root → key > key  
    INSERT(root → left,key)  
  if unbalanced at root? rebalance at root
```



Illustration

- An AVL tree






Illustration (cont.)

- Insert node 54 into the tree → node 78 become unbalanced







Illustration (cont.)

- Case RL imbalance \rightarrow rebalance by making double rotations





Illustration

- An AVL tree






Illustration (cont.)

- Delete node 32 from the tree → node 44 become unbalanced







Illustration (cont.)

- Case RL imbalance \rightarrow rebalance by making double rotations





Red-Black Tree

- Insertion
- Deletion



Red-Black Tree

Concept 4

A red-black (**RB**) tree is a special type of binary search tree that must satisfy

1. Each node is either **red** or **black**.
2. The root is **black**. (sometimes omitted)
3. If a node is **red**, then both its children are **black**.
4. Every path from a given node to any of its descendant null link has the same number of **black** nodes (balance criteria).


Concept 5

A left-leaning red-black (**LLRB**) tree (leveraging Andersson's idea AA tree) is a variant of the red-black tree that has only left red children



Example


- A red-black tree





Example (cont.)


- A left-leaning red-black tree



Example (cont.)



- AA tree





The Height of a RB Tree

Theorem 1


The height of a red-black BST with N nodes is no more than $2\log_2 N$. It means that the height of an RB tree in the worst case

$$h \leq 2\log_2 N \quad (5)$$



Operations

- Case 1: Left rotation to orient a right red node to left red node.





Balanced Tree

Rotations

Strategies in Balancing Tree

AVL Tree

Insertion

Deletion

Red-Black Tree

Insertion

Deletion

Optimal binary search trees

Static


Dynamic

Splaying

Workshop

Operations (cont.)

- Case 2: Right rotation.





Balanced Tree

Rotations

Strategies in Balancing Tree

AVL Tree

Insertion

Deletion

Red-Black Tree

Insertion

Deletion

Optimal binary search trees

Static


Dynamic

Splaying

Workshop

Operations (cont.)

- Case 3: Color flip.



Insertion




```
INSERT(root,key)  
  if root = null  
    root := new NODE(key) red node  
    return  
  if root → key = key  
    return  
  if root → key < key  
    INSERT(root → right,key)  
  if root → key > key  
    INSERT(root → left,key)  
  if ISRED(root → right) and not ISRED(root → left) ROTATELEFT(root)  
  if ISRED(root → left) and ISRED(root → left → left) ROTATERIGHT(root)  
  if ISRED(root → left) and ISRED(root → right) FLIPCOLORS(root)
```



Example

- Typical left-leaning red-black BST built from random keys





Illustration

- A left-leaning red-black tree






Illustration (cont.)

- Insert node 7 into the tree






Illustration (cont.)

- Flip color at 6 into the tree






Illustration (cont.)

- Right rotation at 13 into the tree







Illustration (cont.)

- Flip color at root at change root to black color





Cost summary for symbol-table implementations

Balanced Tree
Rotations
Strategies in Balancing Tree

AVL Tree

Insertion
Deletion

Red-Black Tree

Insertion
Deletion

Optimal binary search trees

Static
Dynamic
Splaying

Workshop

implementation	worst case			average case			key
	search	insert	remove	search hit	insert	remove	
unordered list	N	1	N	$N/2$	1	$N/2$	equal
ordered list	N	N	N	$N/2$	$N/2$	$N/2$	compare
ordered array	$\log_2 N$	N	N	$\log_2 N$	$N/2$	$N/2$	compare
BST	N	N	N	$c \log_2 N$	$c \log_2 N$	\sqrt{N}	compare
AVL	$c_a \log_2 N$	-	-	$\log_2 N$	-	-	compare
RB	$c_r \log_2 N$	-	-	$\log_2 N$	-	-	compare
goal?							

Note: $c = 1.39$, $c_a = 1.44$, $c_r = 2.0$



Optimal binary search trees

- Static
- Dynamic



Introduction

Concept 6

An **optimal binary search tree** (optimal BST), sometimes called a weight-balanced binary tree, is a binary search tree which provides the smallest possible search time (or expected search time) for a given sequence of accesses (or access probabilities)

- Optimal BSTs are generally divided into two types: **static** and **dynamic**
 - In the static optimality problem, the tree cannot be modified after it has been constructed.
 - In the dynamic optimality problem, the tree can be modified at any time, typically by permitting tree rotations.



Introduction

- Suppose that we are designing a binary search tree for a program to translate text from English to Vietnamese, we want words that occur frequently in the text to be placed nearer the root.

Problem

Given a sequence of of n distinct keys in sorted order ($k_1 < k_2 < \dots < k_n$) and their frequencies

$$\mathcal{D} =$$

key	k_1	k_2	\dots	k_n
frequency	f_1	f_2		f_n

What binary search tree has the **lowest search cost**?



Search Cost

- Cost of search for key k_i

$$\text{COST}(k_i) = \text{DEPTH}(k_i) \quad (6)$$

where $\text{DEPTH}(\text{root}) = 1$

- Denote $\text{EXPECTCOST}(l, r)$ be expected cost of search for a BST tree containing $\{k_l, \dots, k_r\}$ given \mathcal{D}

$$\text{EXPECTCOST}(l, r) = \sum_{i=l}^r \text{COST}(k_i) f_i \quad (7)$$




Search Cost (cont.)

- Compute the expected cost for the following binary search tree given

$$\mathcal{D} =$$

key	k_1	k_2	k_3	k_4	k_5	k_6
frequency	0.1	0.2	0.1	0.3	0.2	0.1





Optimal Search Cost

Balanced Tree
Rotations
Strategies in Balancing Tree

AVL Tree

Insertion
Deletion

Red-Black Tree

Insertion
Deletion

Optimal binary search trees

Static
Dynamic
Splaying

Workshop

- Denote $\text{OPTIMALCOST}(l, r)$ be optimal cost of search for a BST tree T containing $\{k_l, \dots, k_r\}$ given \mathcal{D}
- Denote $\text{OPTIMALCOST}(l, m, r)$ be optimal cost of search for a BST tree T containing $\{k_l, \dots, k_r\}$ and k_m be the root node given \mathcal{D}
- We have

$$\begin{aligned}\text{OPTIMALCOST}(l, m, r) &= \sum_{i=l}^r f_i \\ &\quad + \text{OPTIMALCOST}(l, m-1) \\ &\quad + \text{OPTIMALCOST}(m+1, r)\end{aligned}\tag{8}$$

$$\text{OPTIMALCOST}(l, r) = \min_{m \in \{l, \dots, r\}} \{\text{OPTIMALCOST}(l, m, r)\}\tag{9}$$



Optimal Search Cost (cont.)

Balanced Tree

Rotations

Strategies in Balancing Tree

AVL Tree

Insertion

Deletion

Red-Black Tree

Insertion

Deletion


Optimal binary search trees

Static

Dynamic

Splaying

Workshop





Splay tree

Concept 7

A **splay tree** is a binary search tree with the additional property that recently accessed elements are quick to access again.

- All normal operations (*insert*, *look-up*) on a binary search tree are combined with one basic operation, called **splaying**.
- For many sequences of non-random operations, splay trees perform better than other binary search trees.

Splaying



- When a node x is accessed, a splay operation is performed on x to move it to the root.
- To perform a splay operation we carry out a sequence of *splay steps*, each of which moves x closer to the root.
- There are three types of splay steps, each of which has two symmetric variants:
 - zig step
 - zig-zig step
 - zig-zag step



Balanced Tree

Rotations

Strategies in Balancing Tree

AVL Tree

Insertion

Deletion

Red-Black Tree

Insertion

Deletion

Optimal binary search trees


Static

Dynamic

Splaying

Workshop

Zig step





Balanced Tree

Rotations

Strategies in Balancing Tree

AVL Tree

Insertion

Deletion

Red-Black Tree

Insertion

Deletion

Optimal binary search trees

Static

Dynamic

Splaying

Workshop

Zig-zig





Balanced Tree

Rotations

Strategies in Balancing Tree

AVL Tree

Insertion

Deletion

Red-Black Tree

Insertion

Deletion

Optimal binary search trees


Static

Dynamic

Splaying

Workshop

Zig-zag



Workshop



Balanced Tree

Rotations

Strategies in Balancing Tree

AVL Tree

Insertion

Deletion

Red-Black Tree

Insertion

Deletion


Optimal binary search trees


Static

Dynamic

Splaying

Workshop

 Quiz



1. What is an AVL tree?

2. What is a Red-black tree?

60



Exercises



- Programming exercises in [[Cormen, 2009](#), [Sedgewick, 2002](#)]

References



Cormen, T. H. (2009).
Introduction to algorithms.
MIT press.



Sedgewick, R. (2002).
Algorithms in Java, Parts 1-4, volume 1.
Addison-Wesley Professional.



Walls and Mirrors (2014).
Data Abstraction And Problem Solving with C++.
Pearson.