

Priority Queues and Heaps

Bùi Tiến Lên

2021



KHOA CÔNG NGHỆ THÔNG TIN
TRƯỜNG ĐẠI HỌC KHOA HỌC TỰ NHIÊN

Contents



1. **Binary Heaps**

2. **d-Heaps**

3. **Application**

4. **Workshop**



Priority queue applications

d-Heaps

Application

Simulation
Time-driven Simulation
Event-driven
Simulation

Workshop

Event-driven simulation.	(customers in a line, colliding particles)
Numerical computation.	(reducing roundoff error)
Data compression.	(Huffman codes)
Graph searching.	(Dijkstra's algorithm, Prim's algorithm)
Number theory.	(sum of powers)
Artificial intelligence.	(A* search)
Statistics.	(maintain largest M values in a sequence)
Operating systems.	(load balancing, interrupt handling)
Discrete optimization.	(bin packing, scheduling)
Spam filtering.	(Bayesian spam filter)





Priority queue

Concept 1

A **priority queue** is a data structure for maintaining a set S of elements, each with an associated value called a **key**. There are two kinds of priority queues: **max-priority queues** and **min-priority queues**.

Max-priority queue supports the following operations.

- $\text{INSERT}(S, x)$: insert the element x into the set S or $S \leftarrow S \cup \{x\}$.
- $\text{MAX}(S)$: return the element of S with the largest key.
- $\text{REMOVE-MAX}(S)$: remove and return the element of S with the largest key.
- $\text{INCREASE-KEY}(S, x, \Delta k)$: increase the value of element x 's key by the Δk .

Unordered and ordered array implementation



d-Heaps

Application

Simulation
Time-driven Simulation
Event-driven
Simulation

Workshop

- A sequence of operations on a max-priority queue

operation	argument	return value	size	contents (unordered)					contents (ordered)							
insert	P		1	P					P							
insert	Q		2	P	Q				P	Q						
insert	E		3	P	Q	E			E	P	Q					
remove max		Q	2	P	E				E	P						
insert	X		3	P	E	X			E	P	X					
insert	A		4	P	E	X	A		A	E	P	X				
insert	M		5	P	E	X	A	M	A	E	M	P	X			
remove max		X	4	P	E	M	A		A	E	M	P				
insert	P		5	P	E	M	A	P	A	E	M	P	P			
insert	L		6	P	E	M	A	P	L	A	E	L	M	P	P	
insert	E		7	P	E	M	A	P	L	E	A	E	L	M	P	P
remove max		P	6	E	M	A	P	L	E	A	E	E	L	M	P	

Analysis



Challenge Implement **all** operations efficiently.

implementation	insert	remove max	max
unordered array	1	N	N
ordered array	N	1	1
goal			



Binary Heaps

- Insert Key
- Remove Max
- Increase Key



Introduction

Concept 2

- **Max heap:** A tree is heap-ordered if the key in each node is larger than or equal to the keys in all of that node's children (if any)
- **Min heap:** A tree is heap-ordered if the key in each node is smaller than or equal to the keys in all of that node's children (if any)

Theorem 1

- *Max heap: No node in a heap-ordered tree has a key larger than the key at the root*
- *Min heap: No node in a heap-ordered tree has a key smaller than the key at the root*




Binary heap representations

Concept 3

A **binary heap** data structure is a complete binary tree that can be represented by an array object.

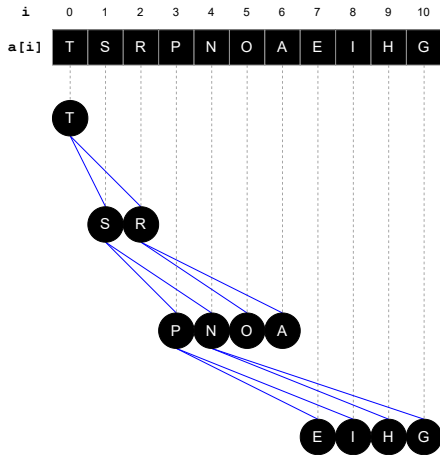

- Complete binary tree is perfectly balanced, except for bottom level.
- Height of complete binary tree with N nodes is $\log_2(N + 1)$.



Binary heap representations (cont.)



- Array representation of a max binary heap.





Binary heap properties

d-Heaps

Application

Workshop

- Largest key is $a[0]$, which is root of max binary heap.
- Can use array indices to move through tree.

```
PARENT( $i$ )  
  return  $\lfloor \frac{i-1}{2} \rfloor$ 
```

```
LEFTCHILD( $i$ )  
  return  $2i + 1$ 
```

```
RIGHTCHILD( $i$ )  
  return  $2i + 2$ 
```



Promotion in a heap

- **Scenario.** Child's key becomes larger key than its parent's key. (violation)
- To eliminate the violation:
 - Exchange key in child with key in parent.
 - Repeat until heap order restored.

```
UP-HEAP( $a, i$ )  
  while  $i > 0$   
     $p \leftarrow \text{PARENT}(i)$   
    if  $a[i] > a[p]$   
      SWAP( $a, i, p$ )  
       $i \leftarrow p$   
    else  
      return
```



Insertion in a heap

Insert. Add node at end, then percolate it up.

Cost. At most $\log_2 N + 1$ compares.

```
INSERT(a, k)  
  n ← a.size  
  a[n] ← k  
  UP-HEAP(a, n)  
  a.size ← a.size + 1
```



Example


d-Heaps

Application

Simulation
Time-driven Simulation
Event-driven
Simulation

Workshop

- Given a max binary heap




0	1	2	3	4	5	6	7	8	9	10	11
68	65	32	31	26	24	21	20	19	13		



Example (cont.)


- Insert key **66** into the heap





Example (cont.)


- Swap 66 and 26





Example (cont.)

- Swap 66 and 65





Demotion in a heap

- **Scenario.** Parent's key becomes smaller than one (or both) of its children's. (violation)
- To eliminate the violation:
 - Exchange key in parent with key in larger child.
 - Repeat until heap order restored.

DOWN-HEAP(a, i)

$l \leftarrow \text{LEFTCHILD}(i)$

$r \leftarrow \text{RIGHTCHILD}(i)$

$largest \leftarrow i$

if $l < a.size$ **and** $a[l] > a[largest]$ **then** $largest \leftarrow l$

if $r < a.size$ **and** $a[r] > a[largest]$ **then** $largest \leftarrow r$

if $largest \neq i$ **then**

 SWAP($a, i, largest$)

 DOWN-HEAP($a, largest$)



Delete the maximum in a heap

- **Delete max.** Exchange root with node at end, then sink it down.
- **Cost.** At most $2 \log_2 N$ compares.

REMOVE-MAX(a)

$n \leftarrow a.size$

SWAP(a , 0, $n - 1$)

DOWN-HEAP(a , 0)

$a.size \leftarrow a.size - 1$

return $a[n - 1]$ and delete $a[n - 1]$



Example

d-Heaps

Application

Simulation


Time-driven Simulation

Event-driven

Simulation

Workshop

- Given a max binary heap




0	1	2	3	4	5	6	7	8	9	10	11
68	65	32	31	26	24	21	20	19	13		



Example (cont.)

- Delete max : swap 68 and 13, delete 68




0	1	2	3	4	5	6	7	8	9	10	11
13	65	32	31	26	24	21	20	19	68		



Example (cont.)

- Swap 13 and 65




0	1	2	3	4	5	6	7	8	9	10	11
65	13	32	31	26	24	21	20	19	68		



Example (cont.)

- Swap 13 and 31




0	1	2	3	4	5	6	7	8	9	10	11
65	31	32	13	26	24	21	20	19	68		



Example (cont.)

- Swap 13 and 20



0	1	2	3	4	5	6	7	8	9	10	11
65	31	32	20	26	24	21	13	19	68		



Binary Heaps

Insert Key

Remove Max

Increase Key

Analysis

d-Heaps

Application

Simulation

Time-driven Simulation

Event-driven

Simulation

Workshop

implementation	insert	remove max	max
unordered array	1	N	N
ordered array	N	1	1
binary heap	$\log_2 N$	$\log_2 N$	1



Increase Key

- To increase the value of a certain key inside the max-heap, we need to reach this key first. In ordinary heaps, we can't search for a specific key inside the heap.
- Therefore, we'll keep a **map (hash table)** beside the original array. This map will store the index of each key inside the heap.

```
INCREASE-KEY( $a, k, \Delta k, map$ )
```

```
   $i \leftarrow map[k]$ 
```

```
   $map.REMOVE(k)$ 
```

```
   $a[i] \leftarrow a[i] + \Delta k$ 
```

```
   $map[a[i]] \leftarrow i$ 
```

```
  UP-HEAP( $a, 0, map$ )
```



d-Heaps



Introduction

d-Heaps

Application

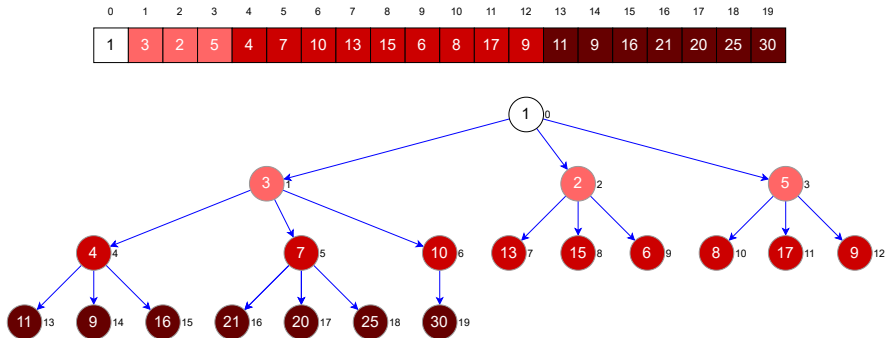
Simulation
Time-driven Simulation
Event-driven
Simulation

Workshop

Concept 4

a ***d*-heap** is a heap, each node of which has *d* children.

- A min 3-heap





Analysis

d-Heaps

Application

Workshop

- A d -heap is much shallower than a binary heap

implementation	insert	remove max	max
unordered array	1	N	N
ordered array	N	1	1
binary heap	$\log_2 N$	$\log_2 N$	1
d -heap	$\log_d N$	$\log_d N$	1



Application

- Simulation
- Time-driven Simulation
- Event-driven Simulation




Molecular dynamics simulation of hard discs

Goal. Simulate the motion of N moving particles that behave according to the laws of elastic collision.

Problem. N bouncing balls in the unit square.

- Moving particles interact via elastic collisions with each other and walls.
- Each particle is a disc with known position, velocity, mass, and radius.
- No other forces.



Bouncing balls




- Check for balls colliding with each other.
 - Physics problems: when? what effect?
 - CS problems: which object does the check? too many checks?




Time-driven simulation

Discretize time in quanta of size dt .


- Update the position of each particle after every dt units of time, and check for overlaps.
- If overlap, roll back the clock to the time of the collision, update the velocities of the colliding particles, and continue the simulation.




t



$t + dt$



$t + 2 dt$
(collision detected)



$t + \Delta t$
(roll back clock)




Time-driven simulation (cont.)


Main drawbacks

- $\sim N^2/2$ overlap checks per time quantum.
- Simulation is too slow if dt is very small.
- May miss collisions if dt is too large. (if colliding particles fail to overlap when we are looking)

dt too small: excessive computation



dt too large: may miss collisions






Event-driven simulation

Change **state** only when something happens.

- Between collisions, particles move in straight-line trajectories.
- Focus only on times when collisions occur.
- Maintain **PQ** of collision events, prioritized by time.
- Remove the min = get next collision.

Collision prediction. Given position, velocity, and radius of a particle. When will it collide next with a wall or another particle?

Collision resolution. If collision occurs, update colliding particle(s) according to laws of elastic collisions.






Particle-wall collision

Collision prediction and resolution.

- Particle of radius s at position (r_x, r_y) .
- Particle moving in unit box with velocity (v_x, v_y) .
- Will it collide with a vertical wall? If so, when?






Particle-particle collision prediction

Collision prediction

- Particle i : radius s_i , position (rx_i, ry_i) , velocity (vx_i, vy_i) .
- Particle j : radius s_j , position (rx_j, ry_j) , velocity (vx_j, vy_j) .



Particle-particle collision prediction (cont.)



- Will particles i and j collide? If so, when?

$$\Delta t = \begin{cases} \infty & \text{if } \Delta v \Delta r \geq 0 \\ \infty & \text{if } d < 0 \\ -\frac{\Delta v \Delta r}{\Delta v \Delta v} & \text{otherwise} \end{cases} \quad (1)$$

$$\sigma = \sigma_i + \sigma_j \quad (2)$$

$$d = (\Delta v \Delta r)^2 - (\Delta v \Delta v)(\Delta r \Delta r - \sigma^2) \quad (3)$$

where

$$\Delta v = (\Delta vx, \Delta vy) = (vx_i - vx_j, vy_i - vy_j) \quad (4)$$

$$\Delta r = (\Delta rx, \Delta ry) = (rx_i - rx_j, ry_i - ry_j) \quad (5)$$



Particle-particle collision resolution

Collision resolution

- When two particles collide, how does velocity change?

$$\begin{aligned}
 vx'_i &= vx_i + \frac{J_x}{m_i} \\
 vy'_i &= vy_i + \frac{J_y}{m_i} \\
 vx'_j &= vx_j - \frac{J_x}{m_j} \\
 vy'_j &= vy_j - \frac{J_y}{m_j}
 \end{aligned} \tag{6}$$

where

$$\begin{aligned}
 J &= \frac{2m_i m_j (\Delta v \Delta r)}{\sigma(m_i + m_j)} \\
 J_x &= \frac{J \Delta r_x}{\sigma} \\
 J_y &= \frac{J \Delta r_y}{\sigma}
 \end{aligned} \tag{7}$$

Collision system: event-driven simulation main loop



Initialization

- Fill PQ with all *potential* particle-wall collisions. (“potential” since collision may not happen if some other collision intervenes)
- Fill PQ with all *potential* particle-particle collisions.

Main loop

- Delete the impending event from PQ (min priority = t).
- If the event has been invalidated, ignore it.
- Advance all particles to time t , on a straight-line trajectory.
- Update the velocities of the colliding particle(s).
- Predict future particle-wall and particle-particle collisions involving the colliding particle(s) and insert events onto PQ.



Workshop

Binary Heaps

Insert Key

Remove Max

Increase Key

d-Heaps


Application

Simulation

Time-driven Simulation

Event-driven Simulation

Workshop




Quiz

1. What is a priority queue?

2. What is a binary heap?

3. What is a *d*-heap?



42



- Programming exercises in [[Cormen, 2009](#), [Sedgewick, 2002](#)]

References



Cormen, T. H. (2009).
Introduction to algorithms.
MIT press.



Sedgewick, R. (2002).
Algorithms in Java, Parts 1-4, volume 1.
Addison-Wesley Professional.



Walls and Mirrors (2014).
Data Abstraction And Problem Solving with C++.
Pearson.